2 Plan

Practice Lab Goals

Step-by-step guided learning

- → Create an an Amazon EC2 Auto Scaling group.
- → Assign EC2 instances to the Auto Scaling group.

DIY

Build on what you have learned.

- → Configure an auto scaling policy to scale down to 0 resources at 01:00 AM every day.

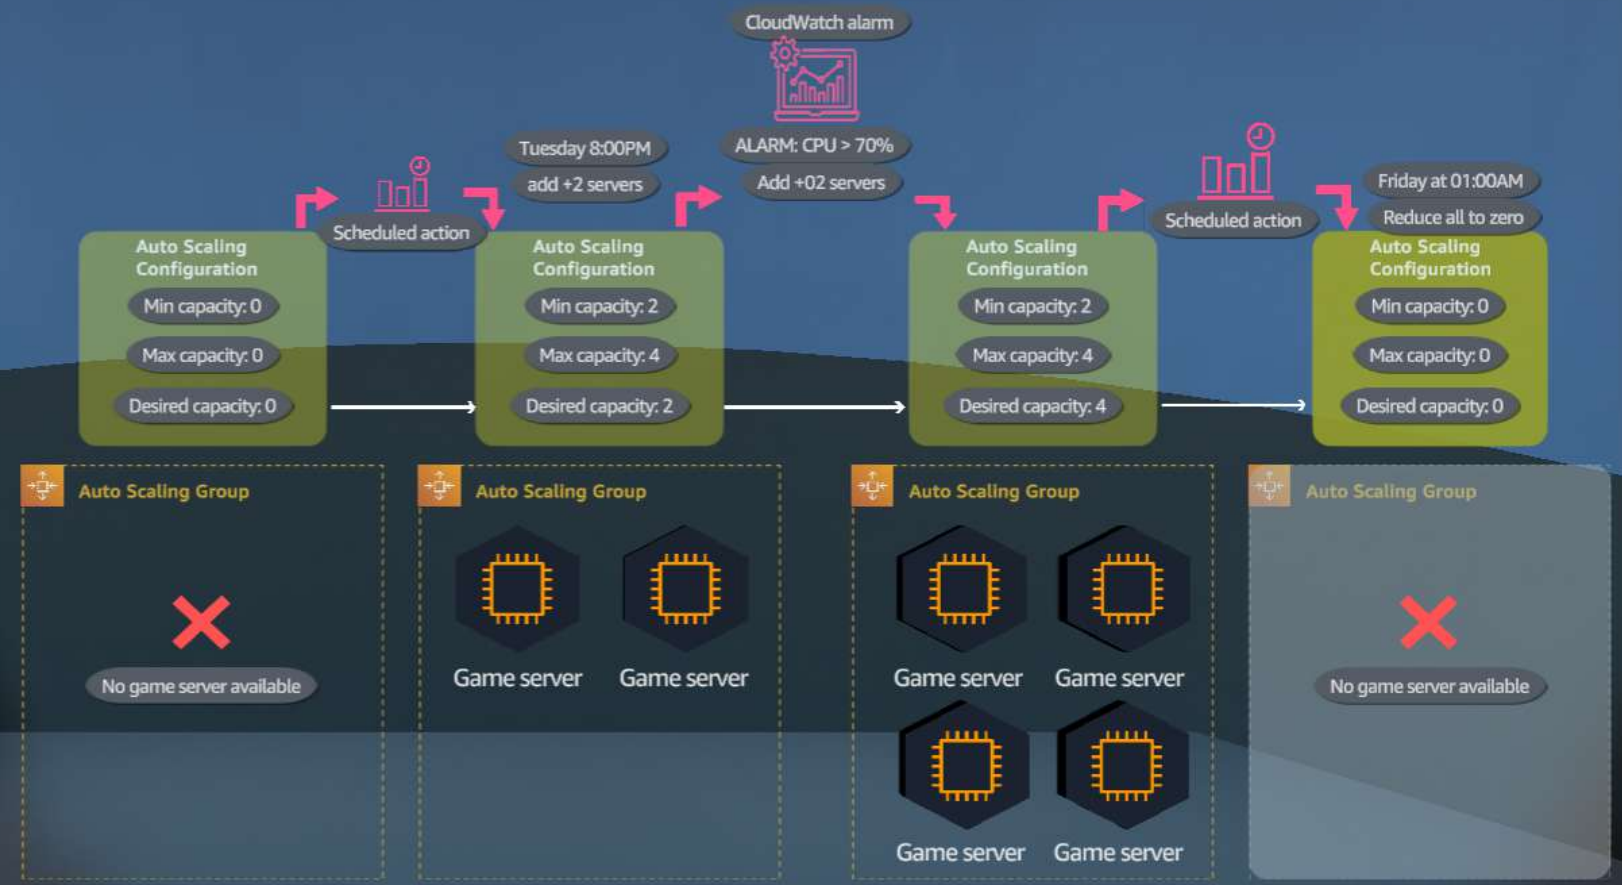

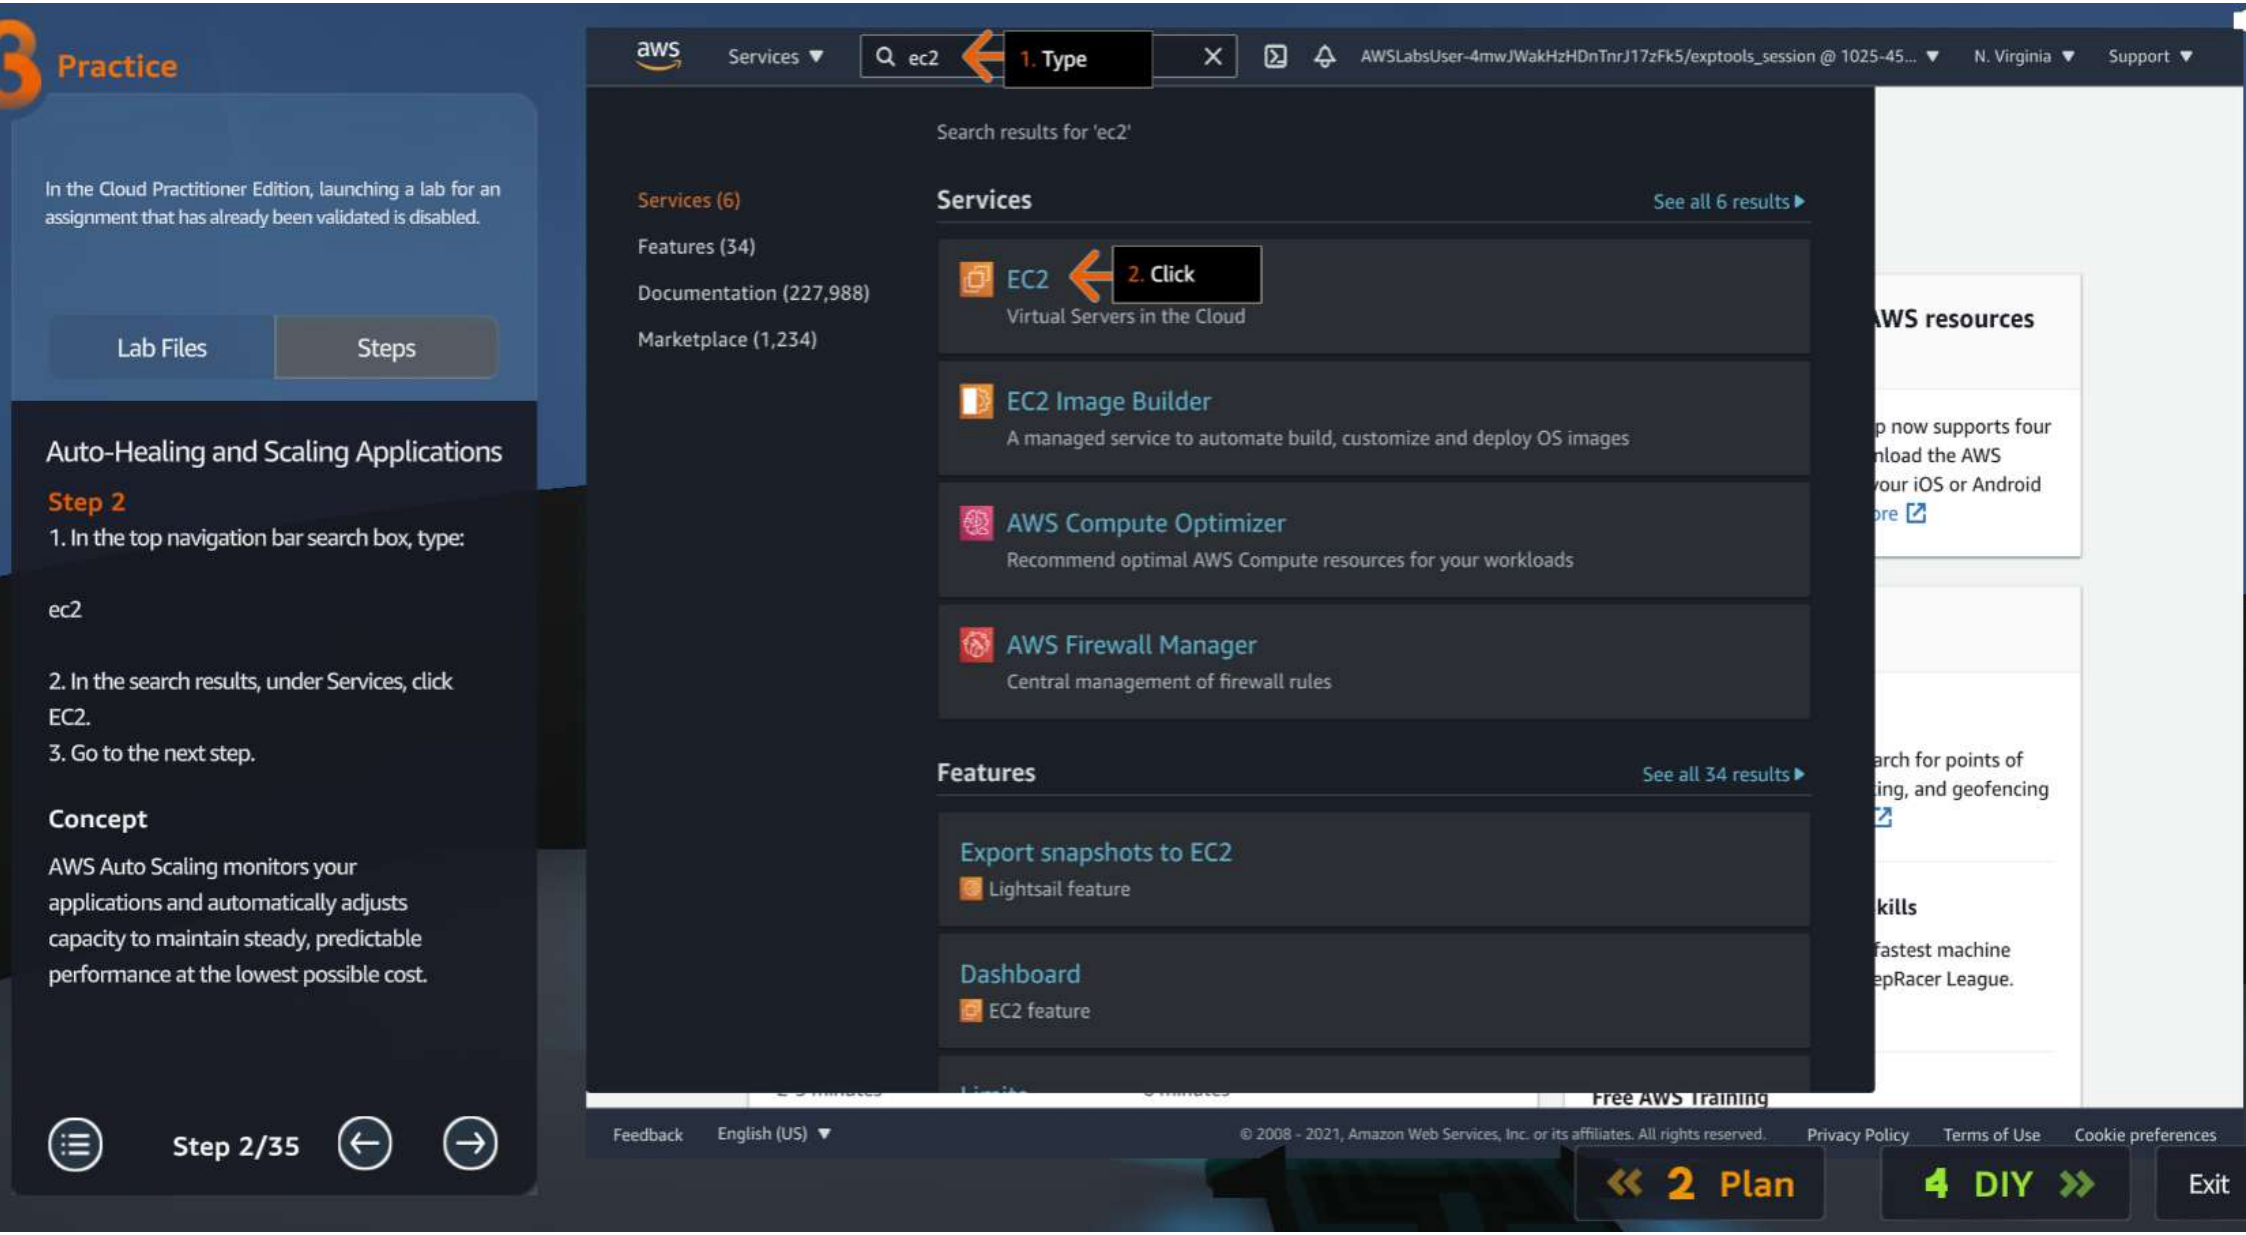
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 2
1. In the top navigation bar search box, type:
ec2
2. In the search results, under Services, click EC2.
3. Go to the next step.
Concept
AWS Auto Scaling monitors your applications and automatically adjusts capacity to maintain steady, predictable performance at the lowest possible cost.
Step 2/35
aws
Services
ec2
1. Type
AWSLabsUser-4mwJWakHzHDnTnrJ17zFk5/exptools_session @ 1025-45...
N. Virginia
Support
Search results for 'ec2'
Services (6)
Features (34)
Documentation (227,988)
Marketplace (1,234)
Services
See all 6 results
EC2
2. Click
Virtual Servers in the Cloud
EC2 Image Builder
A managed service to automate build, customize and deploy OS images
AWS Compute Optimizer
Recommend optimal AWS Compute resources for your workloads
AWS Firewall Manager
Central management of firewall rules
Features
See all 34 results
Export snapshots to EC2
Lightsail feature
Dashboard
EC2 feature
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
2 Plan
4 DIY
Exit

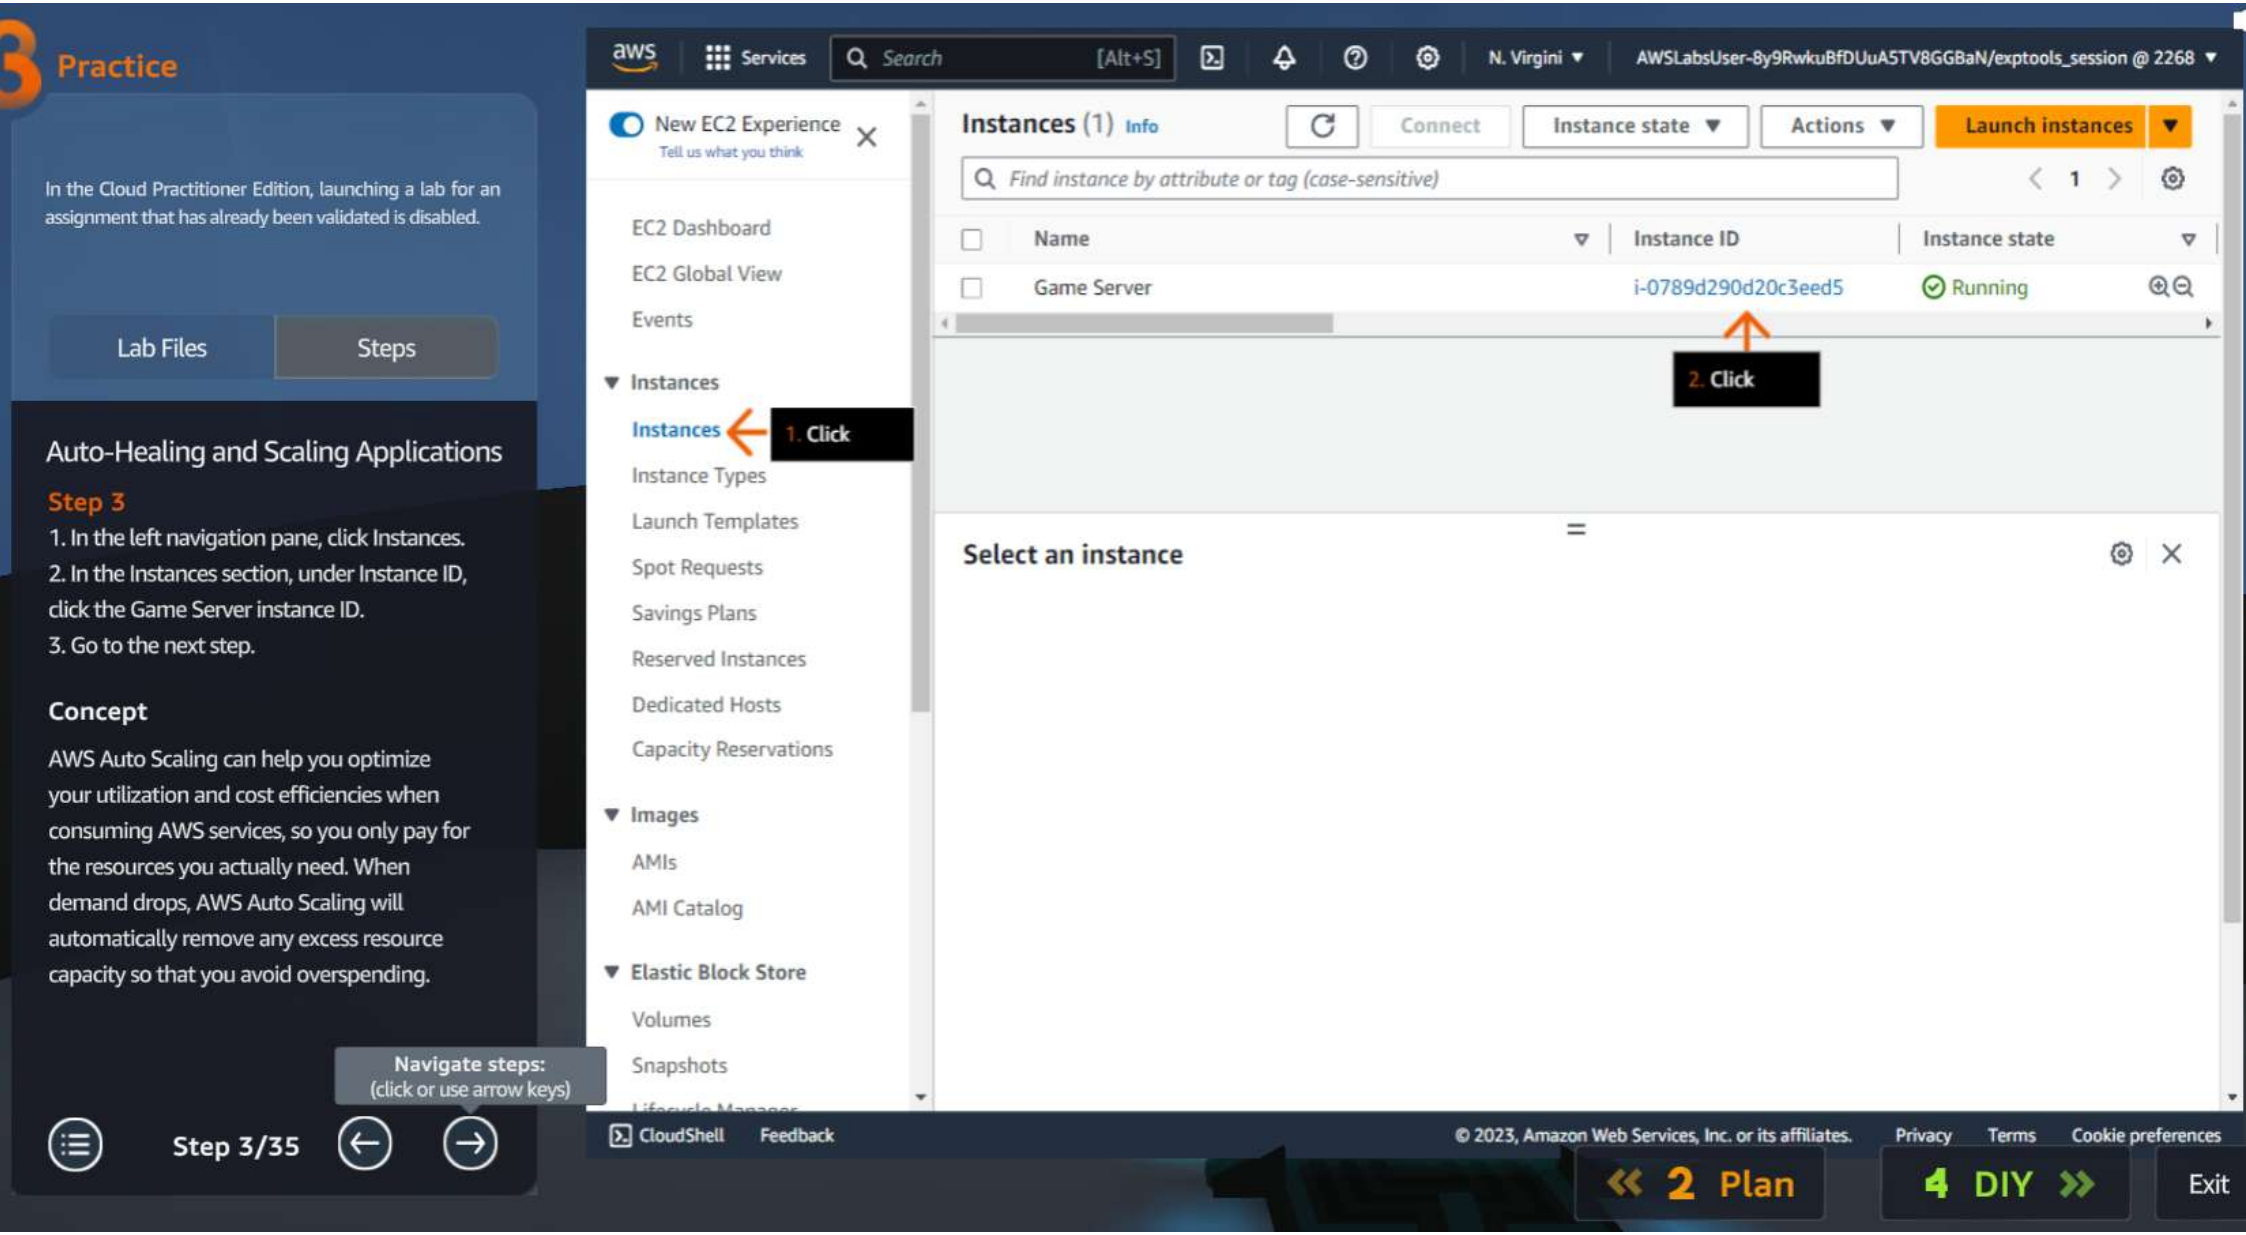
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 3
1. In the left navigation pane, click Instances.
2. In the Instances section, under Instance ID, click the Game Server instance ID.
3. Go to the next step.
Concept
AWS Auto Scaling can help you optimize your utilization and cost efficiencies when consuming AWS services, so you only pay for the resources you actually need. When demand drops, AWS Auto Scaling will automatically remove any excess resource capacity so that you avoid overspending.
Navigate steps: (click or use arrow keys)
Step 3/35
aws
Services
Search
[Alt+S]
N. Virgini
AWSLabsUser-8y9RwkuBfDUuA5TV8GGBaN/exptools_session @ 2268
New EC2 Experience
Tell us what you think
EC2 Dashboard
EC2 Global View
Events
Instances
Instances
1. Click
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Instances (1) Info
Connect
Instance state
Actions
Launch instances
Find instance by attribute or tag (case-sensitive)
Name
Instance ID
Instance state
Game Server
i-0789d290d20c3eed5
Running
2. Click
Select an instance
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
« 2 Plan
4 DIY »
Exit

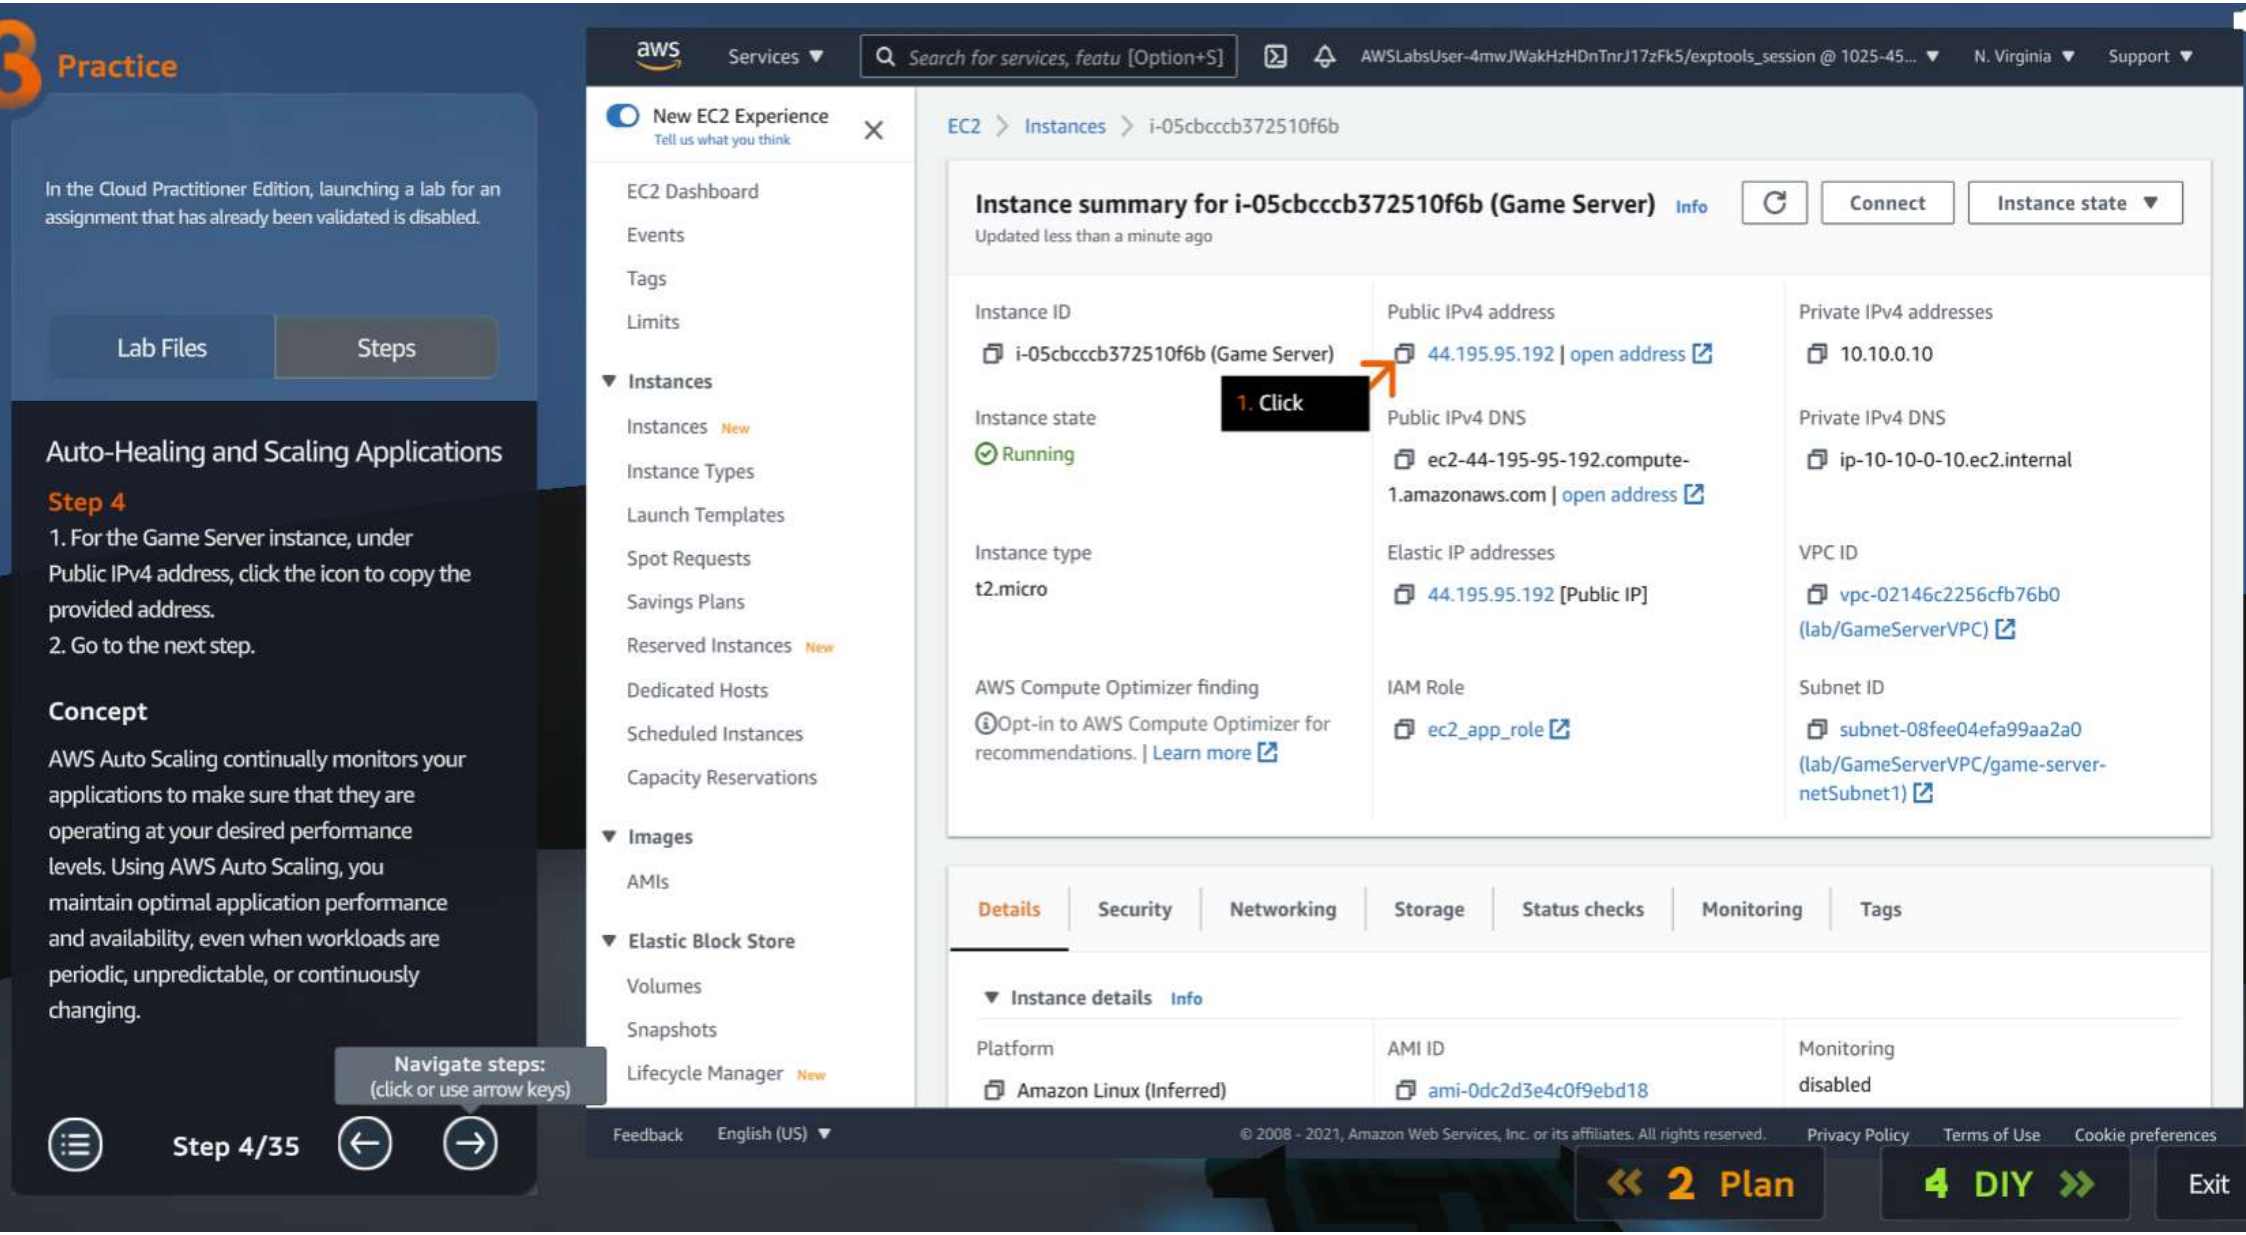
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 4
1. For the Game Server instance, under Public IPv4 address, click the icon to copy the provided address.
2. Go to the next step.
Concept
AWS Auto Scaling continually monitors your applications to make sure that they are operating at your desired performance levels. Using AWS Auto Scaling, you maintain optimal application performance and availability, even when workloads are periodic, unpredictable, or continuously changing.
Navigate steps: (click or use arrow keys)
Step 4/35
aws
Services
Search for services, featu [Option+S]
AWSLabsUser-4mwJWakHzHDnTnrJ17zFk5/exptools_session @ 1025-45...
N. Virginia
Support
New EC2 Experience
Tell us what you think
EC2 Dashboard
Events
Tags
Limits
Instances
Instances New
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances New
Dedicated Hosts
Scheduled Instances
Capacity Reservations
Images
AMIs
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager New
EC2 > Instances > i-05cbcccb372510f6b
Instance summary for i-05cbcccb372510f6b (Game Server) Info
Updated less than a minute ago
Connect
Instance state
Instance ID
i-05cbcccb372510f6b (Game Server)
1. Click
Public IPv4 address
44.195.95.192 | open address
Private IPv4 addresses
10.10.0.10
Instance state
Running
Public IPv4 DNS
ec2-44-195-95-192.compute-1.amazonaws.com | open address
Private IPv4 DNS
ip-10-10-0-10.ec2.internal
Instance type
t2.micro
Elastic IP addresses
44.195.95.192 [Public IP]
VPC ID
vpc-02146c2256cfb76b0 (lab/GameServerVPC)
AWS Compute Optimizer finding
Opt-in to AWS Compute Optimizer for recommendations. | Learn more
IAM Role
ec2_app_role
Subnet ID
subnet-08fee04efa99aa2a0 (lab/GameServerVPC/game-server-netSubnet1)
Details
Security
Networking
Storage
Status checks
Monitoring
Tags
Instance details Info
Platform
Amazon Linux (Inferred)
AMI ID
ami-0dc2d3e4c0f9ebd18
Monitoring
disabled
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

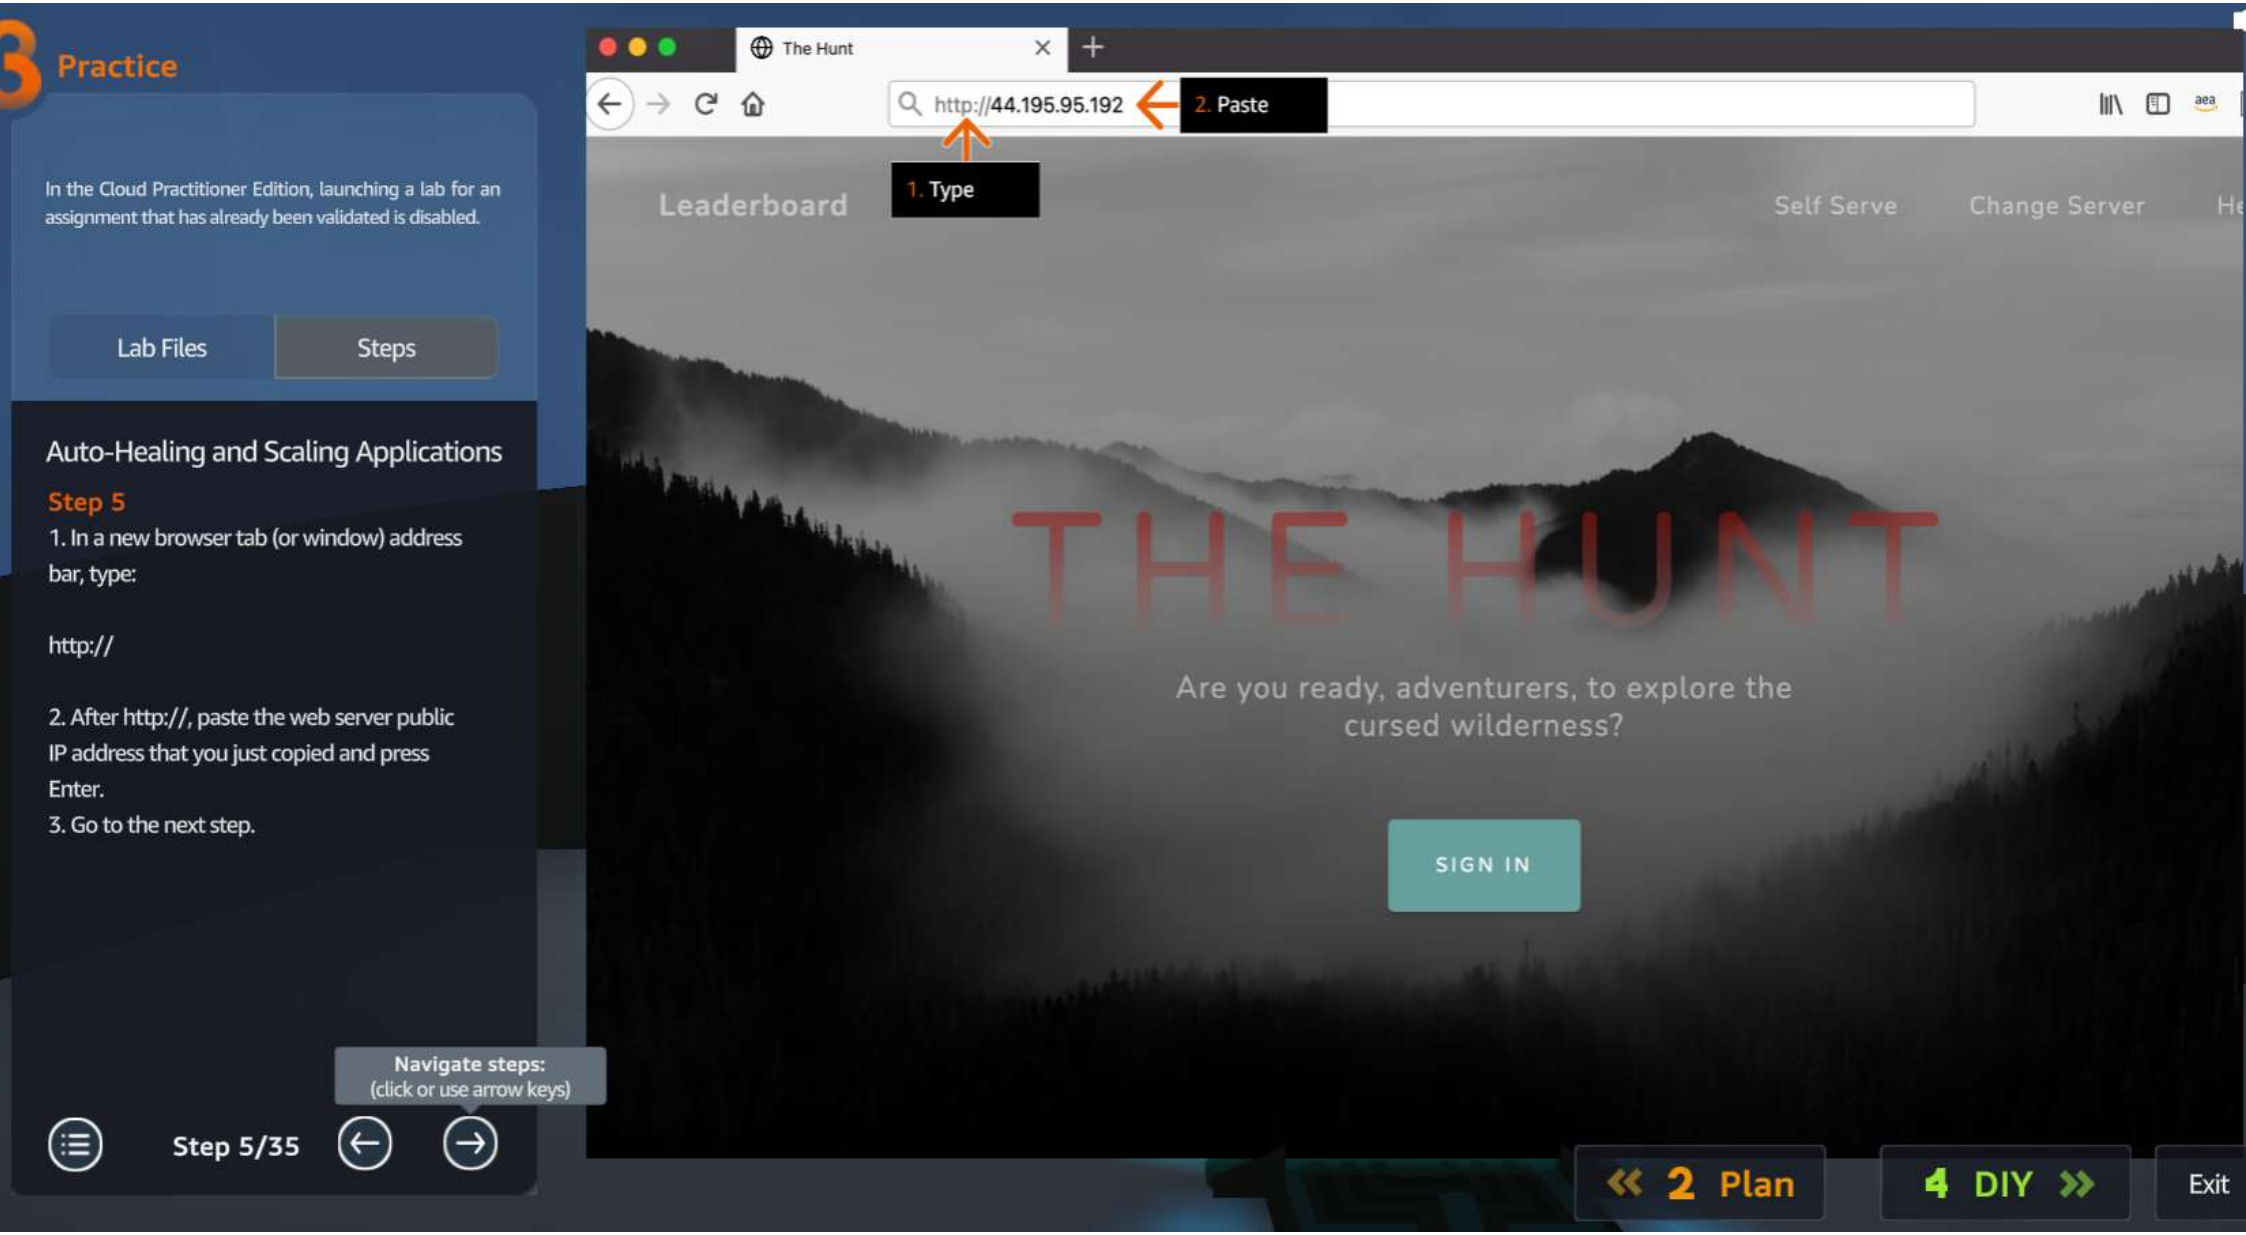
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 5
1. In a new browser tab (or window) address bar, type:
http://
2. After http://, paste the web server public IP address that you just copied and press Enter.
3. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 5/35
The Hunt
http://44.195.95.192
2. Paste
1. Type
Leaderboard
Self Serve
Change Server
THE HUNT
Are you ready, adventurers, to explore the cursed wilderness?
SIGN IN
2 Plan
4 DIY
Exit

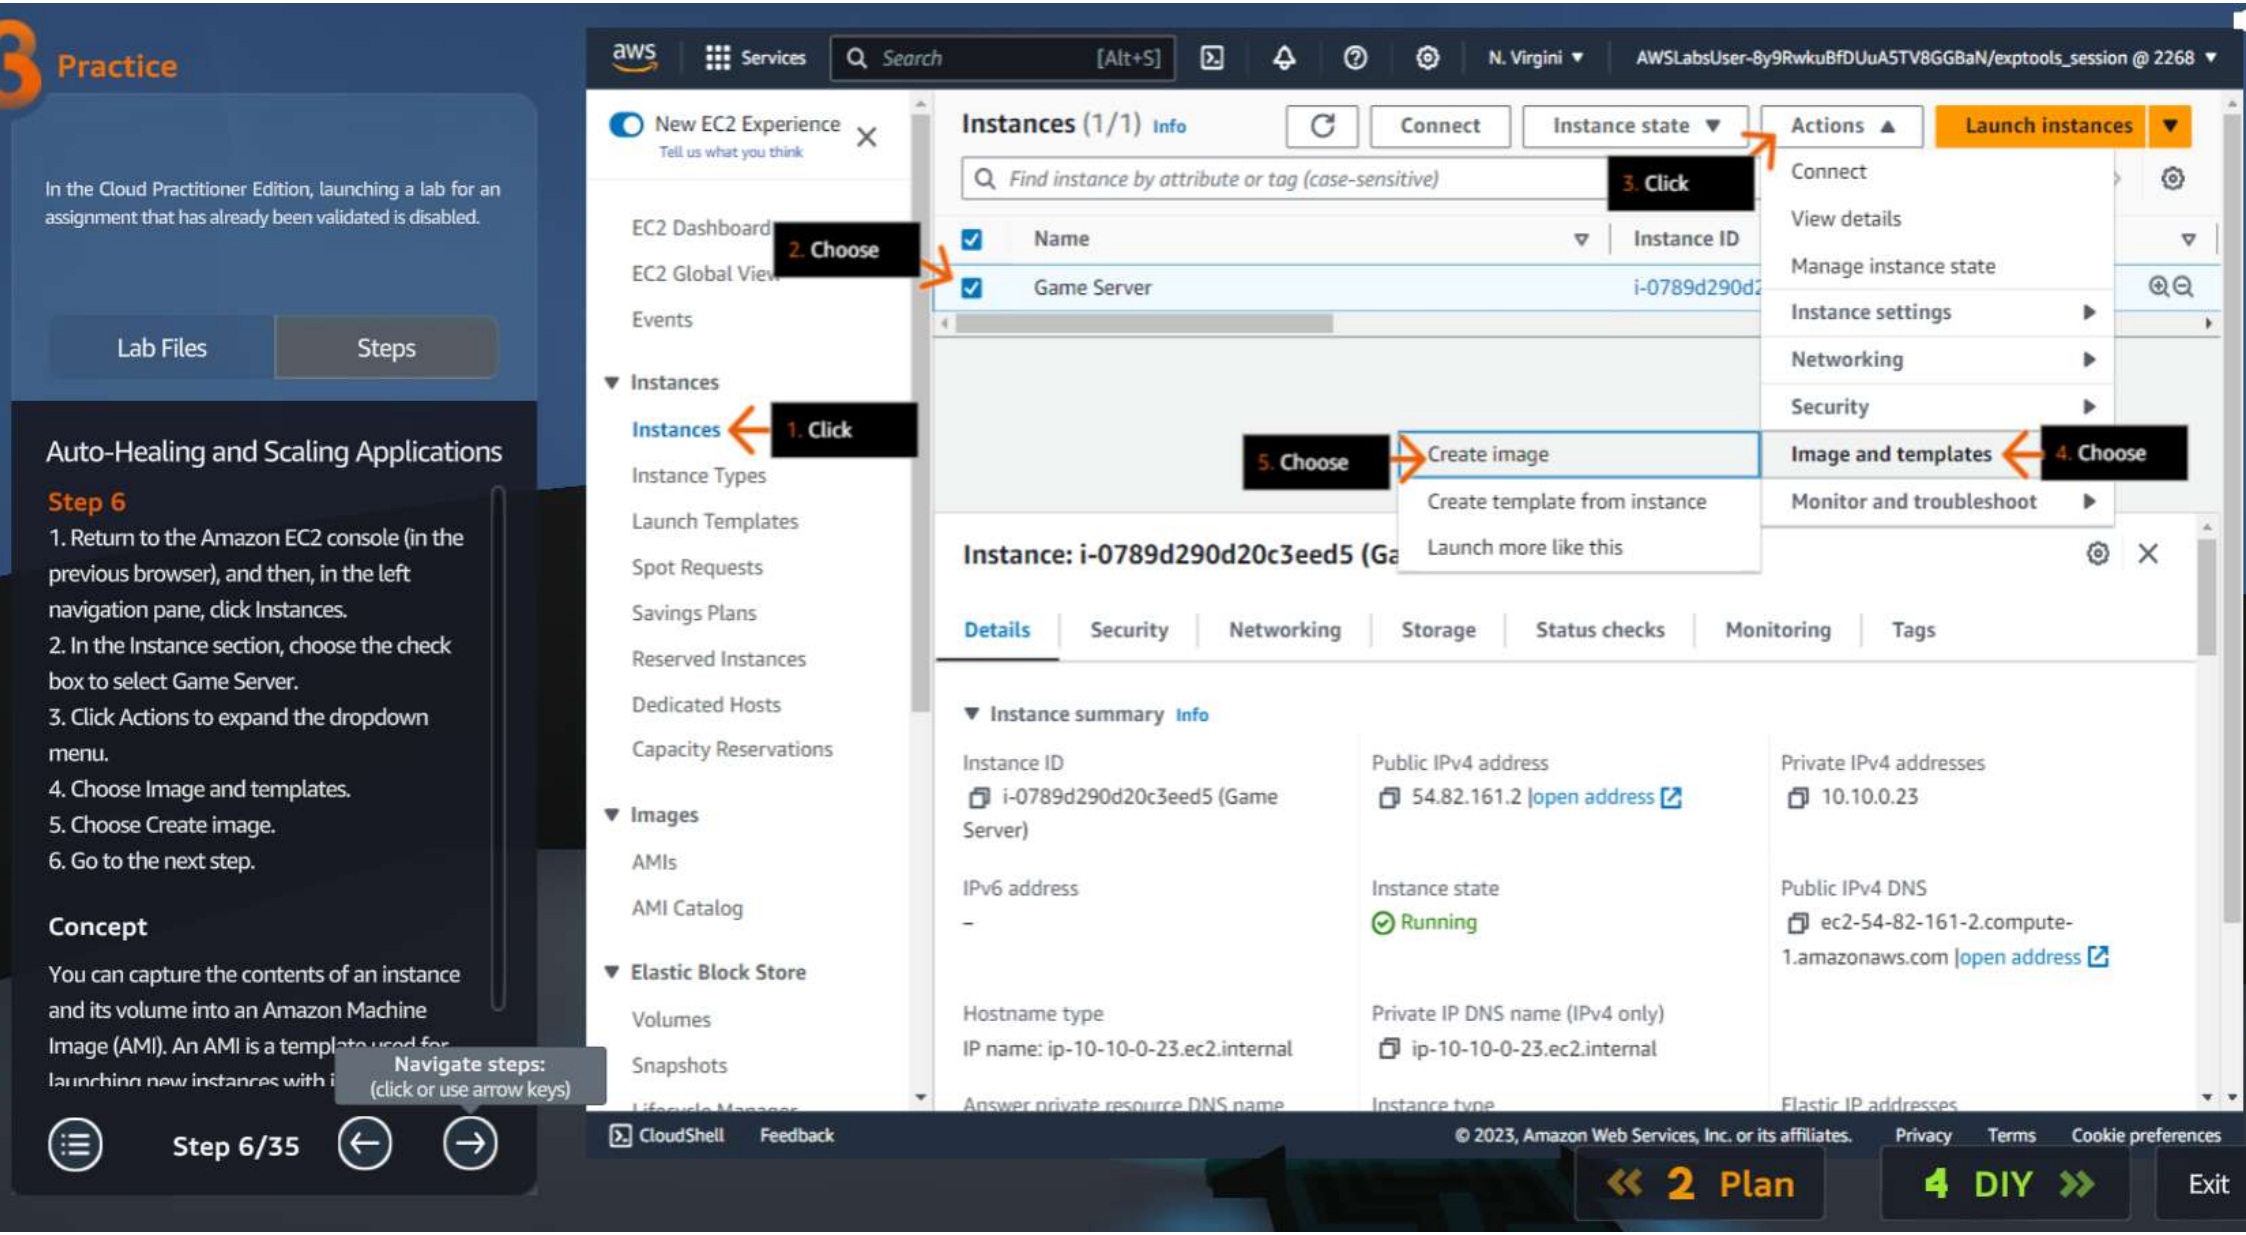
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 6
1. Return to the Amazon EC2 console (in the previous browser), and then, in the left navigation pane, click Instances.
2. In the Instance section, choose the check box to select Game Server.
3. Click Actions to expand the dropdown menu.
4. Choose Image and templates.
5. Choose Create image.
6. Go to the next step.
Concept
You can capture the contents of an instance and its volume into an Amazon Machine Image (AMI). An AMI is a template used for
launching new instances with i
Navigate steps: (click or use arrow keys)
Step 6/35
aws
Services
Search
[Alt+S]
N. Virgini
AWSLabsUser-8y9RwkuBfDUuA5TV8GGBaN/exptools_session @ 2268
New EC2 Experience
Tell us what you think
EC2 Dashboard
EC2 Global View
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
1. Click
2. Choose
Instances (1/1) Info
Connect
Instance state
Actions
Launch instances
Find instance by attribute or tag (case-sensitive)
3. Click
Name
Instance ID
Game Server
i-0789d290d2
Connect
View details
Manage instance state
Instance settings
Networking
Security
Image and templates
Monitor and troubleshoot
4. Choose
5. Choose
Create image
Create template from instance
Launch more like this
Instance: i-0789d290d20c3eed5 (Ga
Details
Security
Networking
Storage
Status checks
Monitoring
Tags
Instance summary Info
Instance ID
i-0789d290d20c3eed5 (Game Server)
Public IPv4 address
54.82.161.2 |open address
Private IPv4 addresses
10.10.0.23
IPv6 address
–
Instance state
Running
Public IPv4 DNS
ec2-54-82-161-2.compute-1.amazonaws.com |open address
Hostname type
IP name: ip-10-10-0-23.ec2.internal
Private IP DNS name (IPv4 only)
ip-10-10-0-23.ec2.internal
Answer private resource DNS name
Instance type
Elastic IP addresses
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
« 2 Plan
4 DIY »
Exit

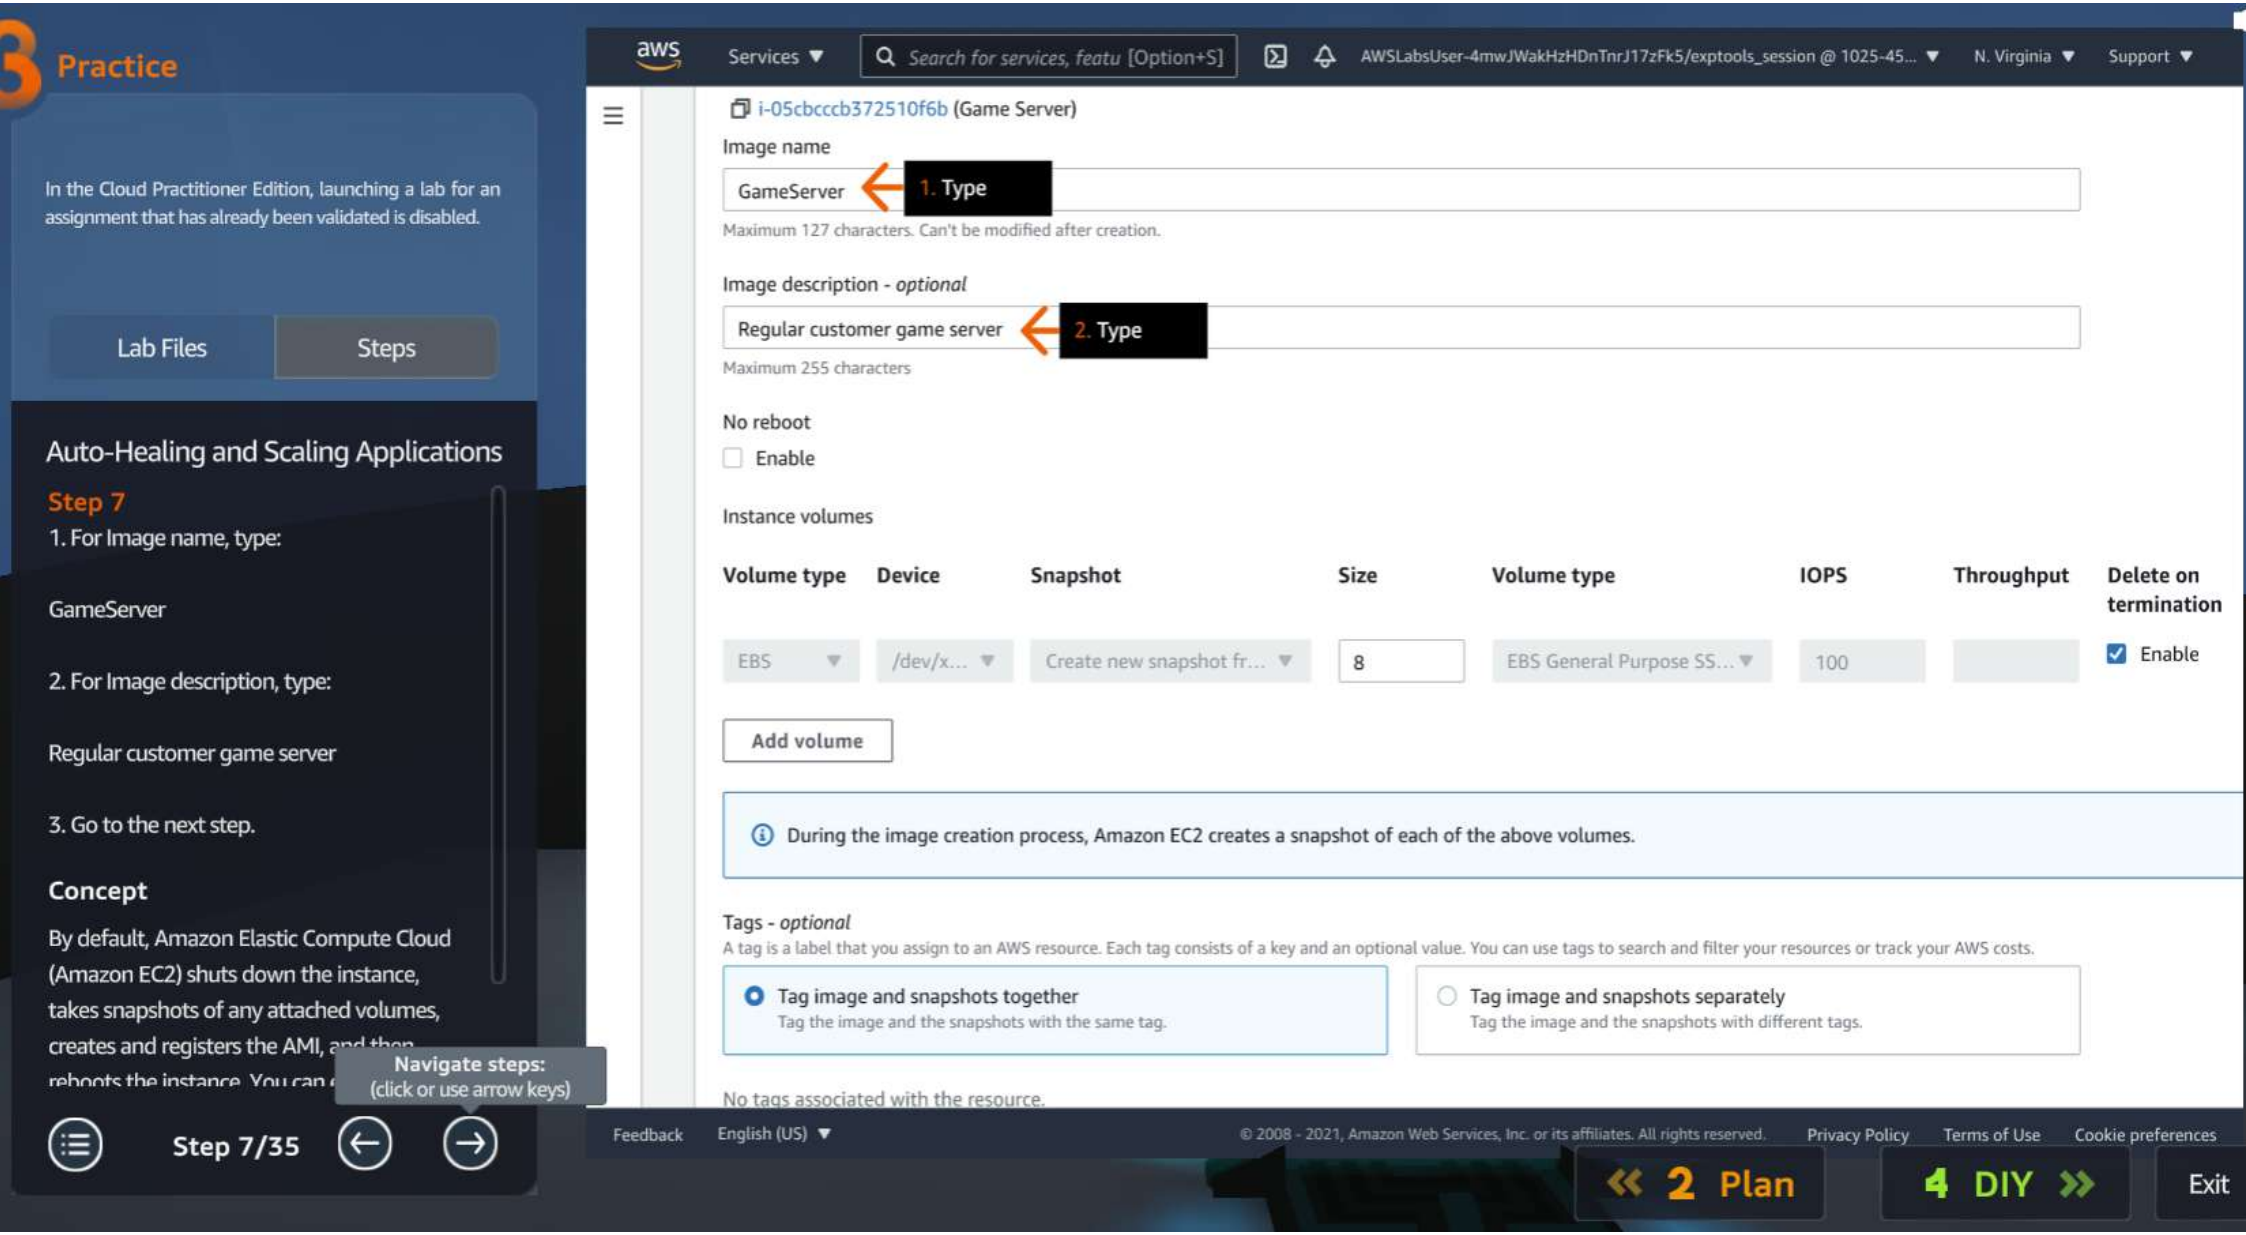
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 7
1. For Image name, type:
GameServer
2. For Image description, type:
Regular customer game server
3. Go to the next step.
Concept
By default, Amazon Elastic Compute Cloud (Amazon EC2) shuts down the instance, takes snapshots of any attached volumes, creates and registers the AMI, and then
Navigate steps:
(click or use arrow keys)
Step 7/35
aws
Services
Search for services, featu [Option+S]
AWSLabsUser-4mwJWakHzHDnTnrJ17zFk5/exptools_session @ 1025-45...
N. Virginia
Support
i-05cbcccb372510f6b (Game Server)
Image name
GameServer
1. Type
Maximum 127 characters. Can't be modified after creation.
Image description - optional
Regular customer game server
2. Type
Maximum 255 characters
No reboot
Enable
Instance volumes
Volume type
Device
Snapshot
Size
Volume type
IOPS
Throughput
Delete on termination
EBS
/dev/x...
Create new snapshot fr...
8
EBS General Purpose SS...
100
Enable
Add volume
During the image creation process, Amazon EC2 creates a snapshot of each of the above volumes.
Tags - optional
A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.
Tag image and snapshots together
Tag the image and the snapshots with the same tag.
Tag image and snapshots separately
Tag the image and the snapshots with different tags.
No tags associated with the resource.
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

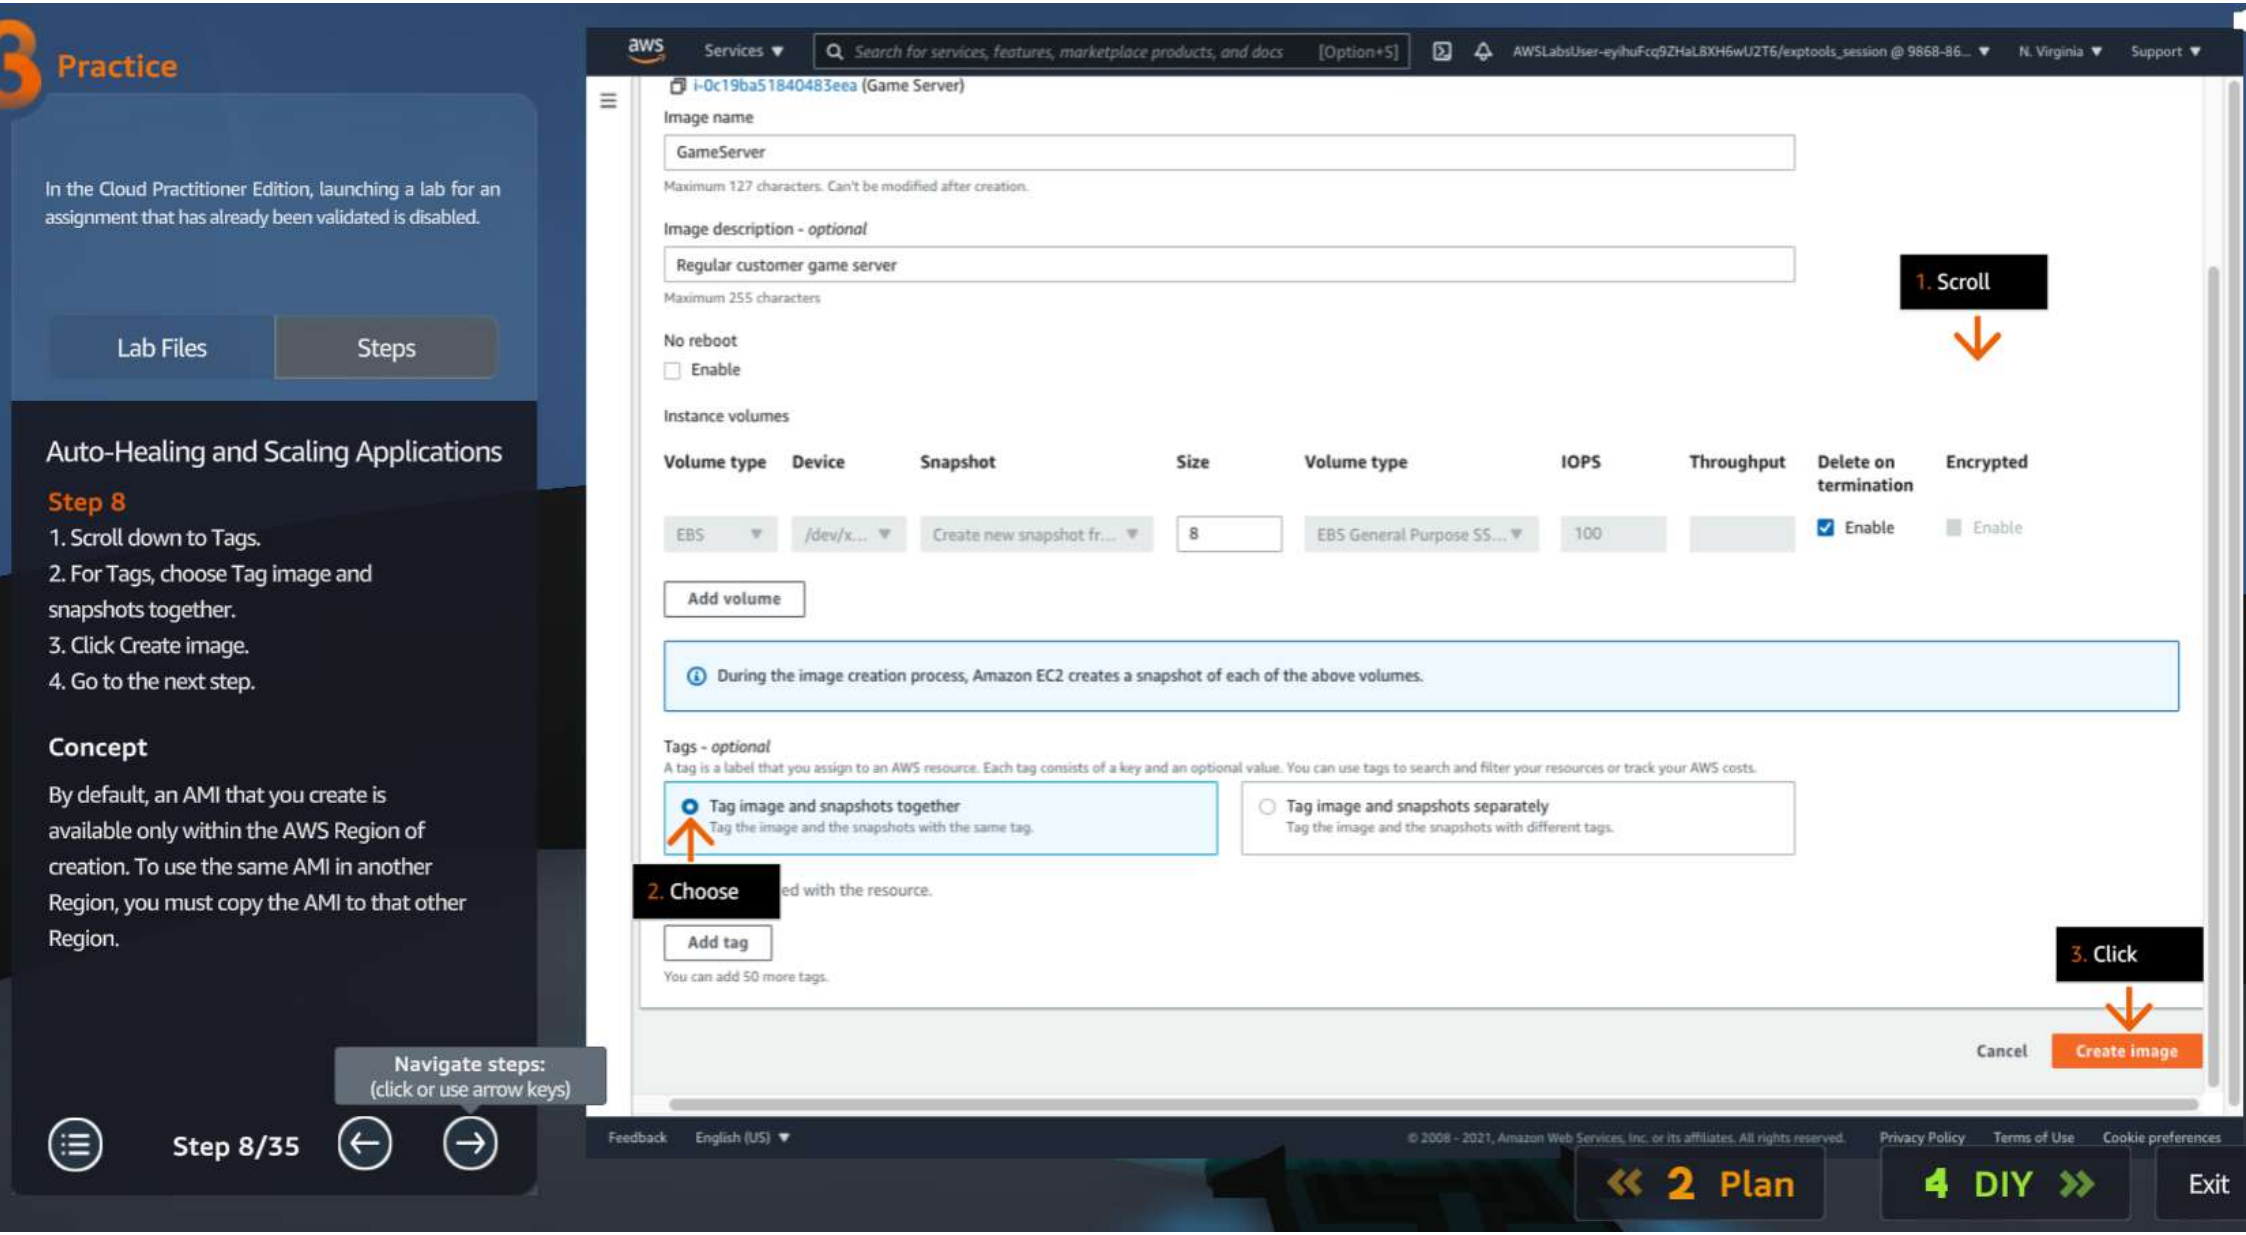
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 8
1. Scroll down to Tags.
2. For Tags, choose Tag image and snapshots together.
3. Click Create image.
4. Go to the next step.
Concept
By default, an AMI that you create is available only within the AWS Region of creation. To use the same AMI in another Region, you must copy the AMI to that other Region.
Navigate steps: (click or use arrow keys)
Step 8/35
aws
Services
Search for services, features, marketplace products, and docs [Option+S]
N. Virginia
Support
i-0c19ba51840483eea (Game Server)
Image name
GameServer
Maximum 127 characters. Can't be modified after creation.
Image description - optional
Regular customer game server
Maximum 255 characters
1. Scroll
No reboot
Enable
Instance volumes
Volume type
Device
Snapshot
Size
Volume type
IOPS
Throughput
Delete on termination
Encrypted
EBS
/dev/x...
Create new snapshot fr...
8
EBS General Purpose SS...
100
Enable
Enable
Add volume
During the image creation process, Amazon EC2 creates a snapshot of each of the above volumes.
Tags - optional
A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.
Tag image and snapshots together
Tag the image and the snapshots with the same tag.
Tag image and snapshots separately
Tag the image and the snapshots with different tags.
2. Choose
ed with the resource.
Add tag
You can add 50 more tags.
3. Click
Cancel
Create image
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

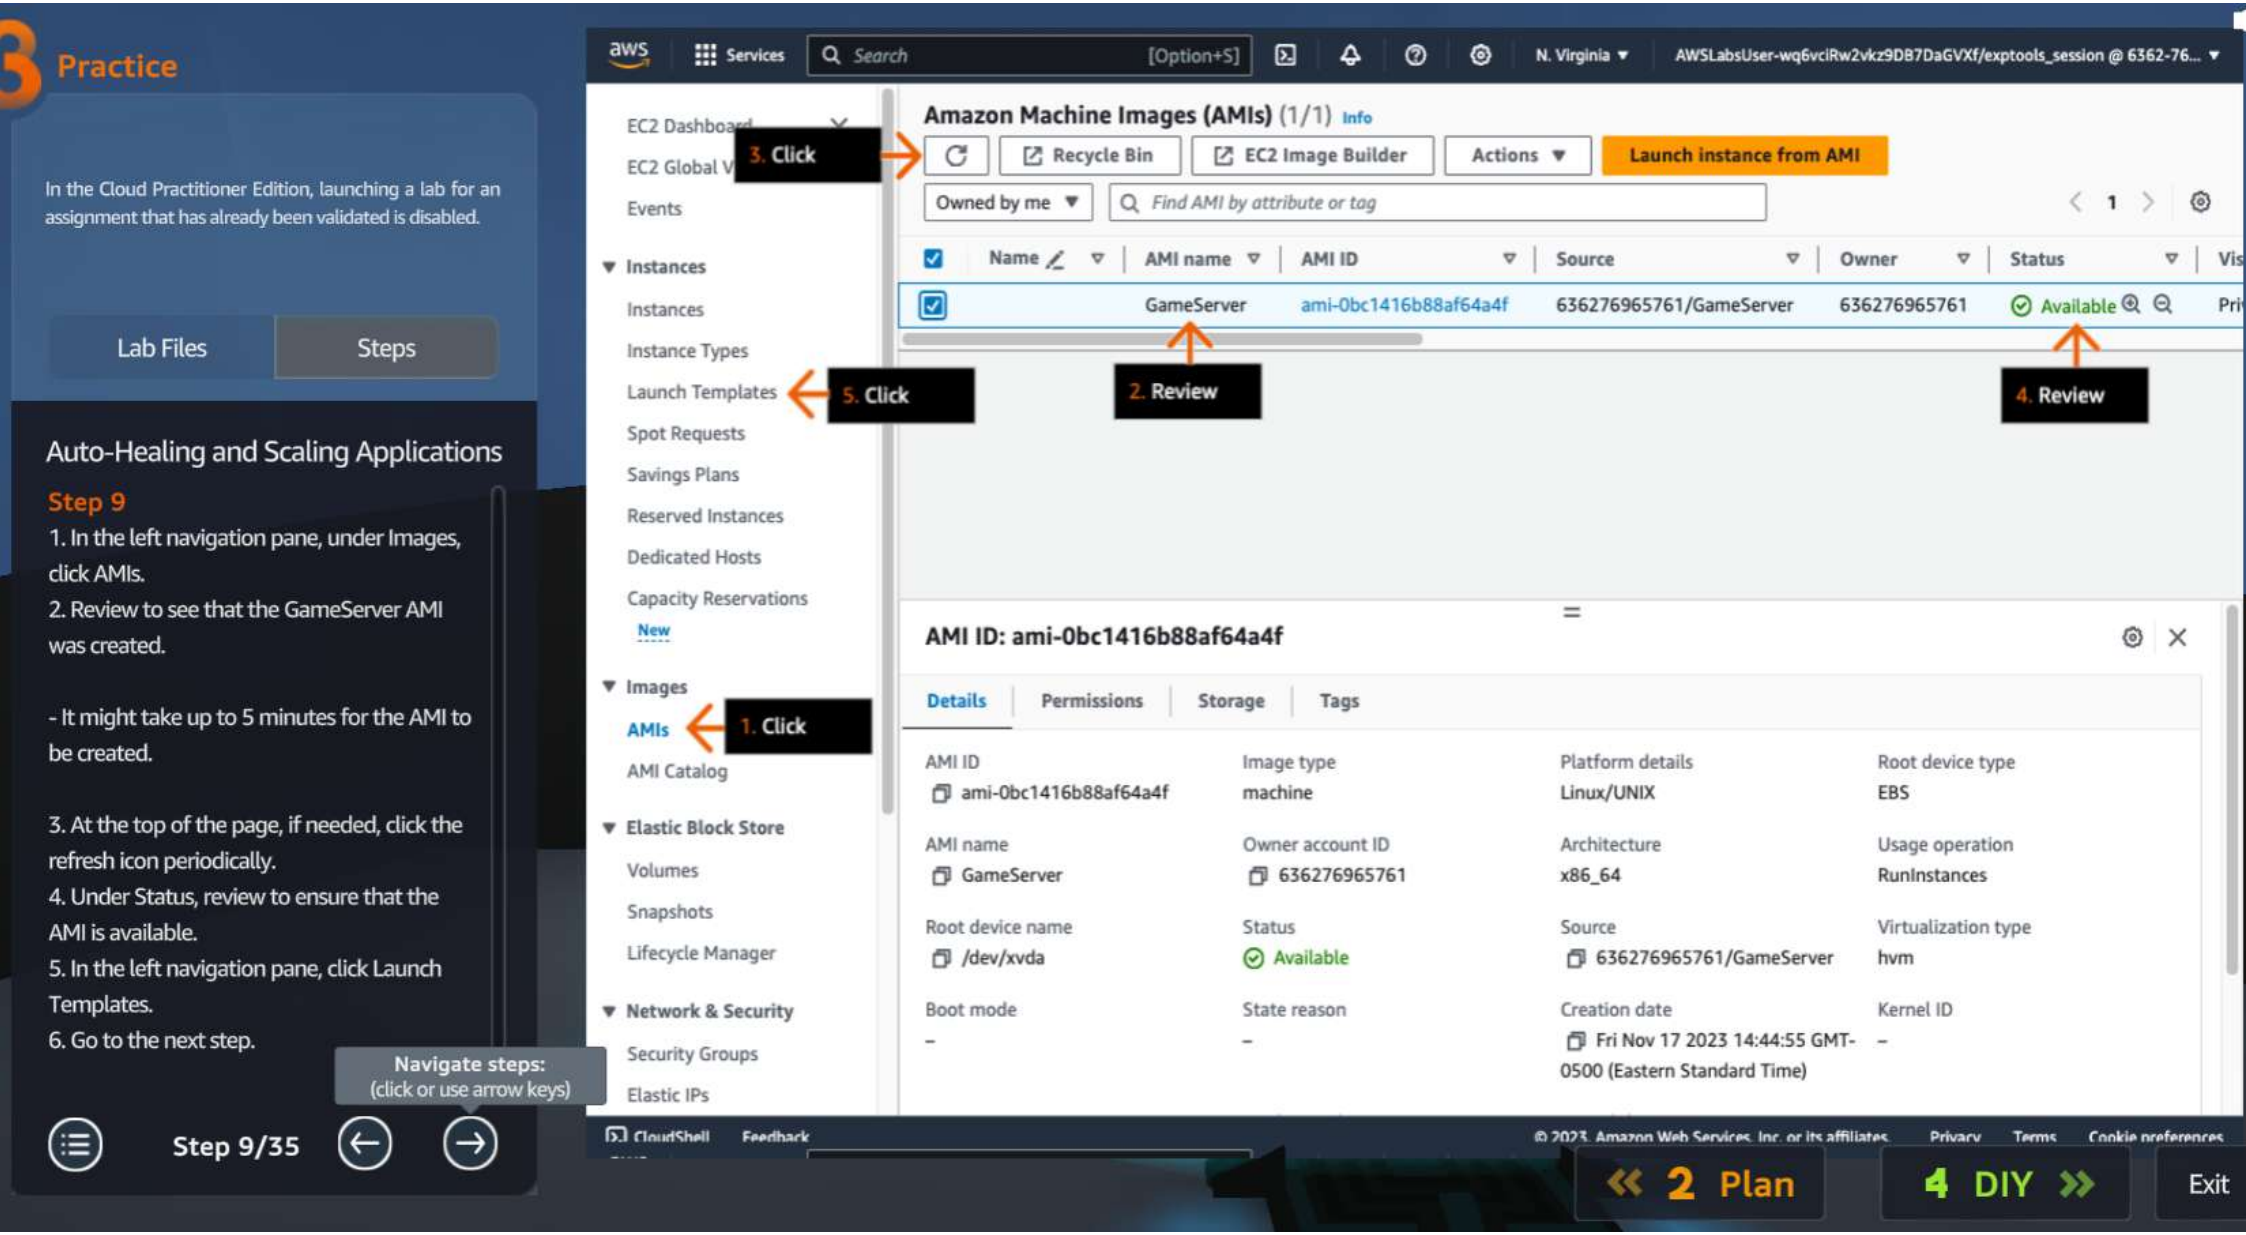
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 9
1. In the left navigation pane, under Images, click AMIs.
2. Review to see that the GameServer AMI was created.
- It might take up to 5 minutes for the AMI to be created.
3. At the top of the page, if needed, click the refresh icon periodically.
4. Under Status, review to ensure that the AMI is available.
5. In the left navigation pane, click Launch Templates.
6. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 9/35
aws
Services
Search
[Option+S]
N. Virginia
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 6362-76...
EC2 Dashboard
EC2 Global V
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
New
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
3. Click
5. Click
1. Click
Amazon Machine Images (AMIs) (1/1) Info
Recycle Bin
EC2 Image Builder
Actions
Launch instance from AMI
Owned by me
Find AMI by attribute or tag
Name
AMI name
AMI ID
Source
Owner
Status
GameServer
ami-0bc1416b88af64a4f
636276965761/GameServer
636276965761
Available
2. Review
4. Review
AMI ID: ami-0bc1416b88af64a4f
Details
Permissions
Storage
Tags
AMI ID
ami-0bc1416b88af64a4f
Image type
machine
Platform details
Linux/UNIX
Root device type
EBS
AMI name
GameServer
Owner account ID
636276965761
Architecture
x86_64
Usage operation
RunInstances
Root device name
/dev/xvda
Status
Available
Source
636276965761/GameServer
Virtualization type
hvm
Boot mode
-
State reason
-
Creation date
Fri Nov 17 2023 14:44:55 GMT-0500 (Eastern Standard Time)
Kernel ID
-
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
« 2 Plan
4 DIY »
Exit

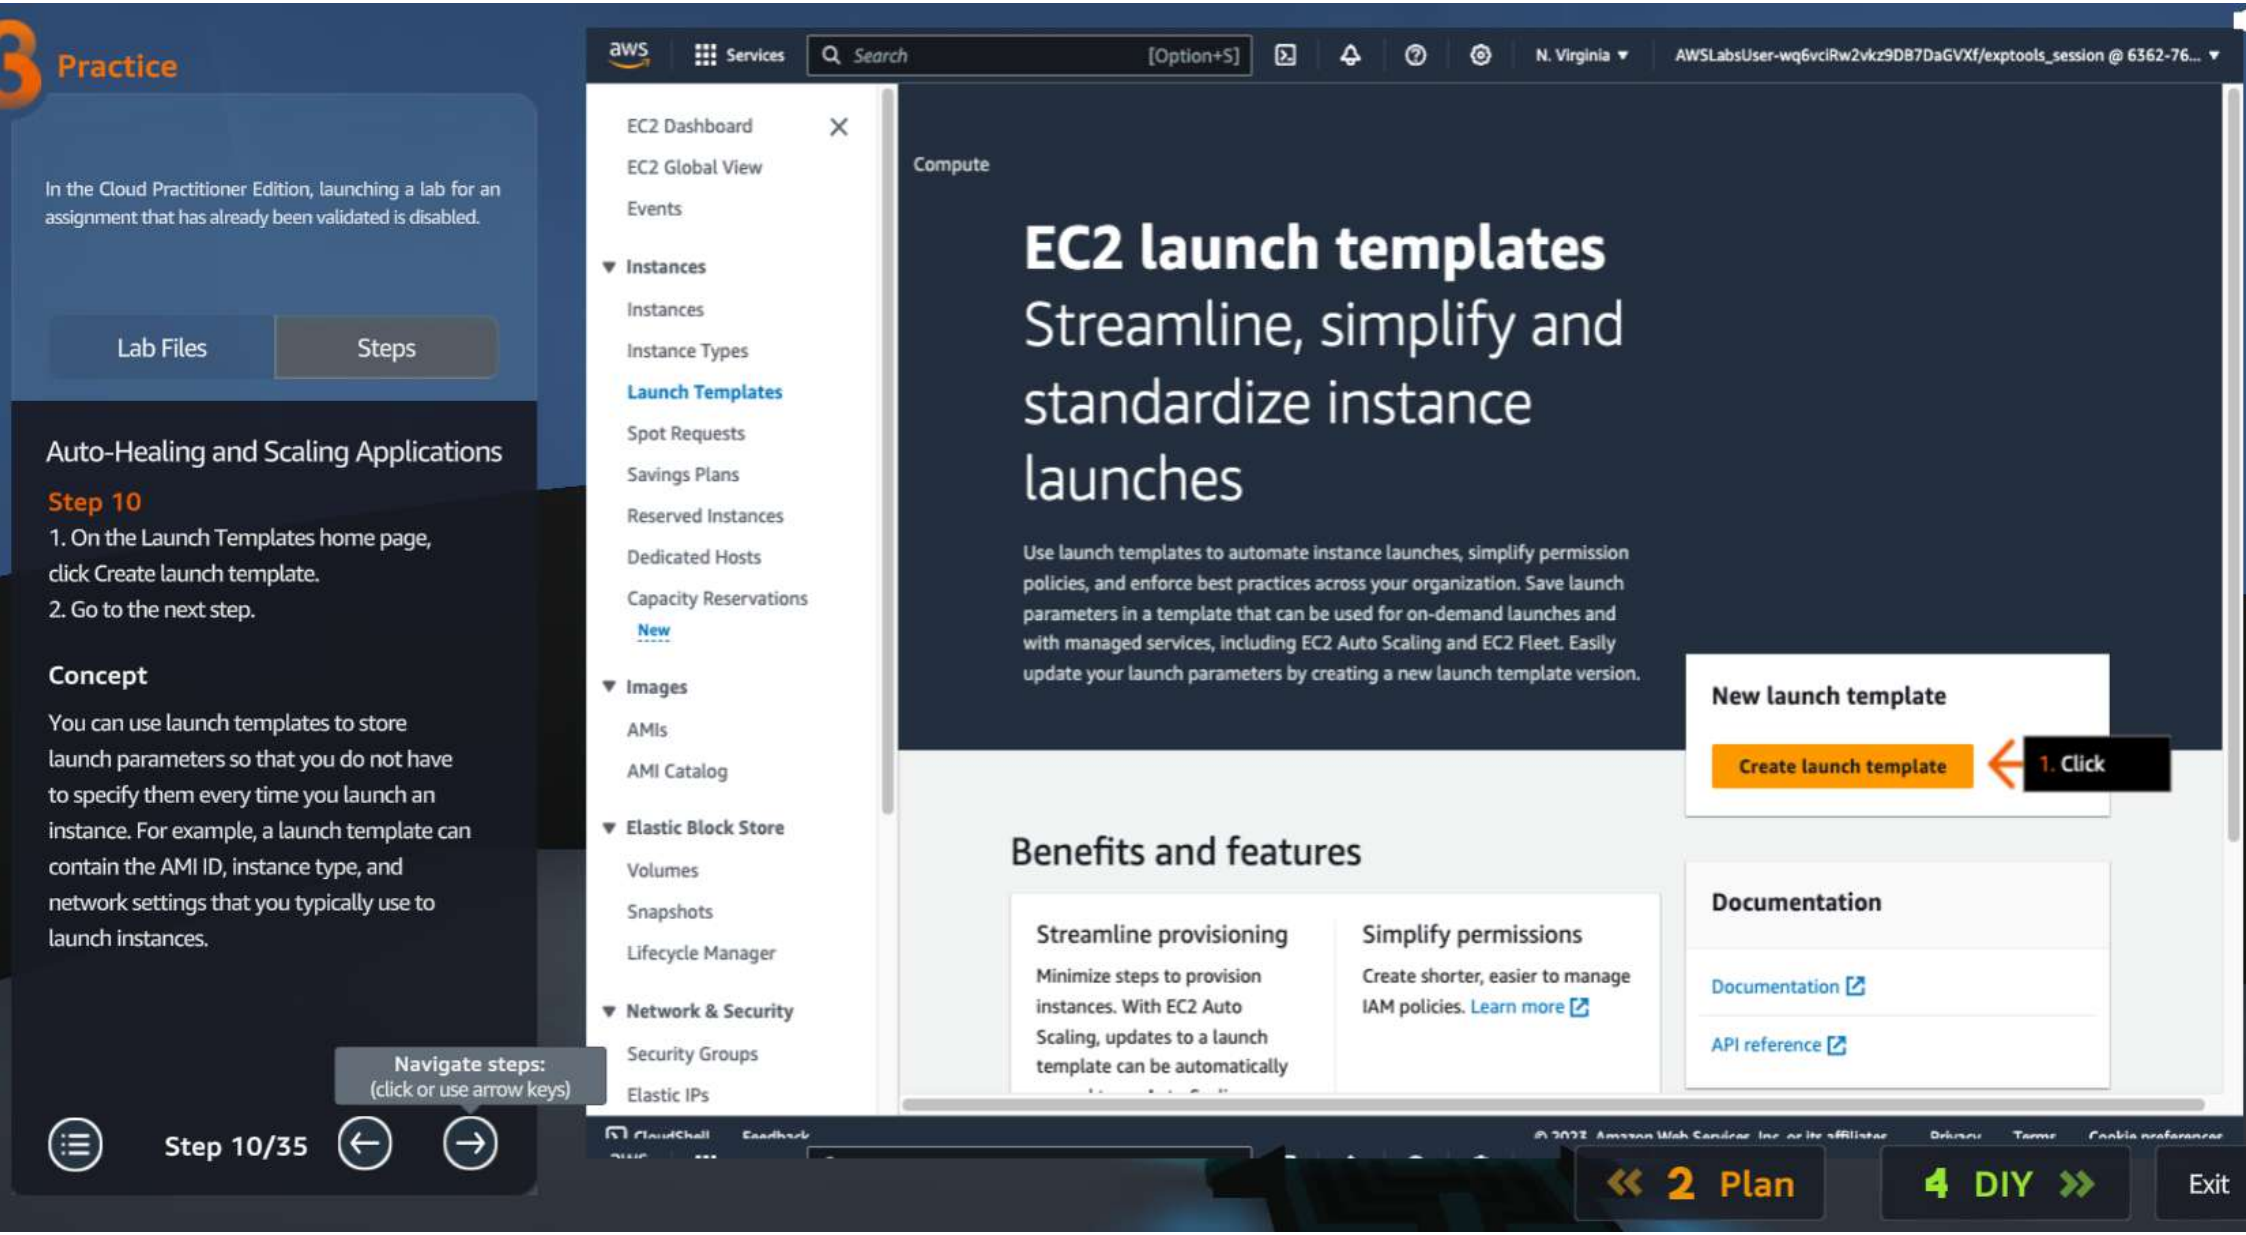
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 10
1. On the Launch Templates home page, click Create launch template.
2. Go to the next step.
Concept
You can use launch templates to store launch parameters so that you do not have to specify them every time you launch an instance. For example, a launch template can contain the AMI ID, instance type, and network settings that you typically use to launch instances.
Navigate steps: (click or use arrow keys)
Step 10/35
aws
Services
Search
[Option+S]
N. Virginia
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 6362-76...
EC2 Dashboard
EC2 Global View
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
New
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Compute
EC2 launch templates
Streamline, simplify and standardize instance launches
Use launch templates to automate instance launches, simplify permission policies, and enforce best practices across your organization. Save launch parameters in a template that can be used for on-demand launches and with managed services, including EC2 Auto Scaling and EC2 Fleet. Easily update your launch parameters by creating a new launch template version.
New launch template
Create launch template
1. Click
Benefits and features
Streamline provisioning
Minimize steps to provision instances. With EC2 Auto Scaling, updates to a launch template can be automatically
Simplify permissions
Create shorter, easier to manage IAM policies. Learn more
Documentation
Documentation
API reference
CloudShell
Feedback
« 2 Plan
4 DIY »
Exit

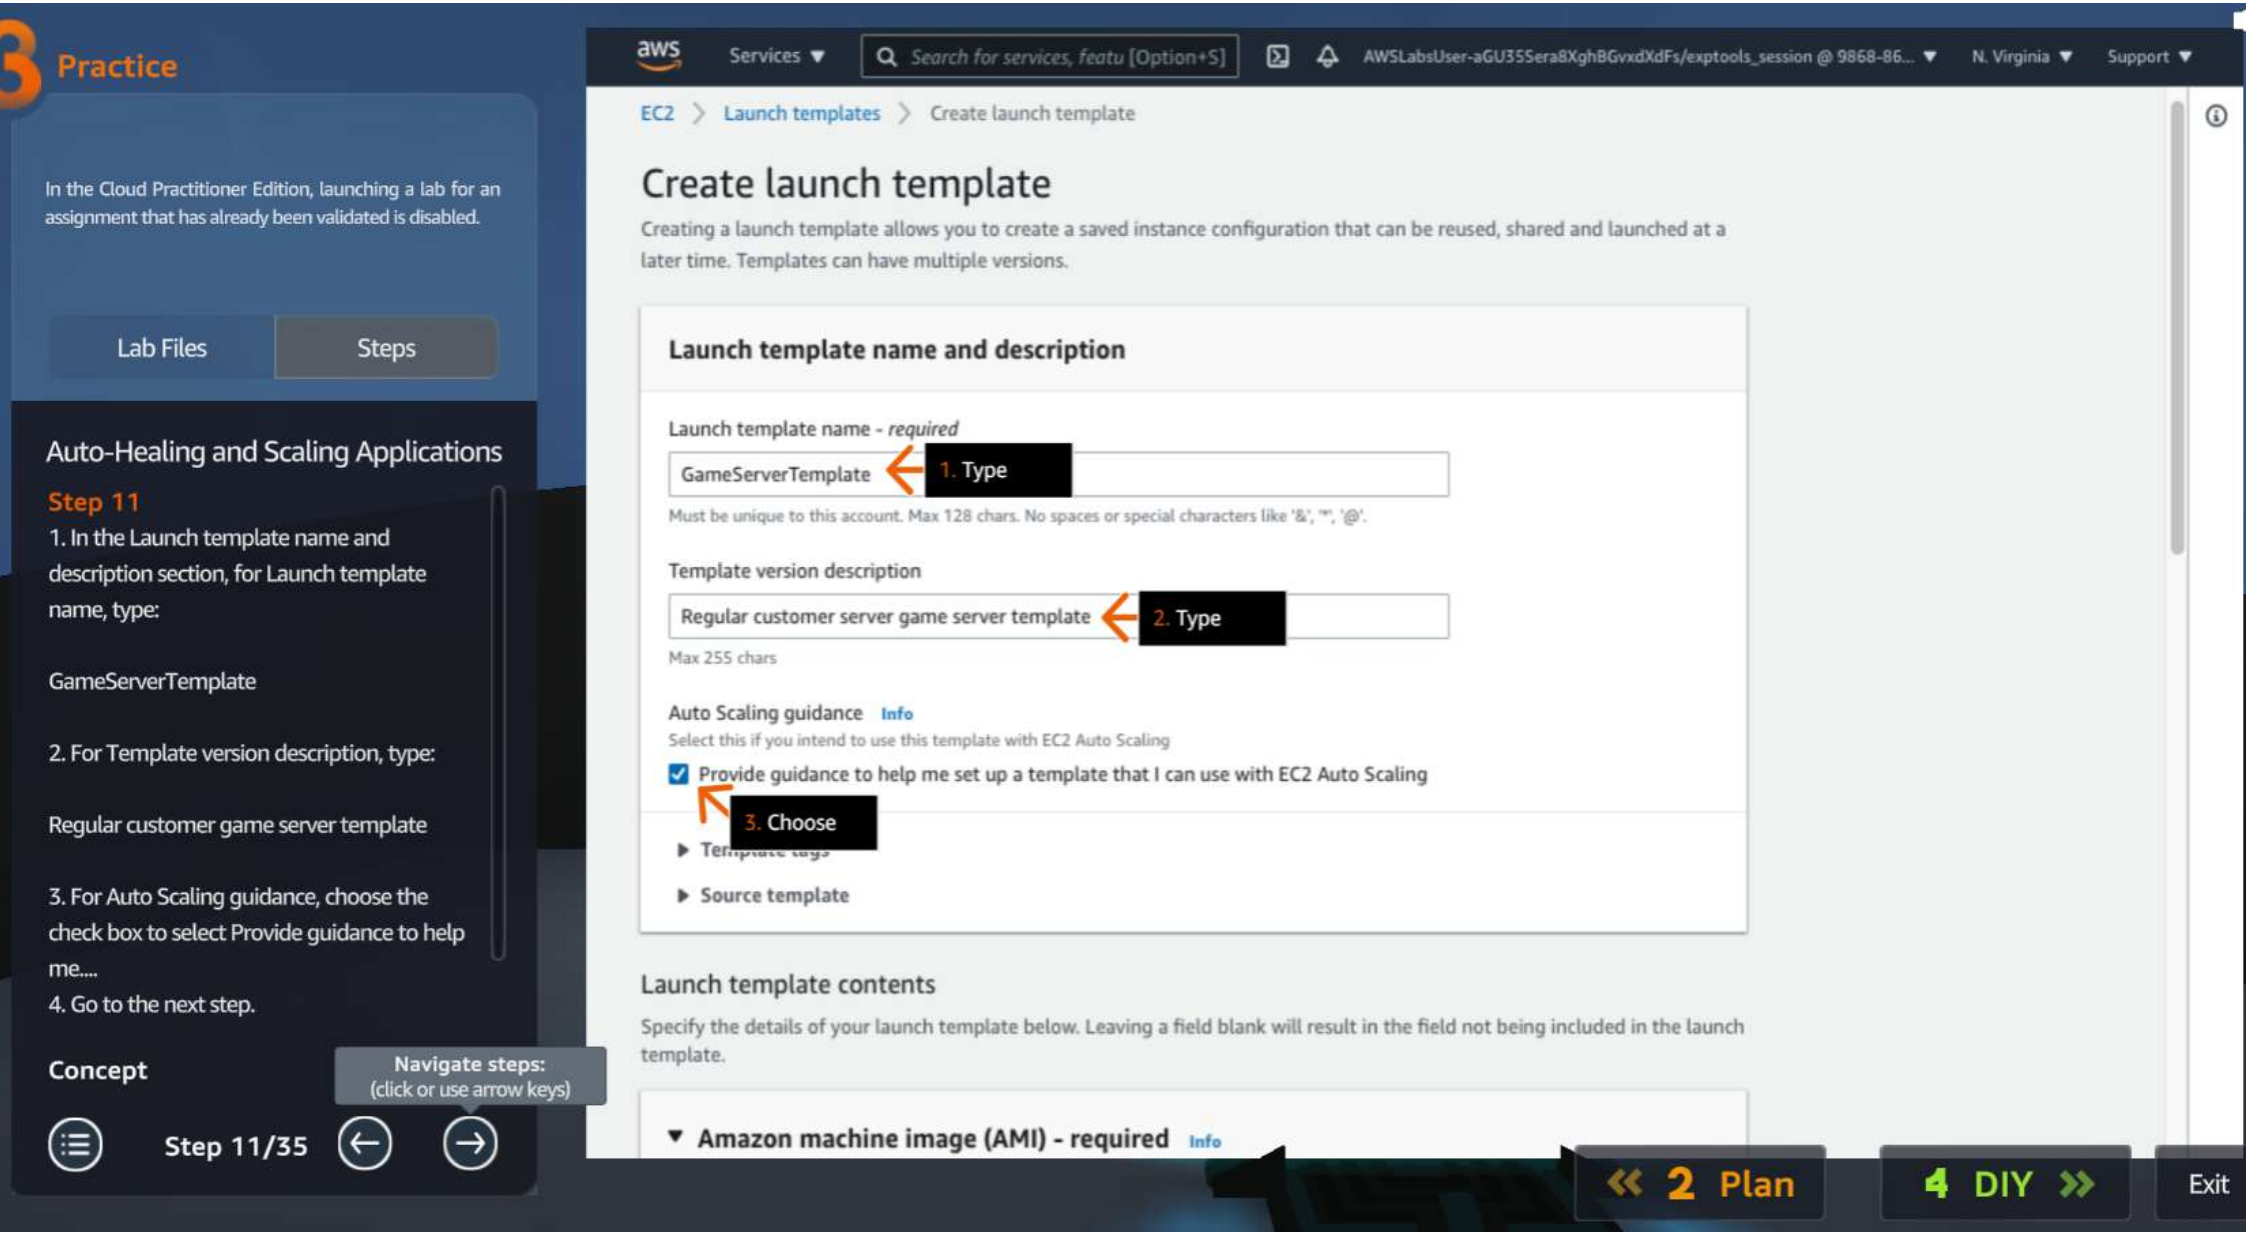
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 11
1. In the Launch template name and description section, for Launch template name, type:
GameServerTemplate
2. For Template version description, type:
Regular customer game server template
3. For Auto Scaling guidance, choose the check box to select Provide guidance to help me....
4. Go to the next step.
Concept
Navigate steps: (click or use arrow keys)
Step 11/35
aws
Services
Search for services, featu [Option+S]
AWSLabsUser-aGU35Sera8XghBGvxdXdFs/exptools_session @ 9868-86...
N. Virginia
Support
EC2 > Launch templates > Create launch template
Create launch template
Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions.
Launch template name and description
Launch template name - required
GameServerTemplate
1. Type
Must be unique to this account. Max 128 chars. No spaces or special characters like '&', '*', '@'.
Template version description
Regular customer server game server template
2. Type
Max 255 chars
Auto Scaling guidance Info
Select this if you intend to use this template with EC2 Auto Scaling
Provide guidance to help me set up a template that I can use with EC2 Auto Scaling
3. Choose
Source template
Launch template contents
Specify the details of your launch template below. Leaving a field blank will result in the field not being included in the launch template.
Amazon machine image (AMI) - required Info
« 2 Plan
4 DIY »
Exit

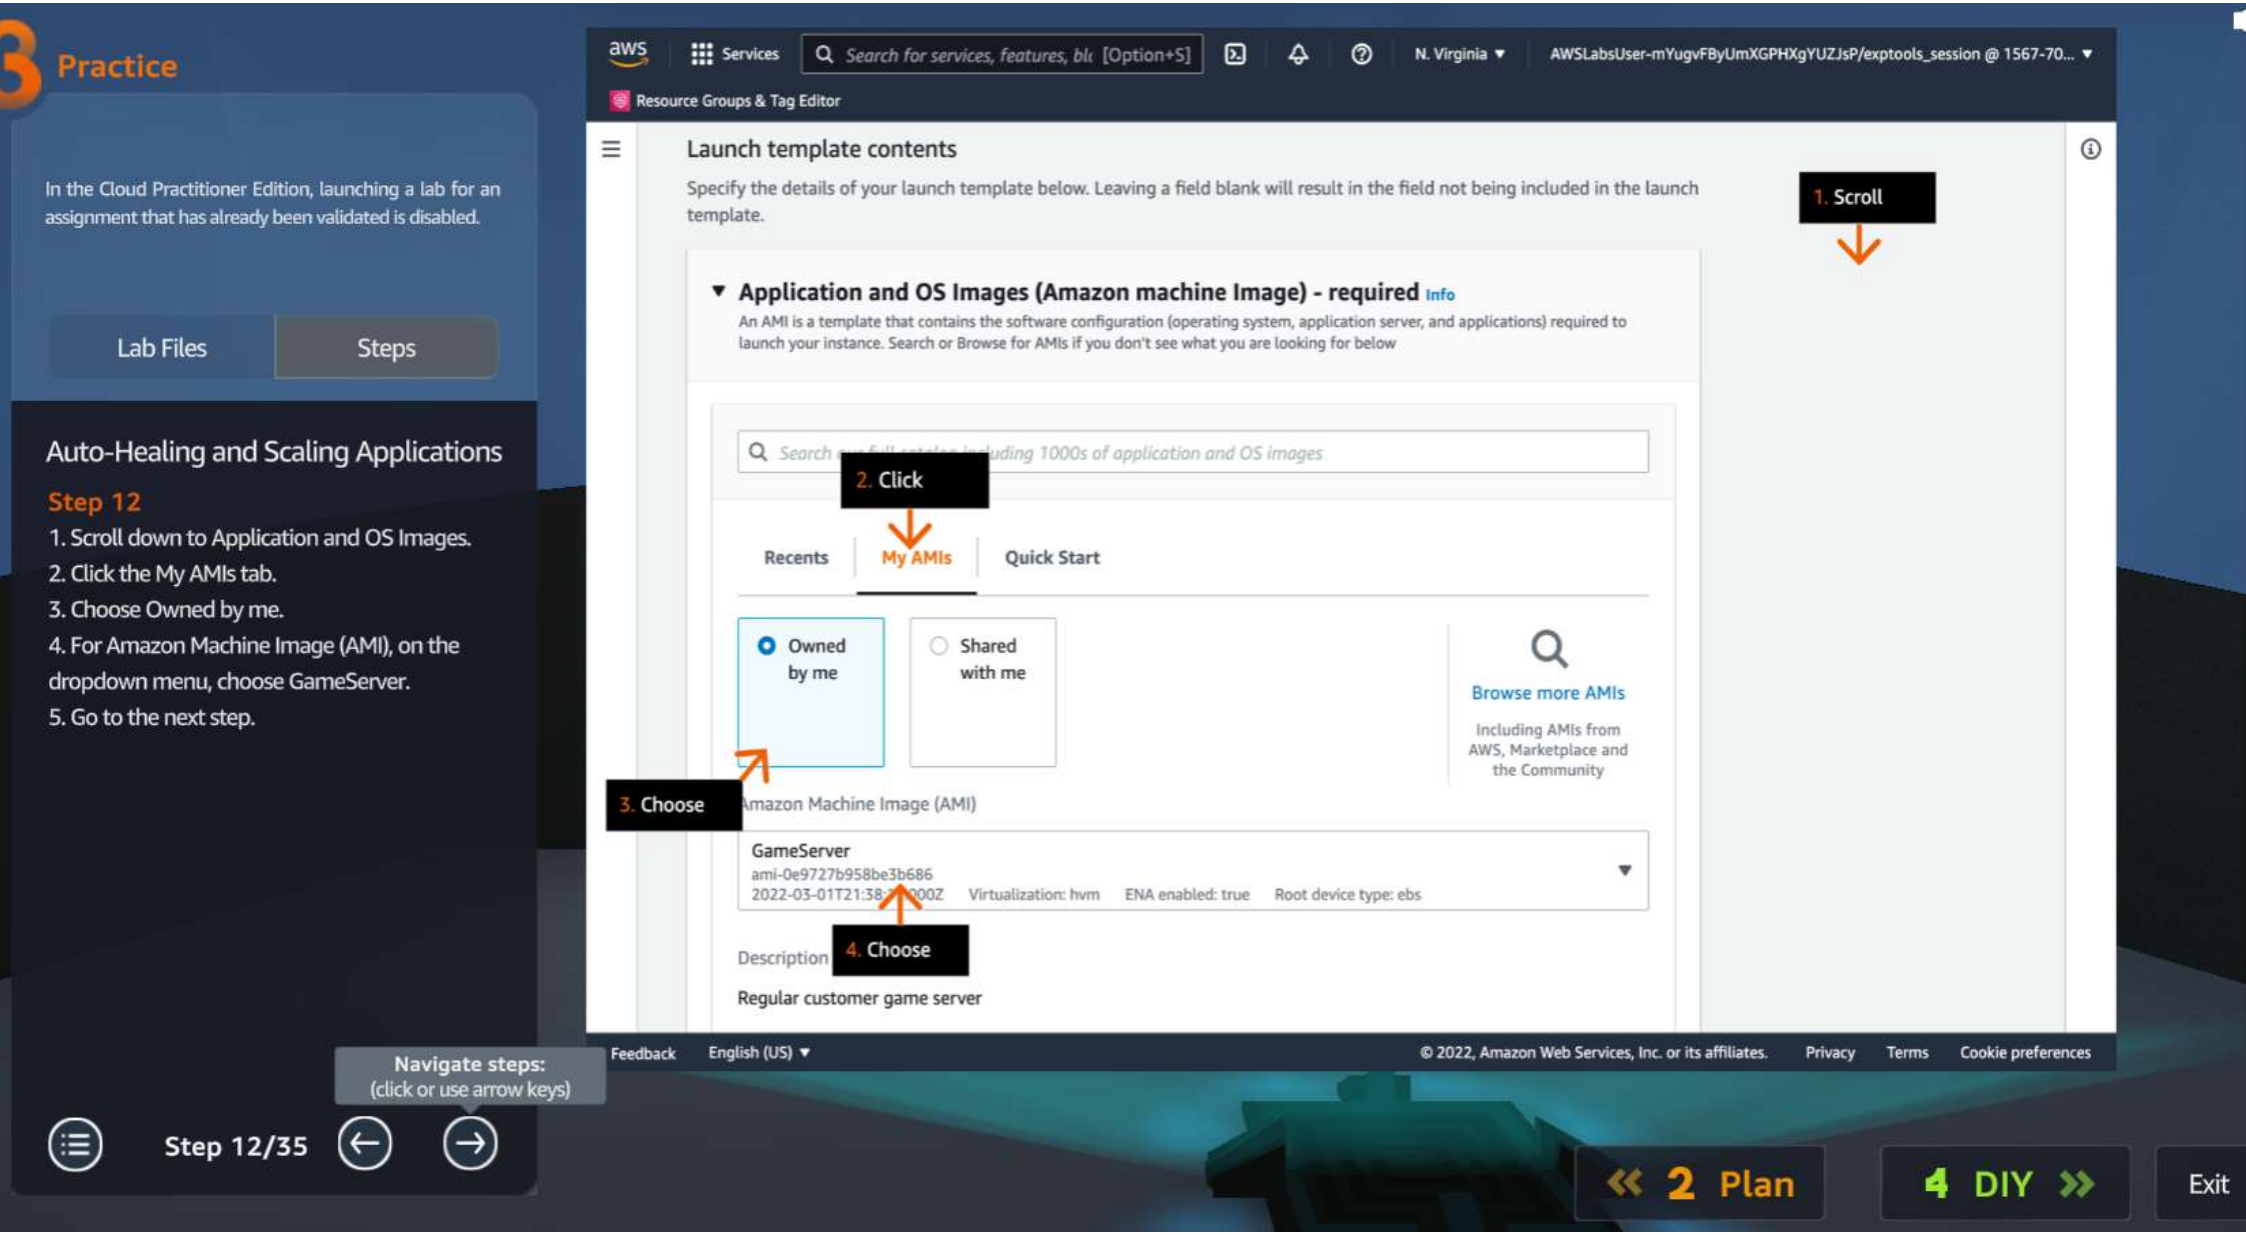
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 12
1. Scroll down to Application and OS Images.
2. Click the My AMIs tab.
3. Choose Owned by me.
4. For Amazon Machine Image (AMI), on the dropdown menu, choose GameServer.
5. Go to the next step.
Navigate steps:
(click or use arrow keys)
Step 12/35
aws
Services
Search for services, features, blc [Option+S]
N. Virginia
AWSLabsUser-mYugvFByUmXGPHXgYUZJsP/exptools_session @ 1567-70...
Resource Groups & Tag Editor
Launch template contents
Specify the details of your launch template below. Leaving a field blank will result in the field not being included in the launch template.
1. Scroll
Application and OS Images (Amazon machine Image) - required Info
An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below
2. Click
Recents
My AMIs
Quick Start
Owned by me
Shared with me
Browse more AMIs
Including AMIs from AWS, Marketplace and the Community
3. Choose
Amazon Machine Image (AMI)
GameServer
ami-0e9727b958be3b686
Virtualization: hvm
ENA enabled: true
Root device type: ebs
4. Choose
Description
Regular customer game server
Feedback
English (US)
© 2022, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
« 2 Plan
4 DIY »
Exit

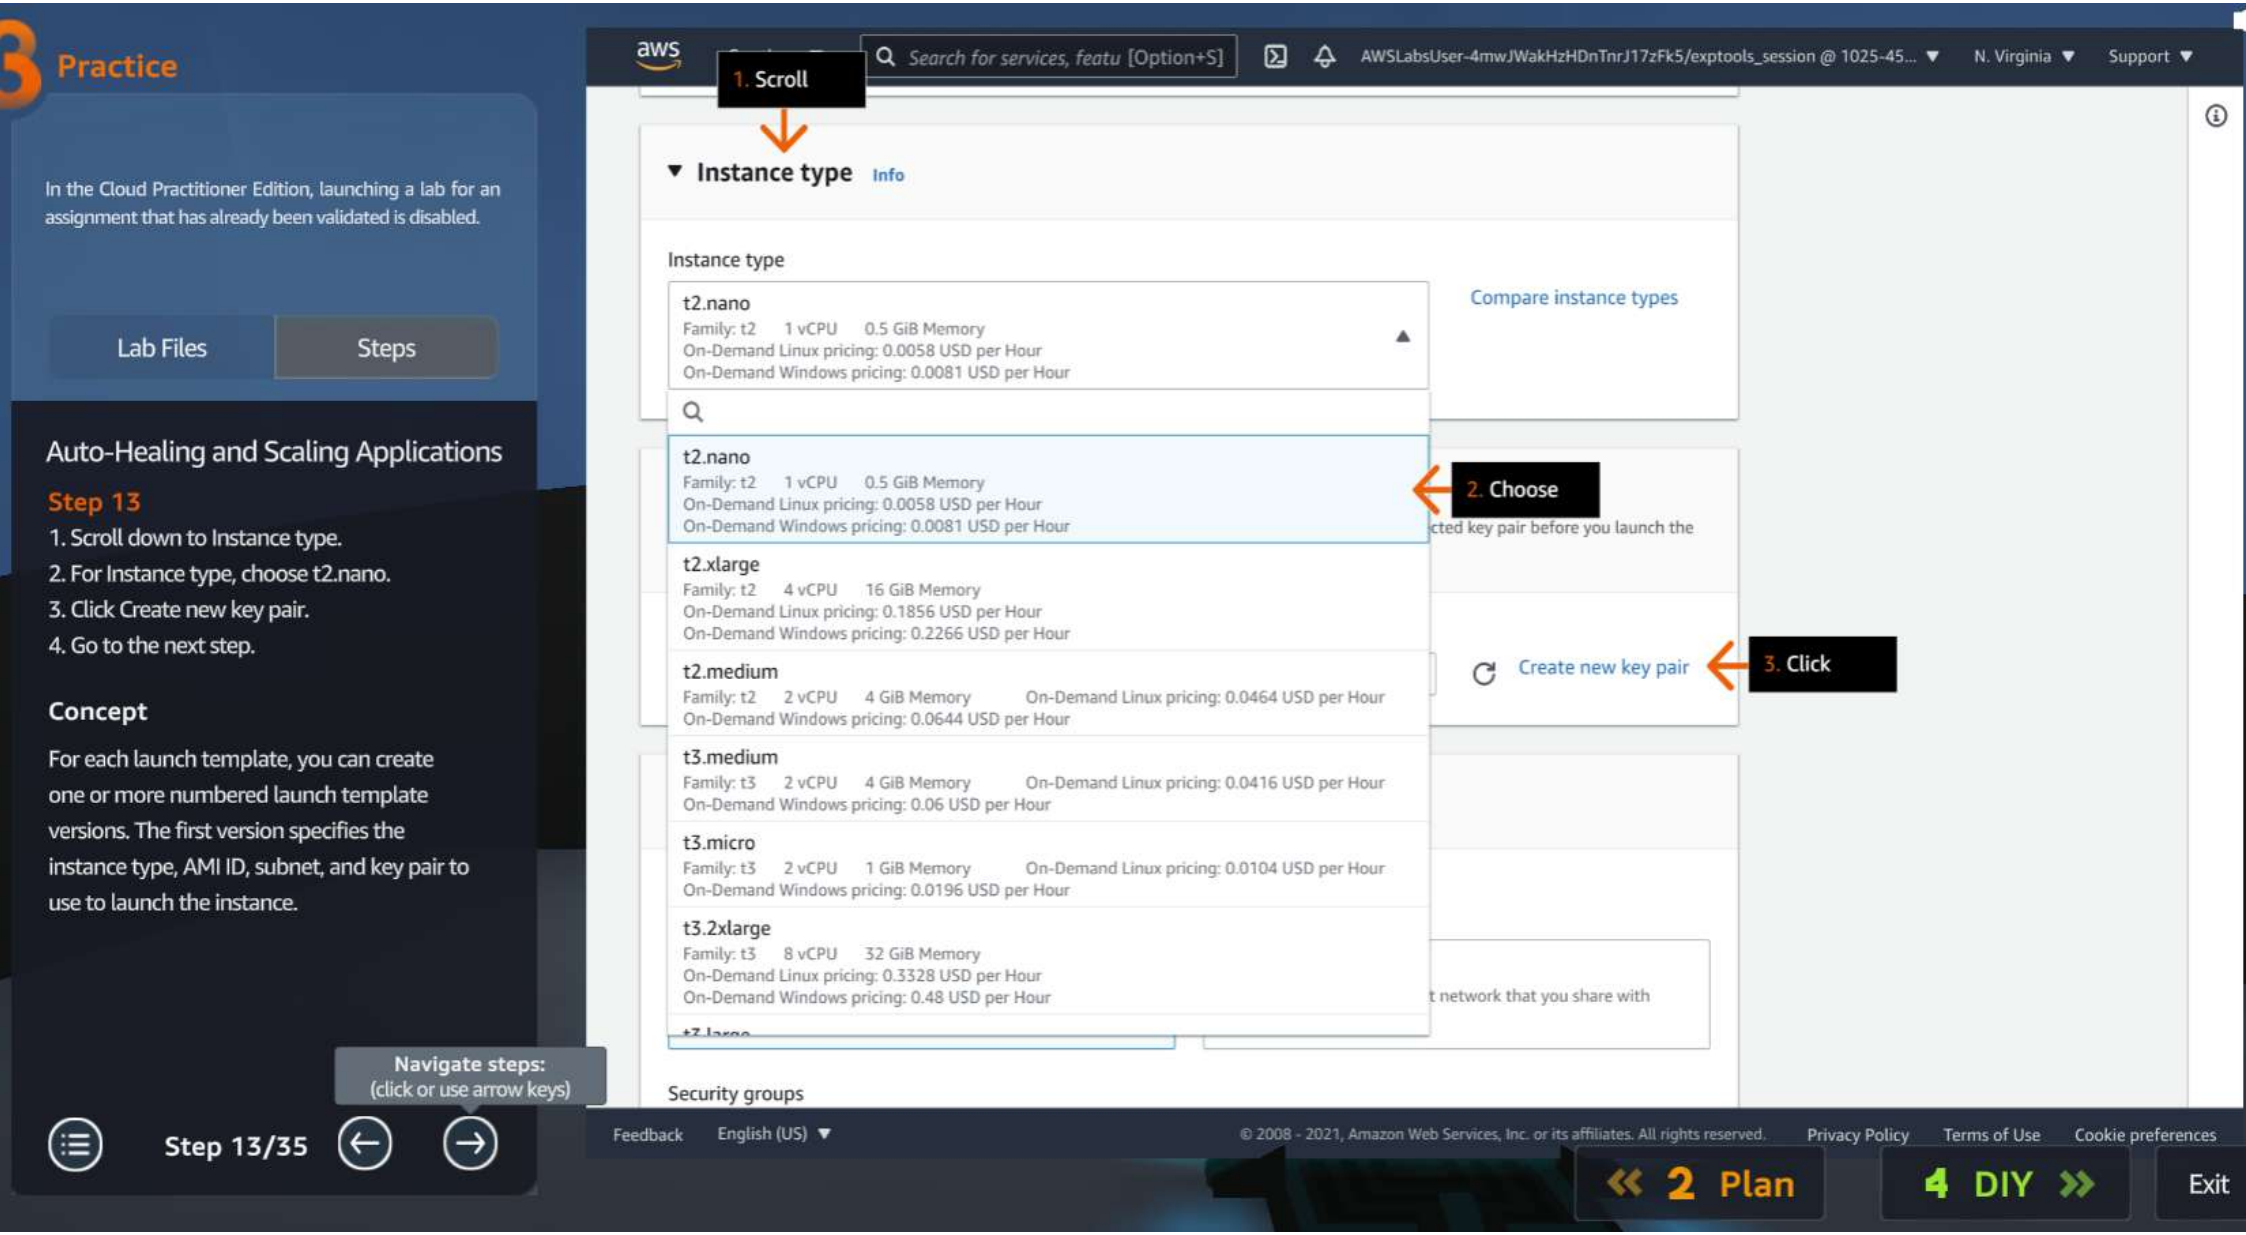
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 13
1. Scroll down to Instance type.
2. For Instance type, choose t2.nano.
3. Click Create new key pair.
4. Go to the next step.
Concept
For each launch template, you can create one or more numbered launch template versions. The first version specifies the instance type, AMI ID, subnet, and key pair to use to launch the instance.
Navigate steps: (click or use arrow keys)
Step 13/35
aws
Search for services, featu [Option+S]
AWSLabsUser-4mwJWakHzHDnTnrJ17zFk5/exptools_session @ 1025-45...
N. Virginia
Support
1. Scroll
Instance type Info
Instance type
t2.nano
Family: t2 1 vCPU 0.5 GiB Memory
On-Demand Linux pricing: 0.0058 USD per Hour
On-Demand Windows pricing: 0.0081 USD per Hour
Compare instance types
t2.nano
Family: t2 1 vCPU 0.5 GiB Memory
On-Demand Linux pricing: 0.0058 USD per Hour
On-Demand Windows pricing: 0.0081 USD per Hour
2. Choose
t2.xlarge
Family: t2 4 vCPU 16 GiB Memory
On-Demand Linux pricing: 0.1856 USD per Hour
On-Demand Windows pricing: 0.2266 USD per Hour
t2.medium
Family: t2 2 vCPU 4 GiB Memory On-Demand Linux pricing: 0.0464 USD per Hour
On-Demand Windows pricing: 0.0644 USD per Hour
t3.medium
Family: t3 2 vCPU 4 GiB Memory On-Demand Linux pricing: 0.0416 USD per Hour
On-Demand Windows pricing: 0.06 USD per Hour
t3.micro
Family: t3 2 vCPU 1 GiB Memory On-Demand Linux pricing: 0.0104 USD per Hour
On-Demand Windows pricing: 0.0196 USD per Hour
t3.2xlarge
Family: t3 8 vCPU 32 GiB Memory
On-Demand Linux pricing: 0.3328 USD per Hour
On-Demand Windows pricing: 0.48 USD per Hour
Create new key pair
3. Click
Security groups
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

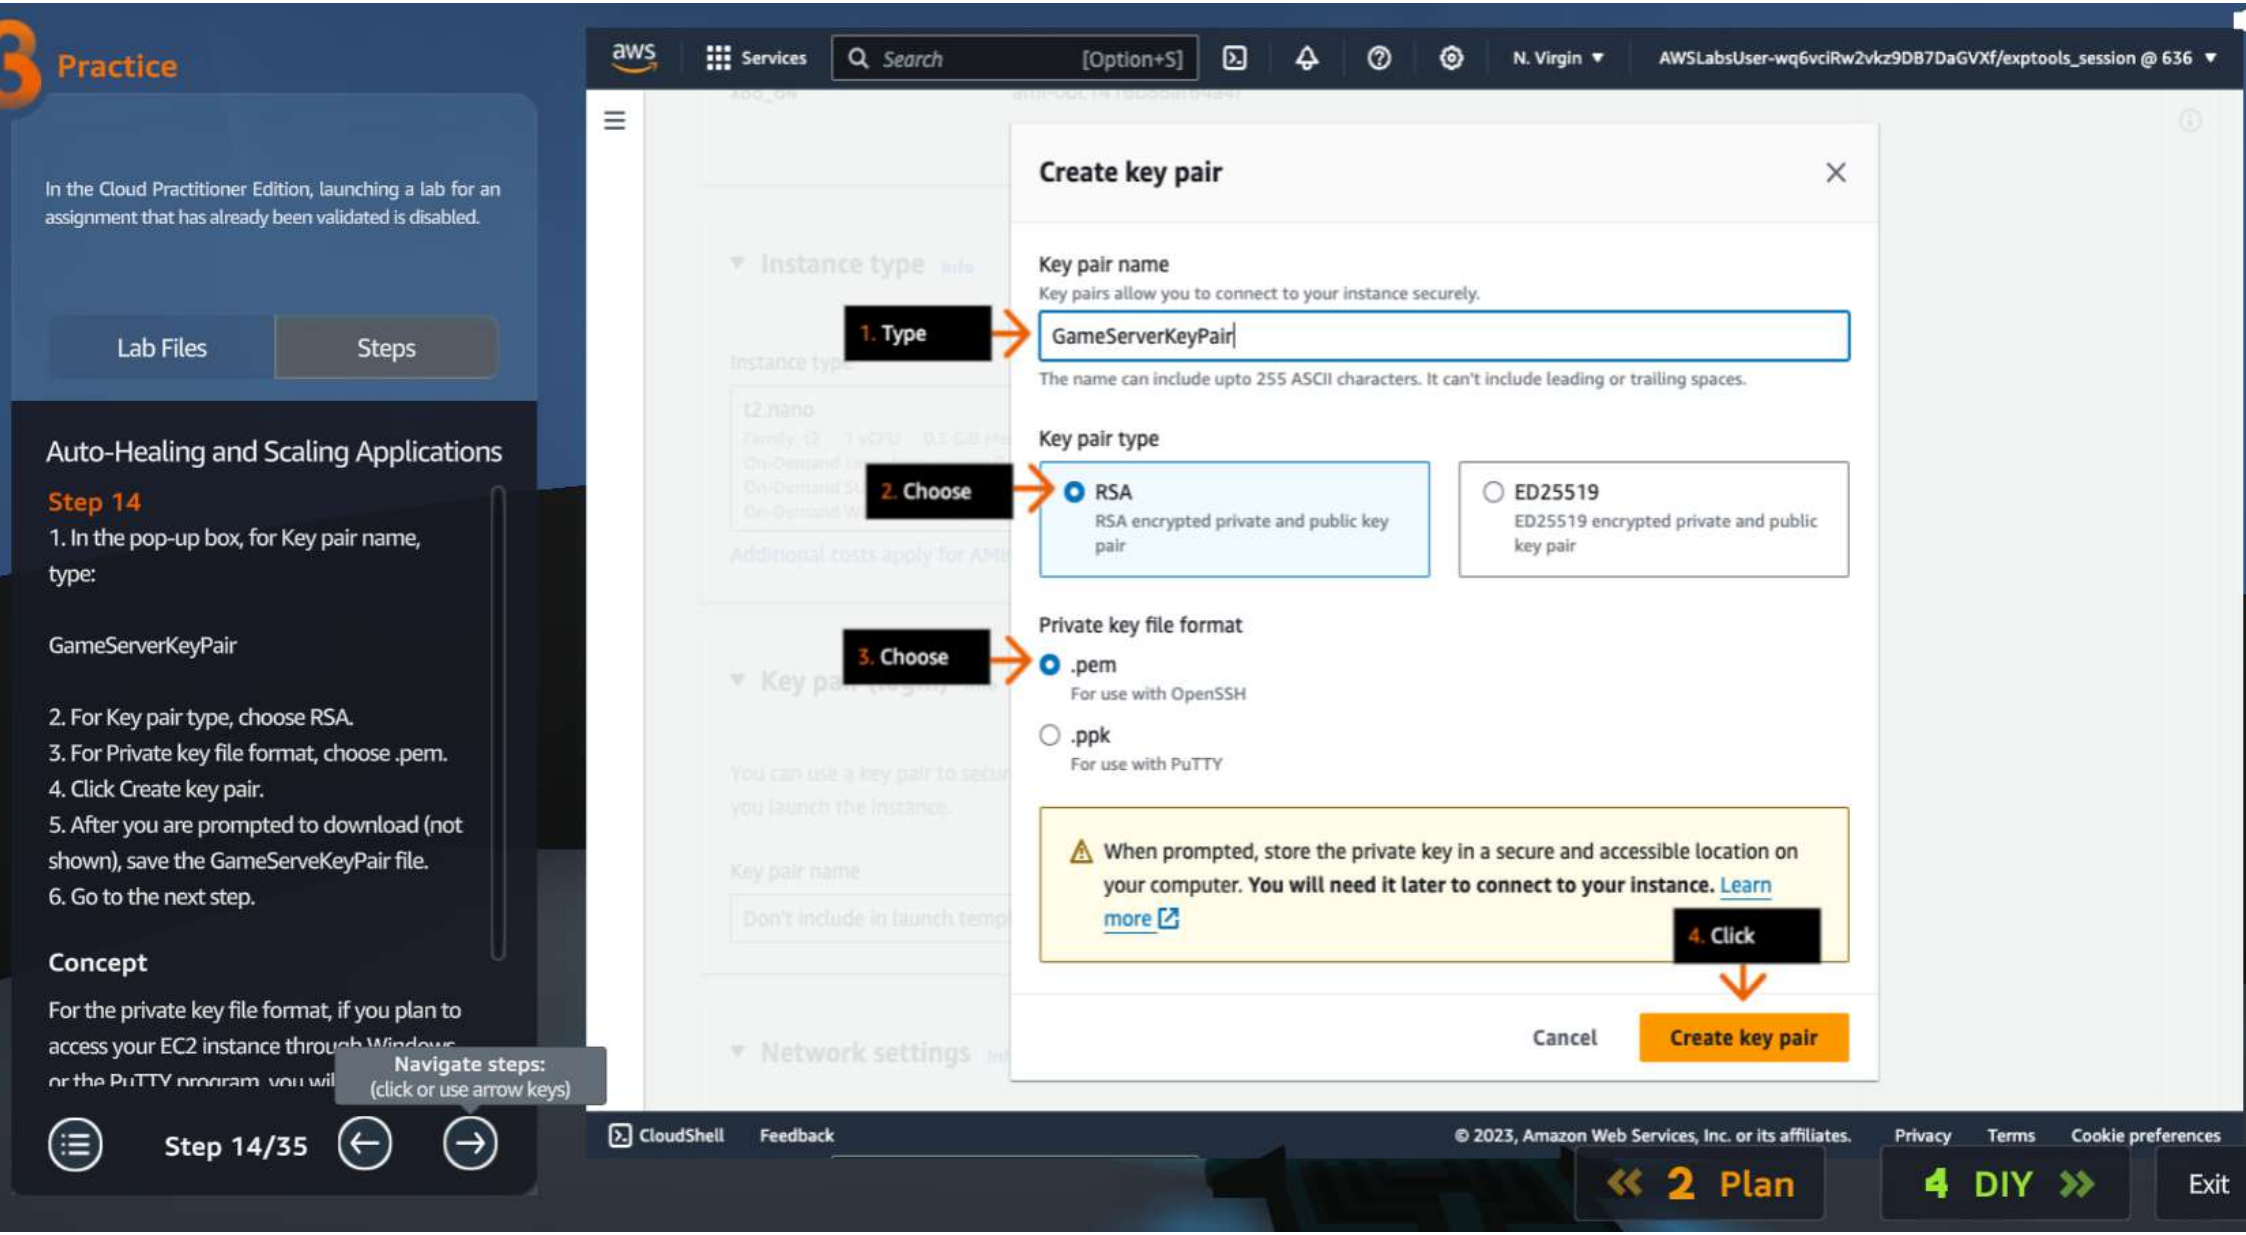
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 14
1. In the pop-up box, for Key pair name, type:
GameServerKeyPair
2. For Key pair type, choose RSA.
3. For Private key file format, choose .pem.
4. Click Create key pair.
5. After you are prompted to download (not shown), save the GameServeKeyPair file.
6. Go to the next step.
Concept
For the private key file format, if you plan to access your EC2 instance through Windows
or the PuTTY program, you wil
Navigate steps:
(click or use arrow keys)
Step 14/35
aws
Services
Search
[Option+S]
N. Virgin
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 636
Create key pair
Key pair name
Key pairs allow you to connect to your instance securely.
1. Type
GameServerKeyPair
The name can include upto 255 ASCII characters. It can't include leading or trailing spaces.
Key pair type
2. Choose
RSA
RSA encrypted private and public key pair
ED25519
ED25519 encrypted private and public key pair
Private key file format
3. Choose
.pem
For use with OpenSSH
.ppk
For use with PuTTY
When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. Learn more
4. Click
Cancel
Create key pair
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
« 2 Plan
4 DIY »
Exit

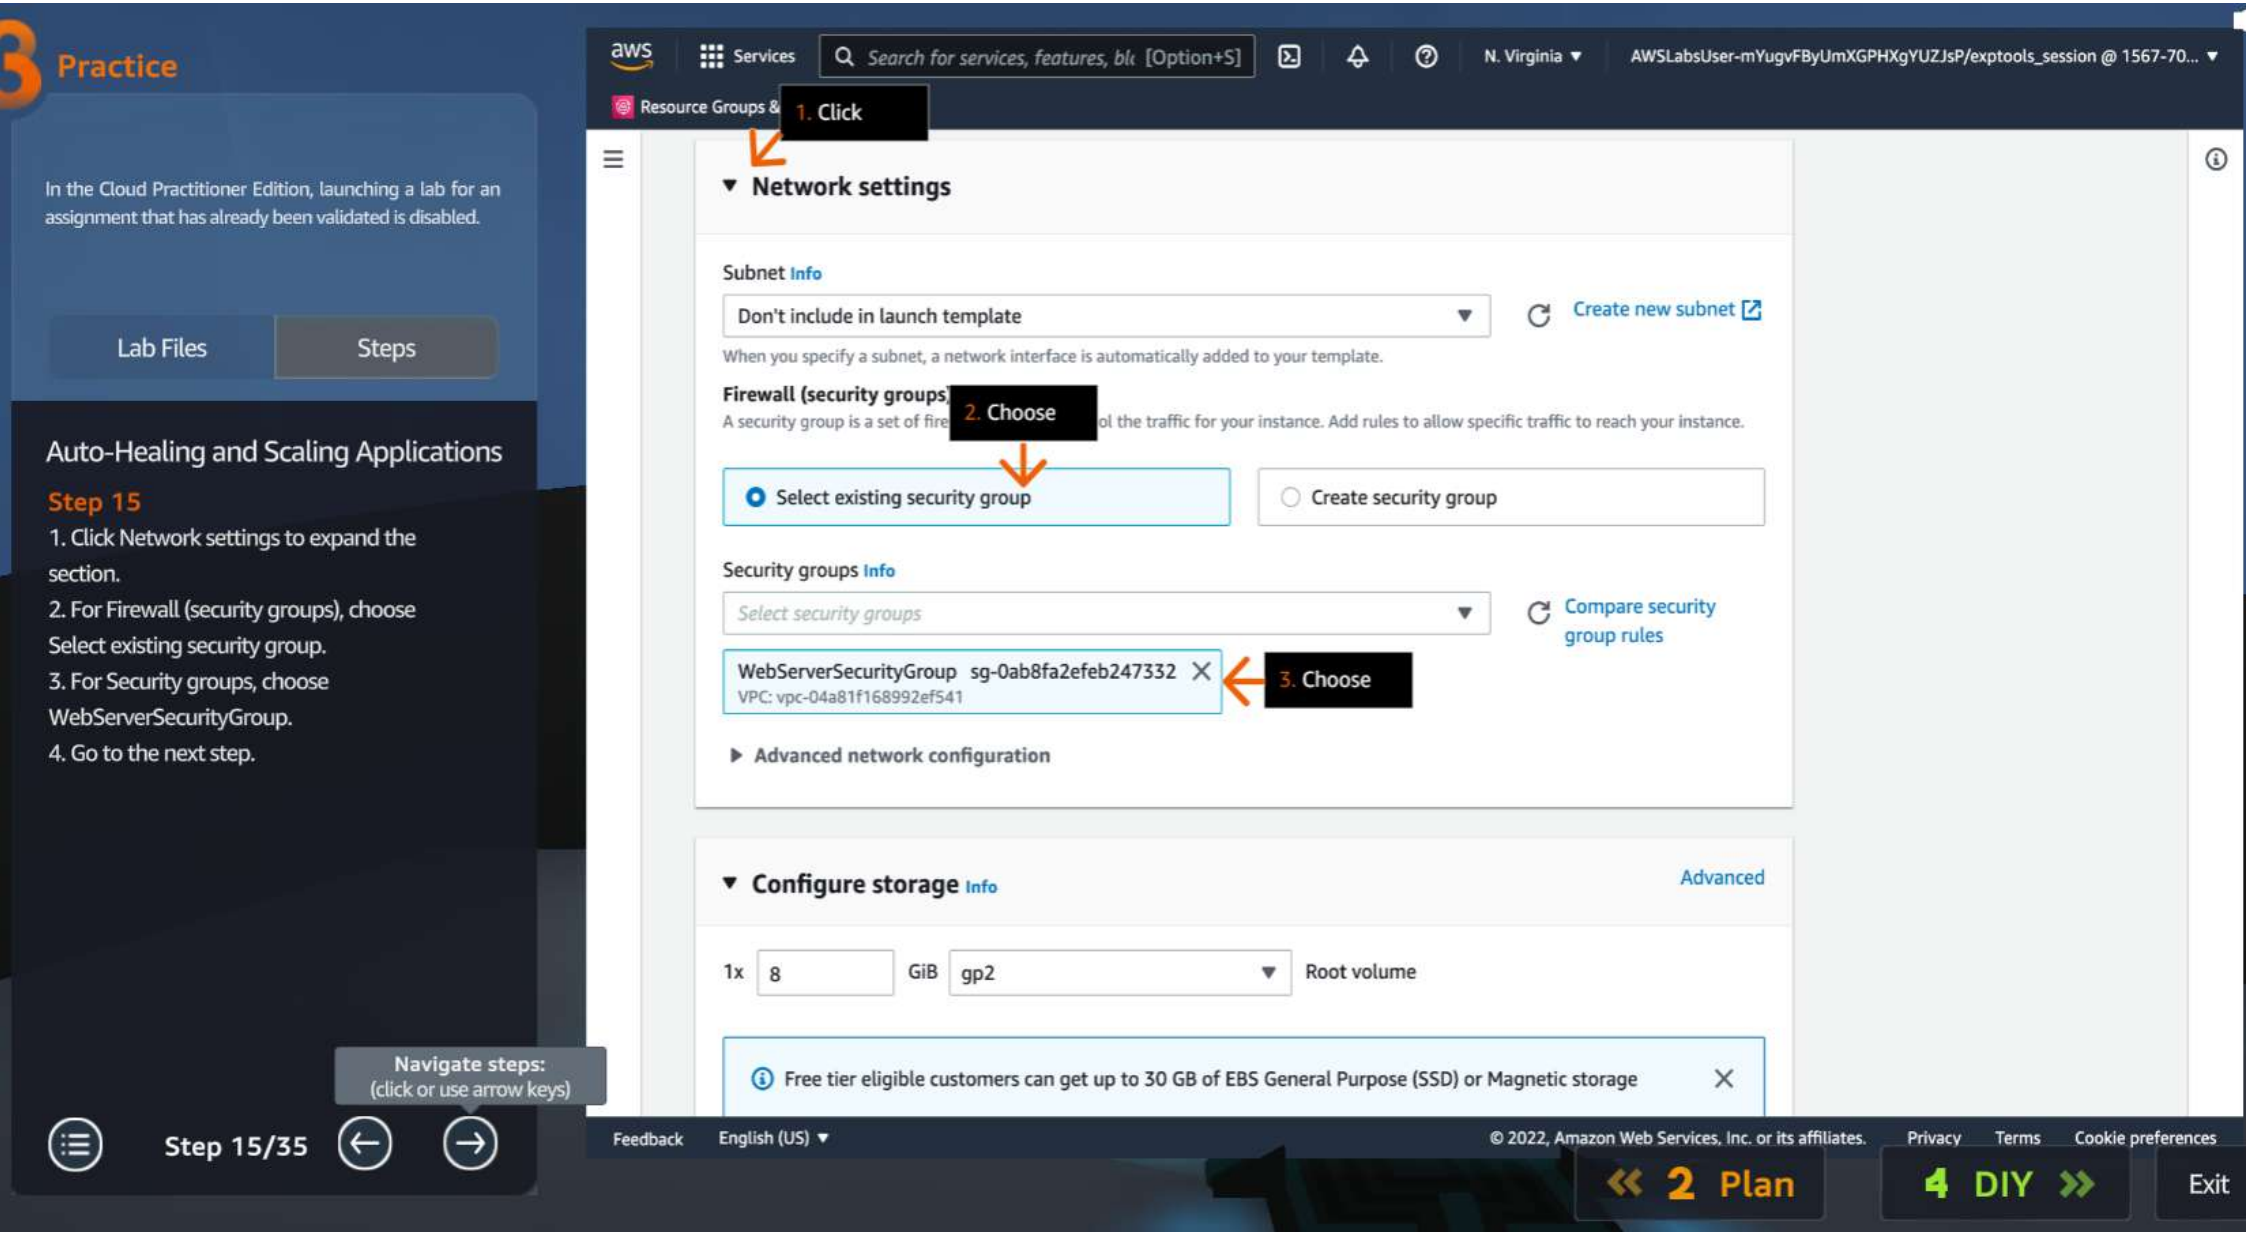
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 15
1. Click Network settings to expand the section.
2. For Firewall (security groups), choose Select existing security group.
3. For Security groups, choose WebServerSecurityGroup.
4. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 15/35
aws
Services
Search for services, features, bl [Option+S]
N. Virginia
AWSLabsUser-mYugvFByUmXGPHXgYUZJsP/exptools_session @ 1567-70...
Resource Groups &
1. Click
Network settings
Subnet Info
Don't include in launch template
Create new subnet
When you specify a subnet, a network interface is automatically added to your template.
Firewall (security groups)
A security group is a set of fire ol the traffic for your instance. Add rules to allow specific traffic to reach your instance.
2. Choose
Select existing security group
Create security group
Security groups Info
Select security groups
Compare security group rules
WebServerSecurityGroup sg-0ab8fa2efeb247332
VPC: vpc-04a81f168992ef541
3. Choose
Advanced network configuration
Configure storage Info
Advanced
1x
8
GiB
gp2
Root volume
Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage
Feedback
English (US)
© 2022, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

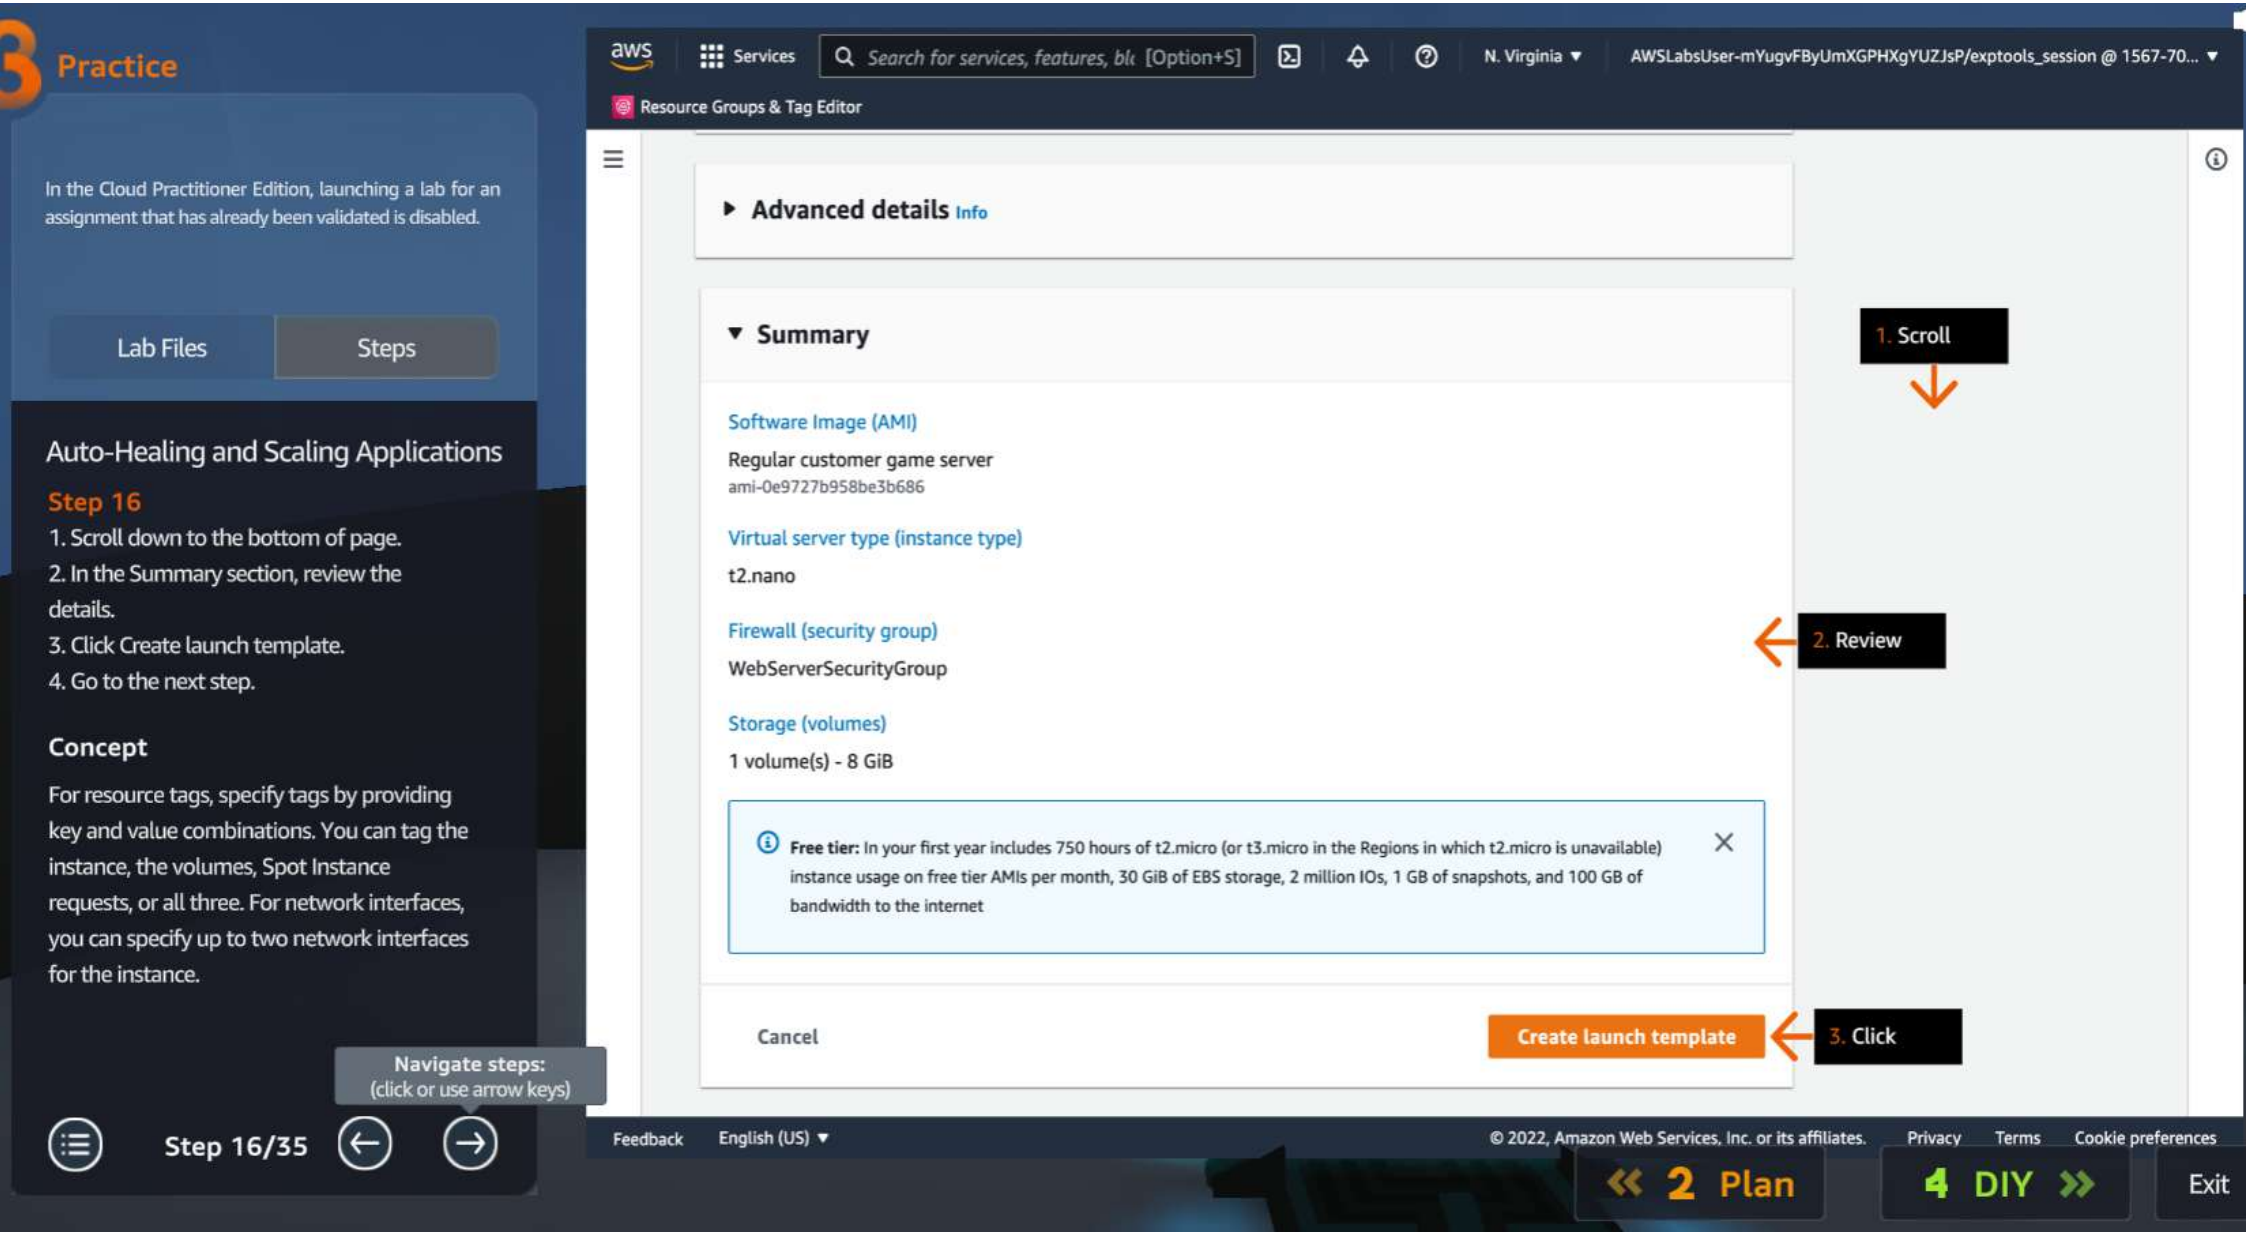
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 16
1. Scroll down to the bottom of page.
2. In the Summary section, review the details.
3. Click Create launch template.
4. Go to the next step.
Concept
For resource tags, specify tags by providing key and value combinations. You can tag the instance, the volumes, Spot Instance requests, or all three. For network interfaces, you can specify up to two network interfaces for the instance.
Navigate steps: (click or use arrow keys)
Step 16/35
aws
Services
Search for services, features, bl [Option+S]
N. Virginia
AWSLabsUser-mYugvFByUmXGPHXgYUZJsP/exptools_session @ 1567-70...
Resource Groups & Tag Editor
Advanced details Info
Summary
1. Scroll
Software Image (AMI)
Regular customer game server
ami-0e9727b958be3b686
Virtual server type (instance type)
t2.nano
Firewall (security group)
WebServerSecurityGroup
2. Review
Storage (volumes)
1 volume(s) - 8 GiB
Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million IOs, 1 GB of snapshots, and 100 GB of bandwidth to the internet
Cancel
Create launch template
3. Click
Feedback
English (US)
© 2022, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
« 2 Plan
4 DIY »
Exit

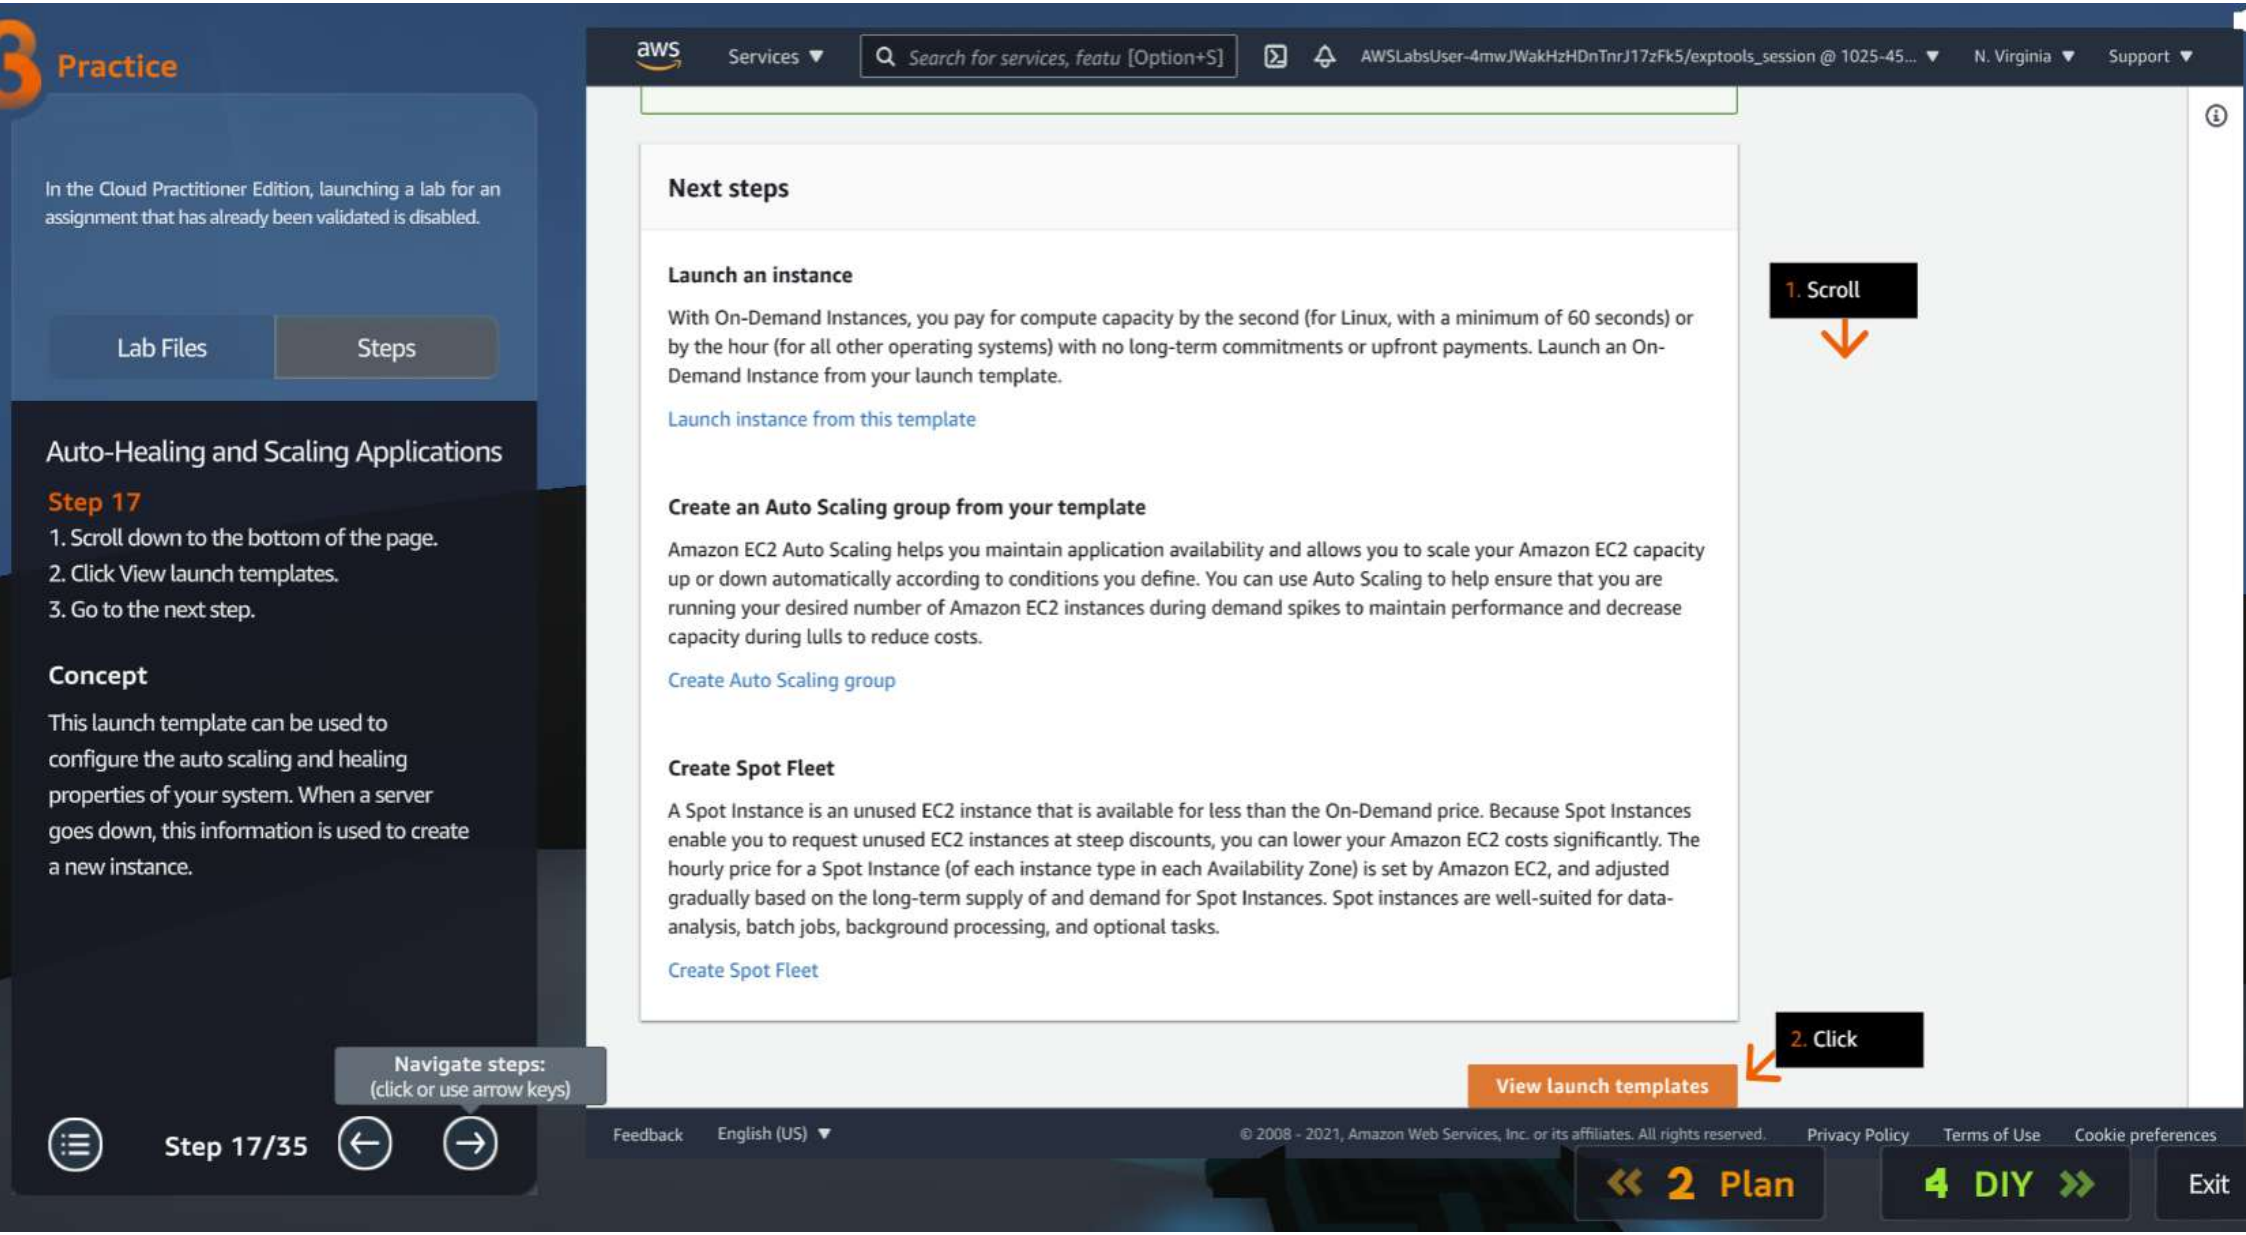
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 17
1. Scroll down to the bottom of the page.
2. Click View launch templates.
3. Go to the next step.
Concept
This launch template can be used to configure the auto scaling and healing properties of your system. When a server goes down, this information is used to create a new instance.
Navigate steps: (click or use arrow keys)
Step 17/35
aws
Services
Search for services, featu [Option+S]
AWSLabsUser-4mwJWakHzHDnTnrJ17zFk5/exptools_session @ 1025-45...
N. Virginia
Support
Next steps
Launch an instance
With On-Demand Instances, you pay for compute capacity by the second (for Linux, with a minimum of 60 seconds) or by the hour (for all other operating systems) with no long-term commitments or upfront payments. Launch an On-Demand Instance from your launch template.
Launch instance from this template
Create an Auto Scaling group from your template
Amazon EC2 Auto Scaling helps you maintain application availability and allows you to scale your Amazon EC2 capacity up or down automatically according to conditions you define. You can use Auto Scaling to help ensure that you are running your desired number of Amazon EC2 instances during demand spikes to maintain performance and decrease capacity during lulls to reduce costs.
Create Auto Scaling group
Create Spot Fleet
A Spot Instance is an unused EC2 instance that is available for less than the On-Demand price. Because Spot Instances enable you to request unused EC2 instances at steep discounts, you can lower your Amazon EC2 costs significantly. The hourly price for a Spot Instance (of each instance type in each Availability Zone) is set by Amazon EC2, and adjusted gradually based on the long-term supply of and demand for Spot Instances. Spot instances are well-suited for data-analysis, batch jobs, background processing, and optional tasks.
Create Spot Fleet
1. Scroll
2. Click
View launch templates
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

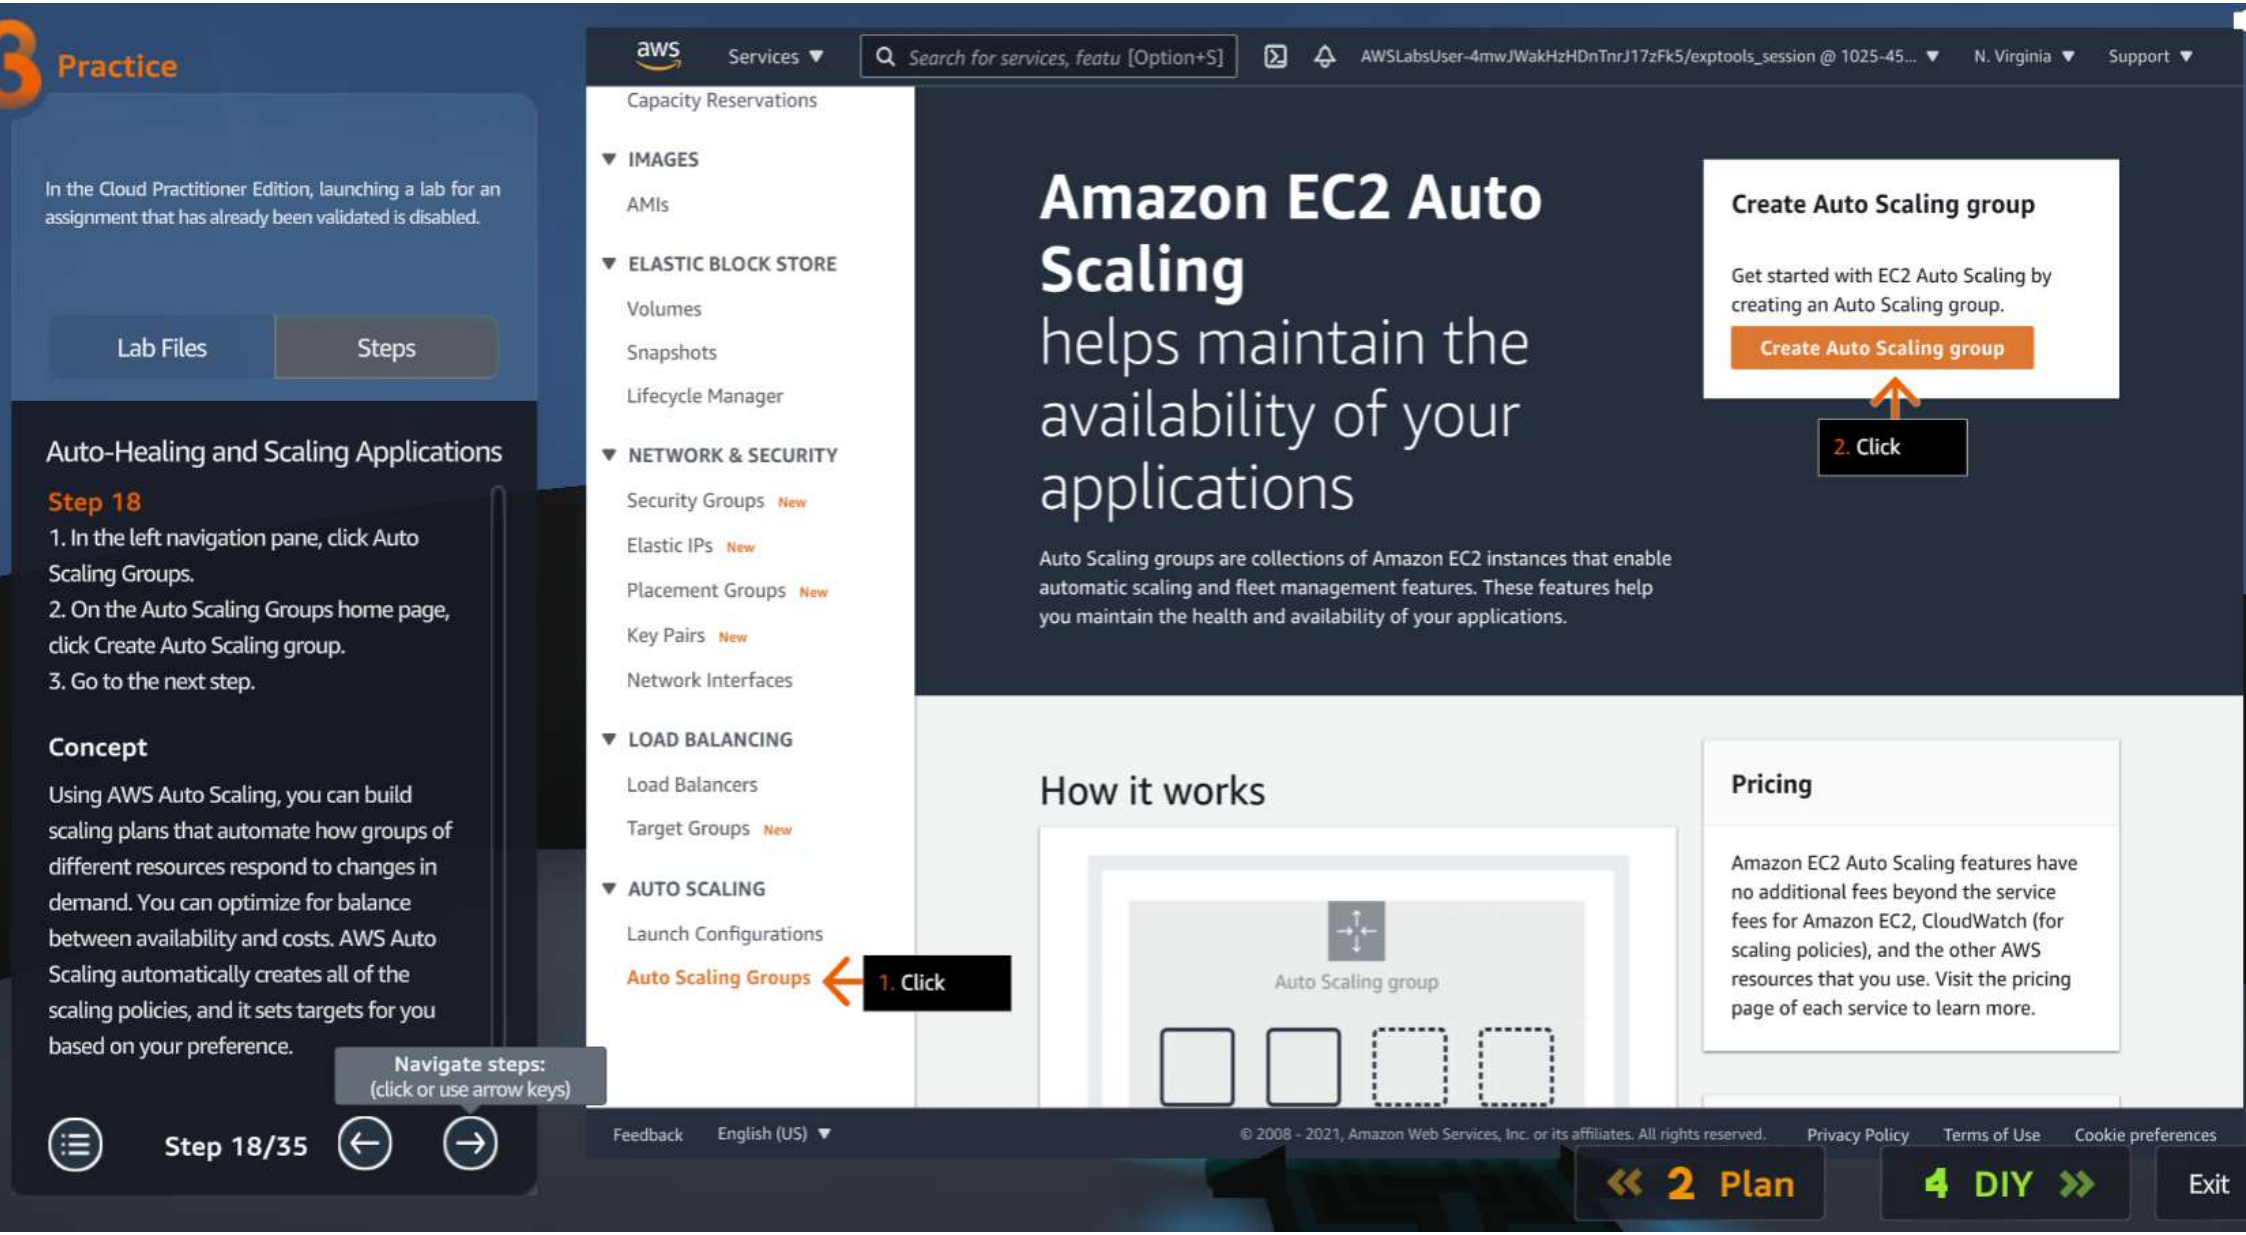
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 18
1. In the left navigation pane, click Auto Scaling Groups.
2. On the Auto Scaling Groups home page, click Create Auto Scaling group.
3. Go to the next step.
Concept
Using AWS Auto Scaling, you can build scaling plans that automate how groups of different resources respond to changes in demand. You can optimize for balance between availability and costs. AWS Auto Scaling automatically creates all of the scaling policies, and it sets targets for you based on your preference.
Navigate steps: (click or use arrow keys)
Step 18/35
aws
Services
Search for services, featu [Option+S]
AWSLabsUser-4mwJWakHzHDnTnrJ17zFk5/exptools_session @ 1025-45...
N. Virginia
Support
Capacity Reservations
IMAGES
AMIs
ELASTIC BLOCK STORE
Volumes
Snapshots
Lifecycle Manager
NETWORK & SECURITY
Security Groups New
Elastic IPs New
Placement Groups New
Key Pairs New
Network Interfaces
LOAD BALANCING
Load Balancers
Target Groups New
AUTO SCALING
Launch Configurations
Auto Scaling Groups
1. Click
Amazon EC2 Auto Scaling
helps maintain the availability of your applications
Auto Scaling groups are collections of Amazon EC2 instances that enable automatic scaling and fleet management features. These features help you maintain the health and availability of your applications.
Create Auto Scaling group
Get started with EC2 Auto Scaling by creating an Auto Scaling group.
Create Auto Scaling group
2. Click
How it works
Auto Scaling group
Pricing
Amazon EC2 Auto Scaling features have no additional fees beyond the service fees for Amazon EC2, CloudWatch (for scaling policies), and the other AWS resources that you use. Visit the pricing page of each service to learn more.
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

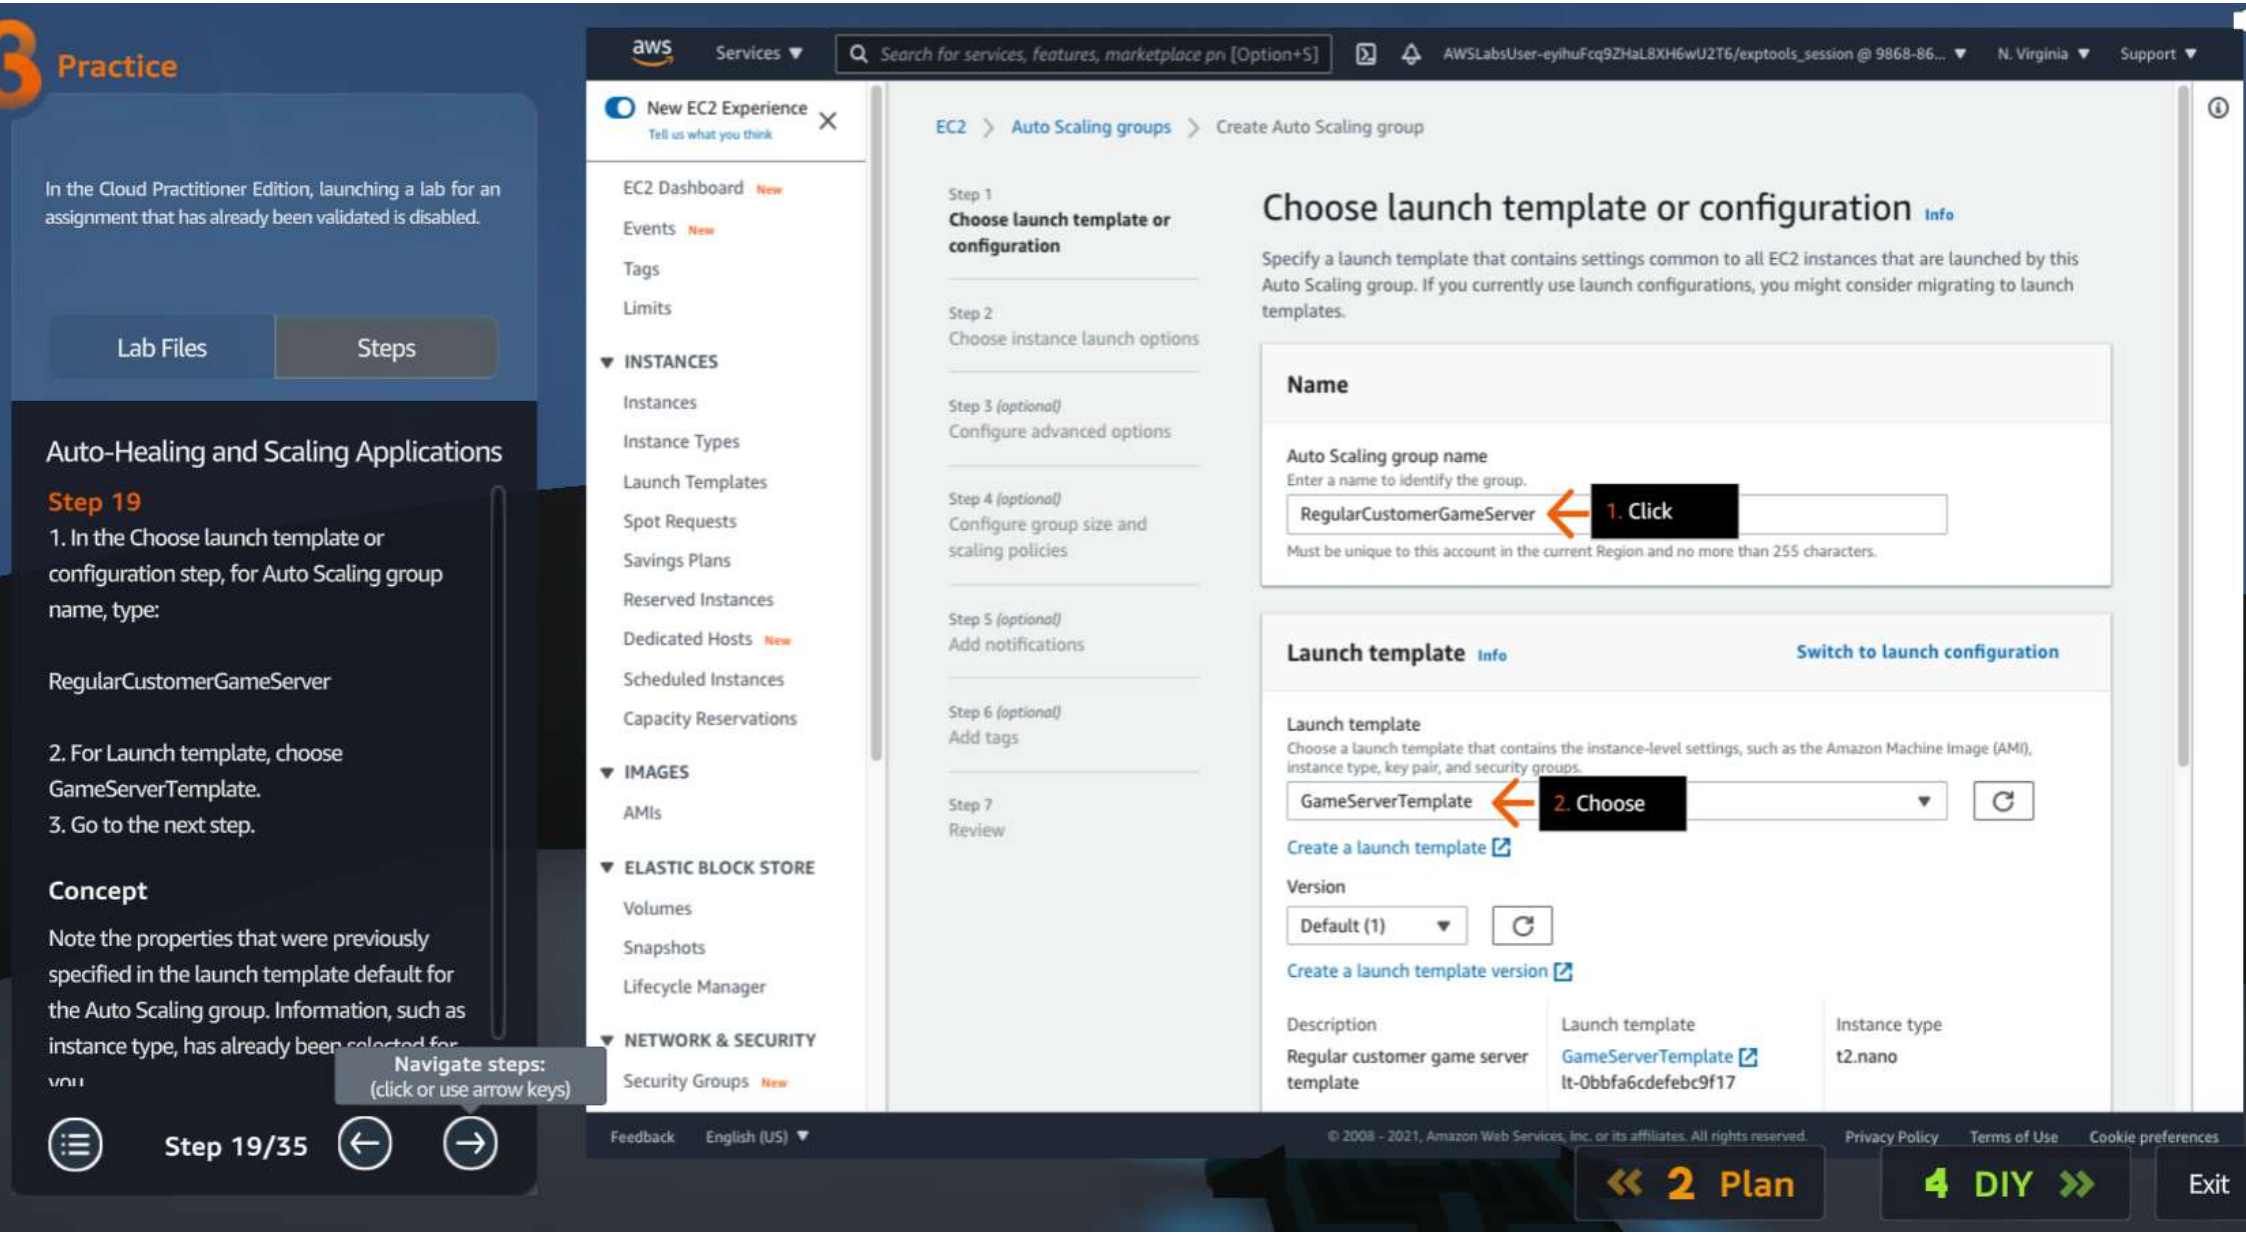
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 19
1. In the Choose launch template or configuration step, for Auto Scaling group name, type:
RegularCustomerGameServer
2. For Launch template, choose GameServerTemplate.
3. Go to the next step.
Concept
Note the properties that were previously specified in the launch template default for the Auto Scaling group. Information, such as instance type, has already been selected for you
Navigate steps: (click or use arrow keys)
Step 19/35
aws
Services
Search for services, features, marketplace pn [Option+S]
AWSLabsUser-eyihuFcq9ZHaL8XH6wU2T6/exptools_session @ 9868-86...
N. Virginia
Support
New EC2 Experience
Tell us what you think
EC2 Dashboard New
Events New
Tags
Limits
INSTANCES
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts New
Scheduled Instances
Capacity Reservations
IMAGES
AMIs
ELASTIC BLOCK STORE
Volumes
Snapshots
Lifecycle Manager
NETWORK & SECURITY
Security Groups New
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1 Choose launch template or configuration
Step 2 Choose instance launch options
Step 3 (optional) Configure advanced options
Step 4 (optional) Configure group size and scaling policies
Step 5 (optional) Add notifications
Step 6 (optional) Add tags
Step 7 Review
Choose launch template or configuration Info
Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group. If you currently use launch configurations, you might consider migrating to launch templates.
Name
Auto Scaling group name
Enter a name to identify the group.
RegularCustomerGameServer
1. Click
Must be unique to this account in the current Region and no more than 255 characters.
Launch template Info
Switch to launch configuration
Launch template
Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.
GameServerTemplate
2. Choose
Create a launch template
Version
Default (1)
Create a launch template version
Description
Regular customer game server template
Launch template
GameServerTemplate
lt-0bbfa6cdefebc9f17
Instance type
t2.nano
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

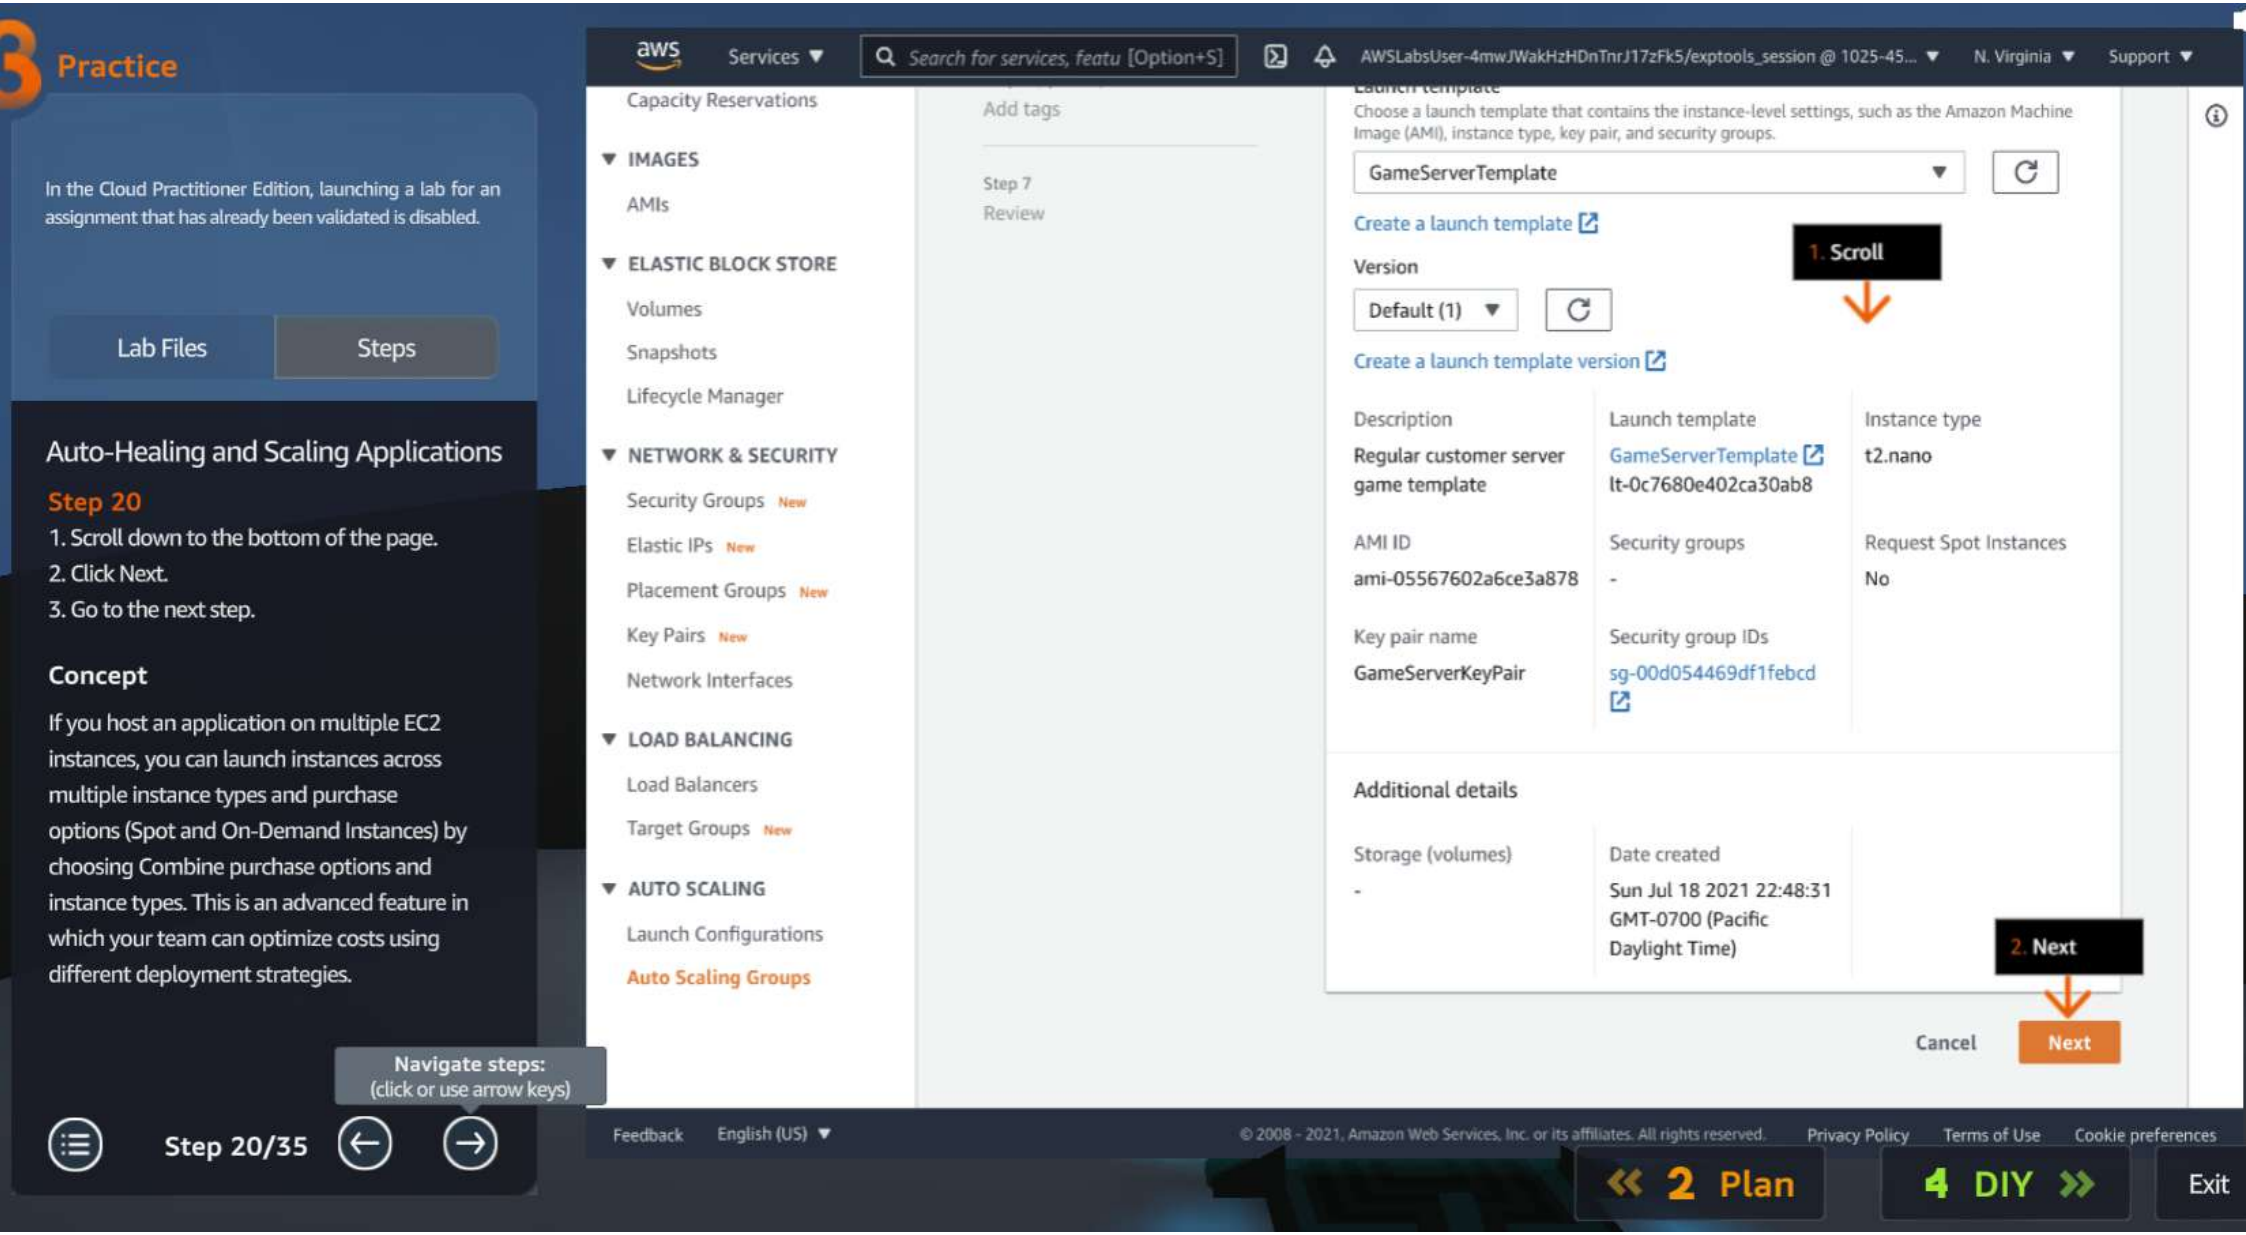
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 20
1. Scroll down to the bottom of the page.
2. Click Next.
3. Go to the next step.
Concept
If you host an application on multiple EC2 instances, you can launch instances across multiple instance types and purchase options (Spot and On-Demand Instances) by choosing Combine purchase options and instance types. This is an advanced feature in which your team can optimize costs using different deployment strategies.
Navigate steps: (click or use arrow keys)
Step 20/35
aws
Services
Search for services, featu [Option+S]
AWSLabsUser-4mwJWakHzHDnTnrJ17zFk5/exptools_session @ 1025-45...
N. Virginia
Support
Capacity Reservations
IMAGES
AMIs
ELASTIC BLOCK STORE
Volumes
Snapshots
Lifecycle Manager
NETWORK & SECURITY
Security Groups New
Elastic IPs New
Placement Groups New
Key Pairs New
Network Interfaces
LOAD BALANCING
Load Balancers
Target Groups New
AUTO SCALING
Launch Configurations
Auto Scaling Groups
Add tags
Step 7
Review
Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.
GameServerTemplate
Create a launch template
1. Scroll
Version
Default (1)
Create a launch template version
Description
Regular customer server game template
Launch template
GameServerTemplate
lt-0c7680e402ca30ab8
Instance type
t2.nano
AMI ID
ami-05567602a6ce3a878
Security groups
-
Request Spot Instances
No
Key pair name
GameServerKeyPair
Security group IDs
sg-00d054469df1febcd
Additional details
Storage (volumes)
-
Date created
Sun Jul 18 2021 22:48:31 GMT-0700 (Pacific Daylight Time)
2. Next
Cancel
Next
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

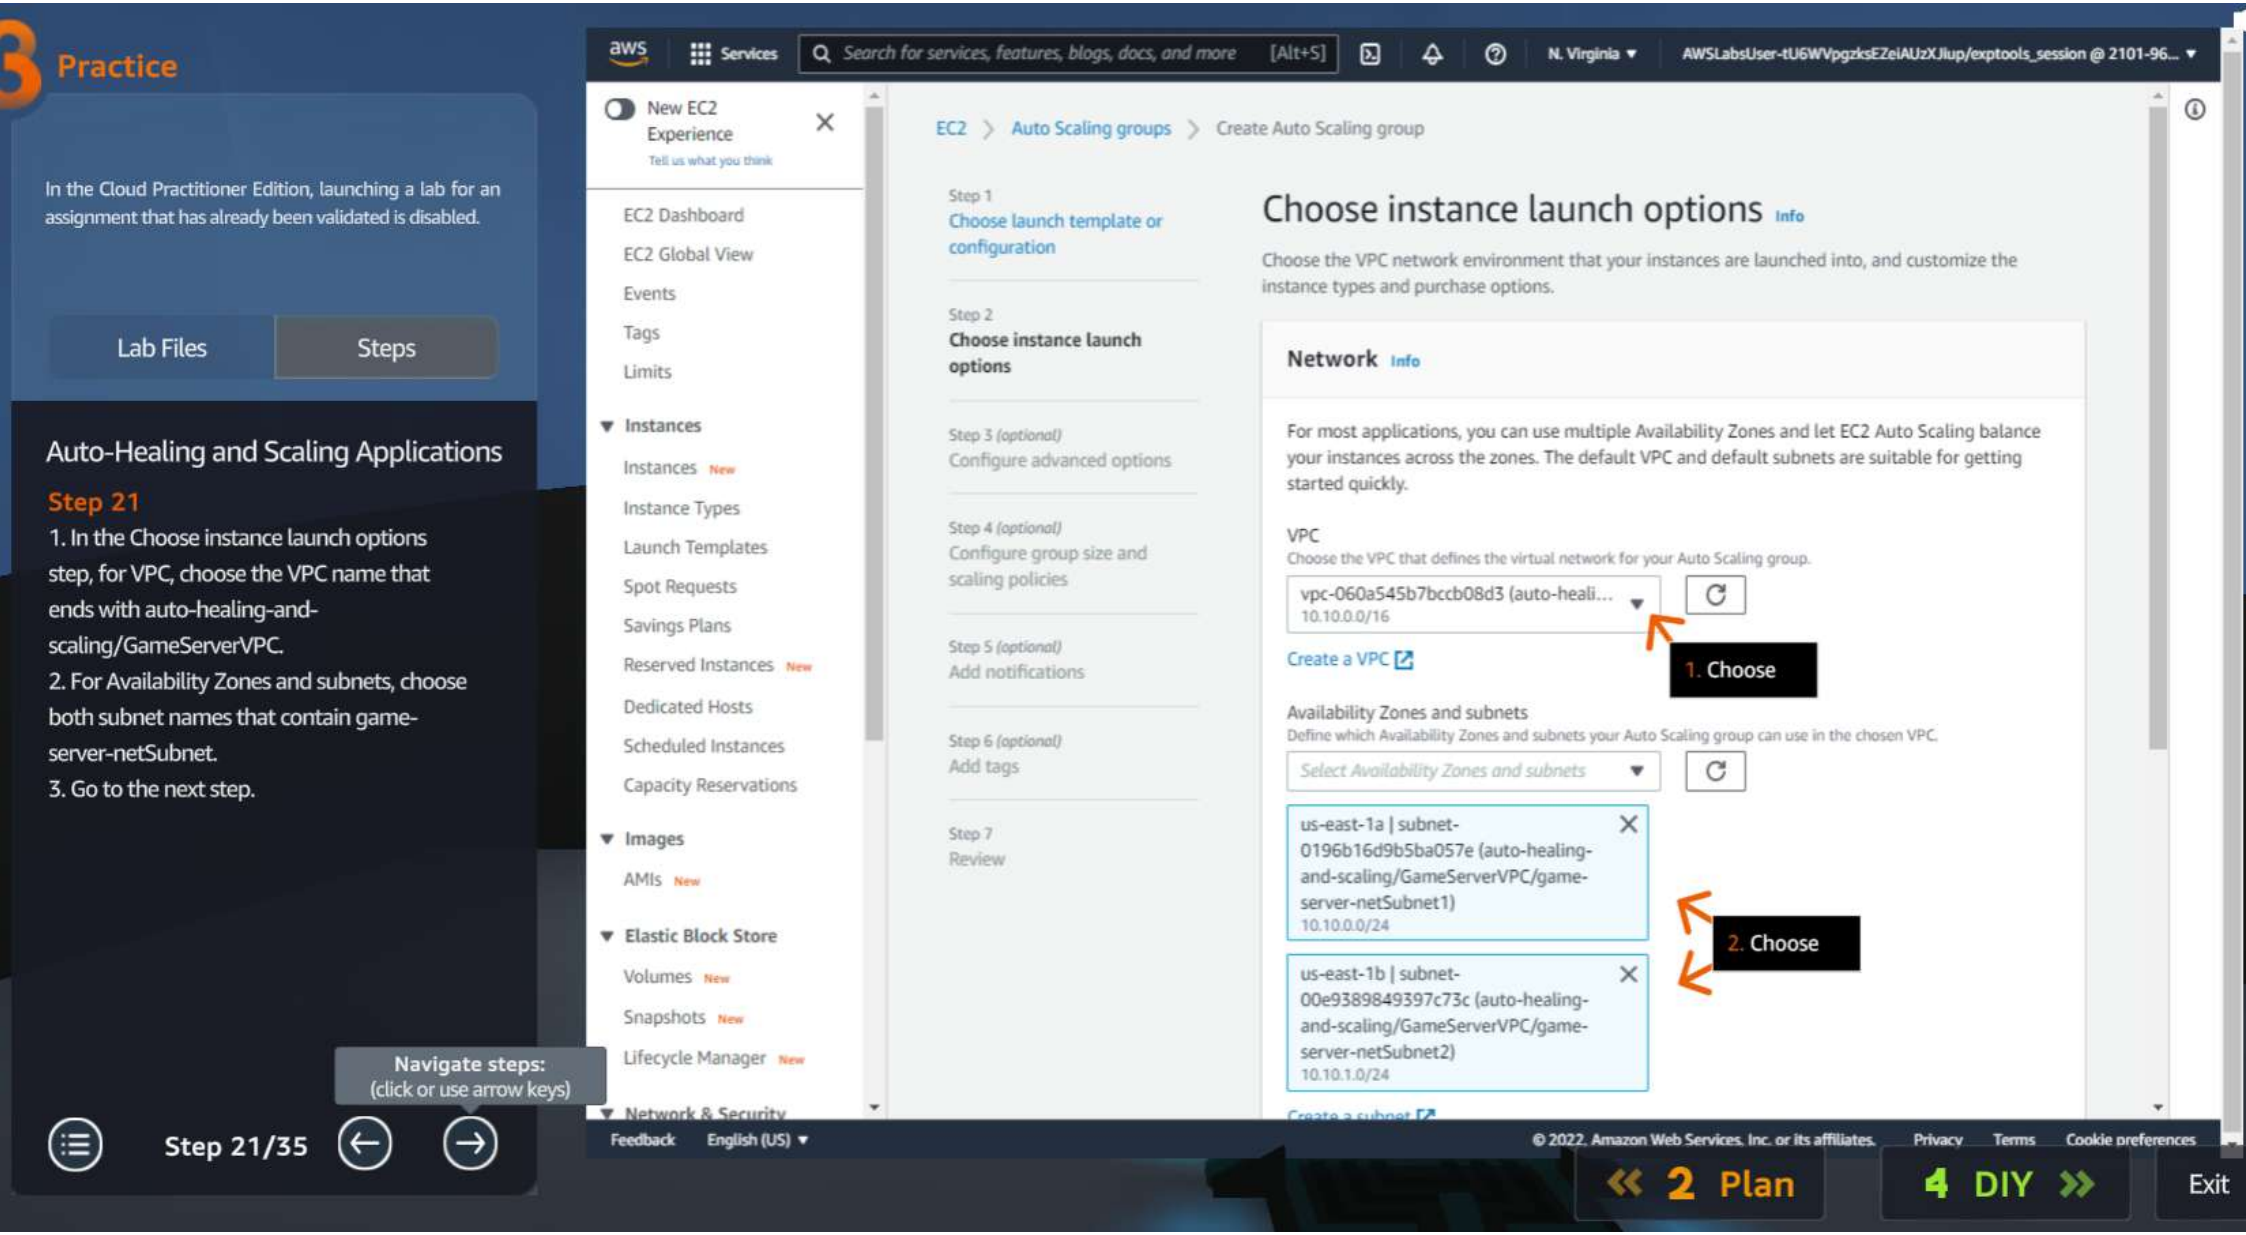
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 21
1. In the Choose instance launch options step, for VPC, choose the VPC name that ends with auto-healing-and-scaling/GameServerVPC.
2. For Availability Zones and subnets, choose both subnet names that contain game-server-netSubnet.
3. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 21/35
aws
Services
Search for services, features, blogs, docs, and more [Alt+S]
N. Virginia
AWSLabsUser-tU6WVpgzksEZeiAUzXJiup/exptools_session @ 2101-96...
New EC2 Experience
Tell us what you think
EC2 Dashboard
EC2 Global View
Events
Tags
Limits
Instances
Instances New
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances New
Dedicated Hosts
Scheduled Instances
Capacity Reservations
Images
AMIs New
Elastic Block Store
Volumes New
Snapshots New
Lifecycle Manager New
Network & Security
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1 Choose launch template or configuration
Step 2 Choose instance launch options
Step 3 (optional) Configure advanced options
Step 4 (optional) Configure group size and scaling policies
Step 5 (optional) Add notifications
Step 6 (optional) Add tags
Step 7 Review
Choose instance launch options Info
Choose the VPC network environment that your instances are launched into, and customize the instance types and purchase options.
Network Info
For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.
VPC
Choose the VPC that defines the virtual network for your Auto Scaling group.
vpc-060a545b7bccb08d3 (auto-heali...
10.10.0.0/16
Create a VPC
1. Choose
Availability Zones and subnets
Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.
Select Availability Zones and subnets
us-east-1a | subnet-0196b16d9b5ba057e (auto-healing-and-scaling/GameServerVPC/game-server-netSubnet1)
10.10.0.0/24
us-east-1b | subnet-00e9389849397c73c (auto-healing-and-scaling/GameServerVPC/game-server-netSubnet2)
10.10.1.0/24
2. Choose
Feedback
English (US)
© 2022, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
« 2 Plan
4 DIY »
Exit

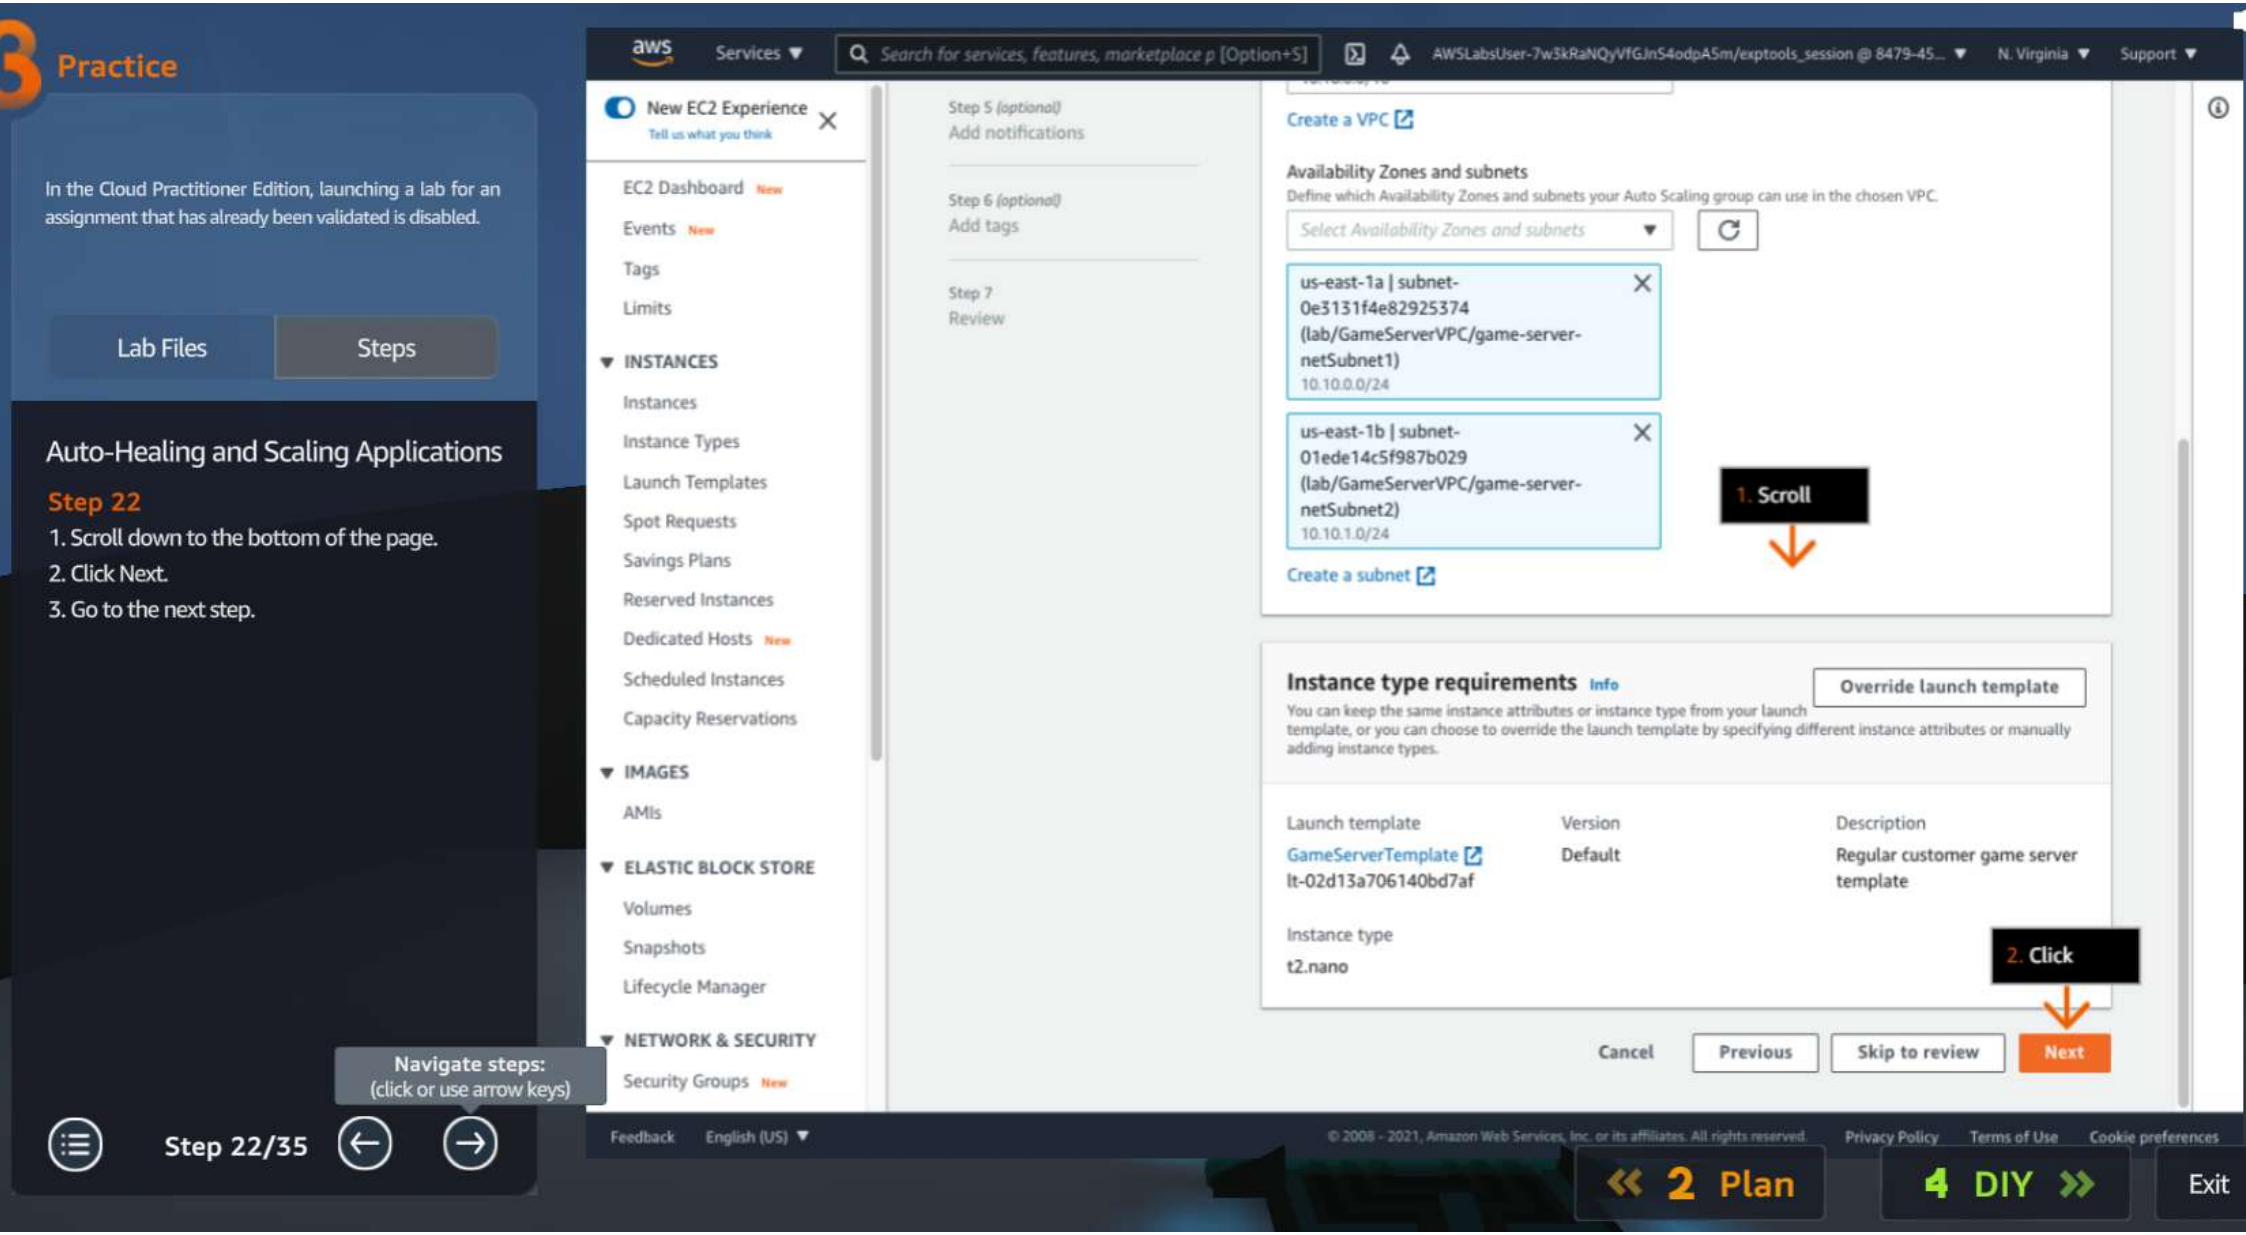
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 22
1. Scroll down to the bottom of the page.
2. Click Next.
3. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 22/35
aws
Services
Search for services, features, marketplace p [Option+S]
AWSLabsUser-7w3kRaNQyVfGJnS4odpA5m/exptools_session @ 8479-45...
N. Virginia
Support
New EC2 Experience
Tell us what you think
EC2 Dashboard New
Events New
Tags
Limits
INSTANCES
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts New
Scheduled Instances
Capacity Reservations
IMAGES
AMIs
ELASTIC BLOCK STORE
Volumes
Snapshots
Lifecycle Manager
NETWORK & SECURITY
Security Groups New
Step 5 (optional)
Add notifications
Step 6 (optional)
Add tags
Step 7
Review
Create a VPC
Availability Zones and subnets
Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.
Select Availability Zones and subnets
us-east-1a | subnet-0e3131f4e82925374 (lab/GameServerVPC/game-server-netSubnet1)
10.10.0.0/24
us-east-1b | subnet-01ede14c5f987b029 (lab/GameServerVPC/game-server-netSubnet2)
10.10.1.0/24
1. Scroll
Create a subnet
Instance type requirements Info
Override launch template
You can keep the same instance attributes or instance type from your launch template, or you can choose to override the launch template by specifying different instance attributes or manually adding instance types.
Launch template
GameServerTemplate
lt-02d13a706140bd7af
Version
Default
Description
Regular customer game server template
Instance type
t2.nano
2. Click
Cancel
Previous
Skip to review
Next
Feedback
English (US)
© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
Cookie preferences
« 2 Plan
4 DIY »
Exit

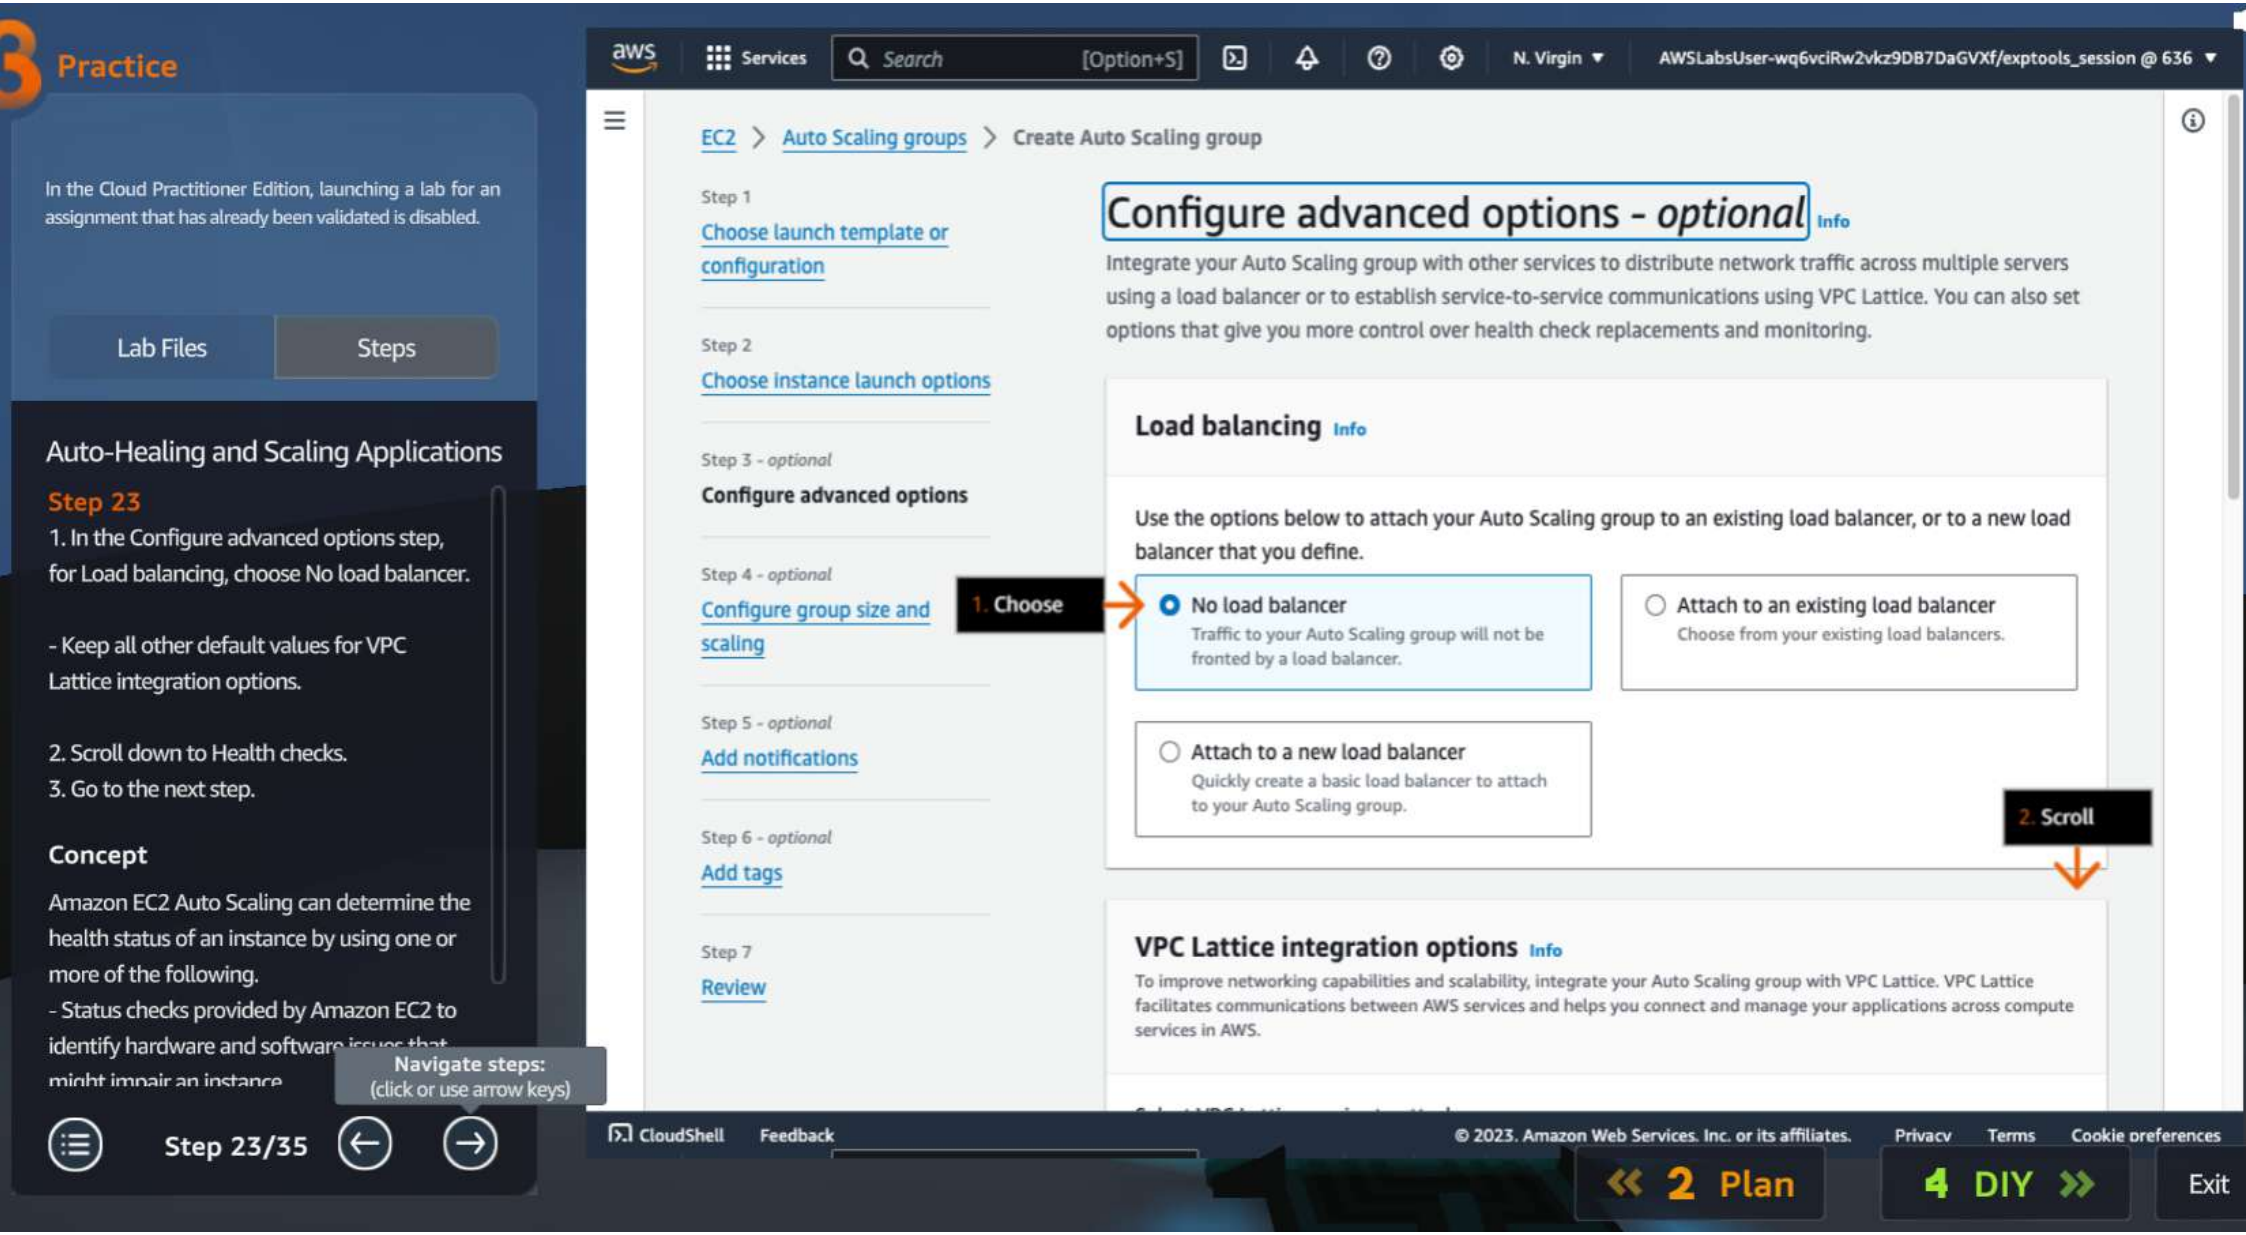
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 23
1. In the Configure advanced options step, for Load balancing, choose No load balancer.
- Keep all other default values for VPC Lattice integration options.
2. Scroll down to Health checks.
3. Go to the next step.
Concept
Amazon EC2 Auto Scaling can determine the health status of an instance by using one or more of the following.
- Status checks provided by Amazon EC2 to identify hardware and software issues that might impair an instance
Navigate steps: (click or use arrow keys)
Step 23/35
aws
Services
Search
[Option+S]
N. Virgin
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 636
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template or configuration
Step 2
Choose instance launch options
Step 3 - optional
Configure advanced options
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Configure advanced options - optional
Info
Integrate your Auto Scaling group with other services to distribute network traffic across multiple servers using a load balancer or to establish service-to-service communications using VPC Lattice. You can also set options that give you more control over health check replacements and monitoring.
Load balancing
Info
Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.
1. Choose
No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.
Attach to an existing load balancer
Choose from your existing load balancers.
Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.
2. Scroll
VPC Lattice integration options
Info
To improve networking capabilities and scalability, integrate your Auto Scaling group with VPC Lattice. VPC Lattice facilitates communications between AWS services and helps you connect and manage your applications across compute services in AWS.
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
« 2 Plan
4 DIY »
Exit

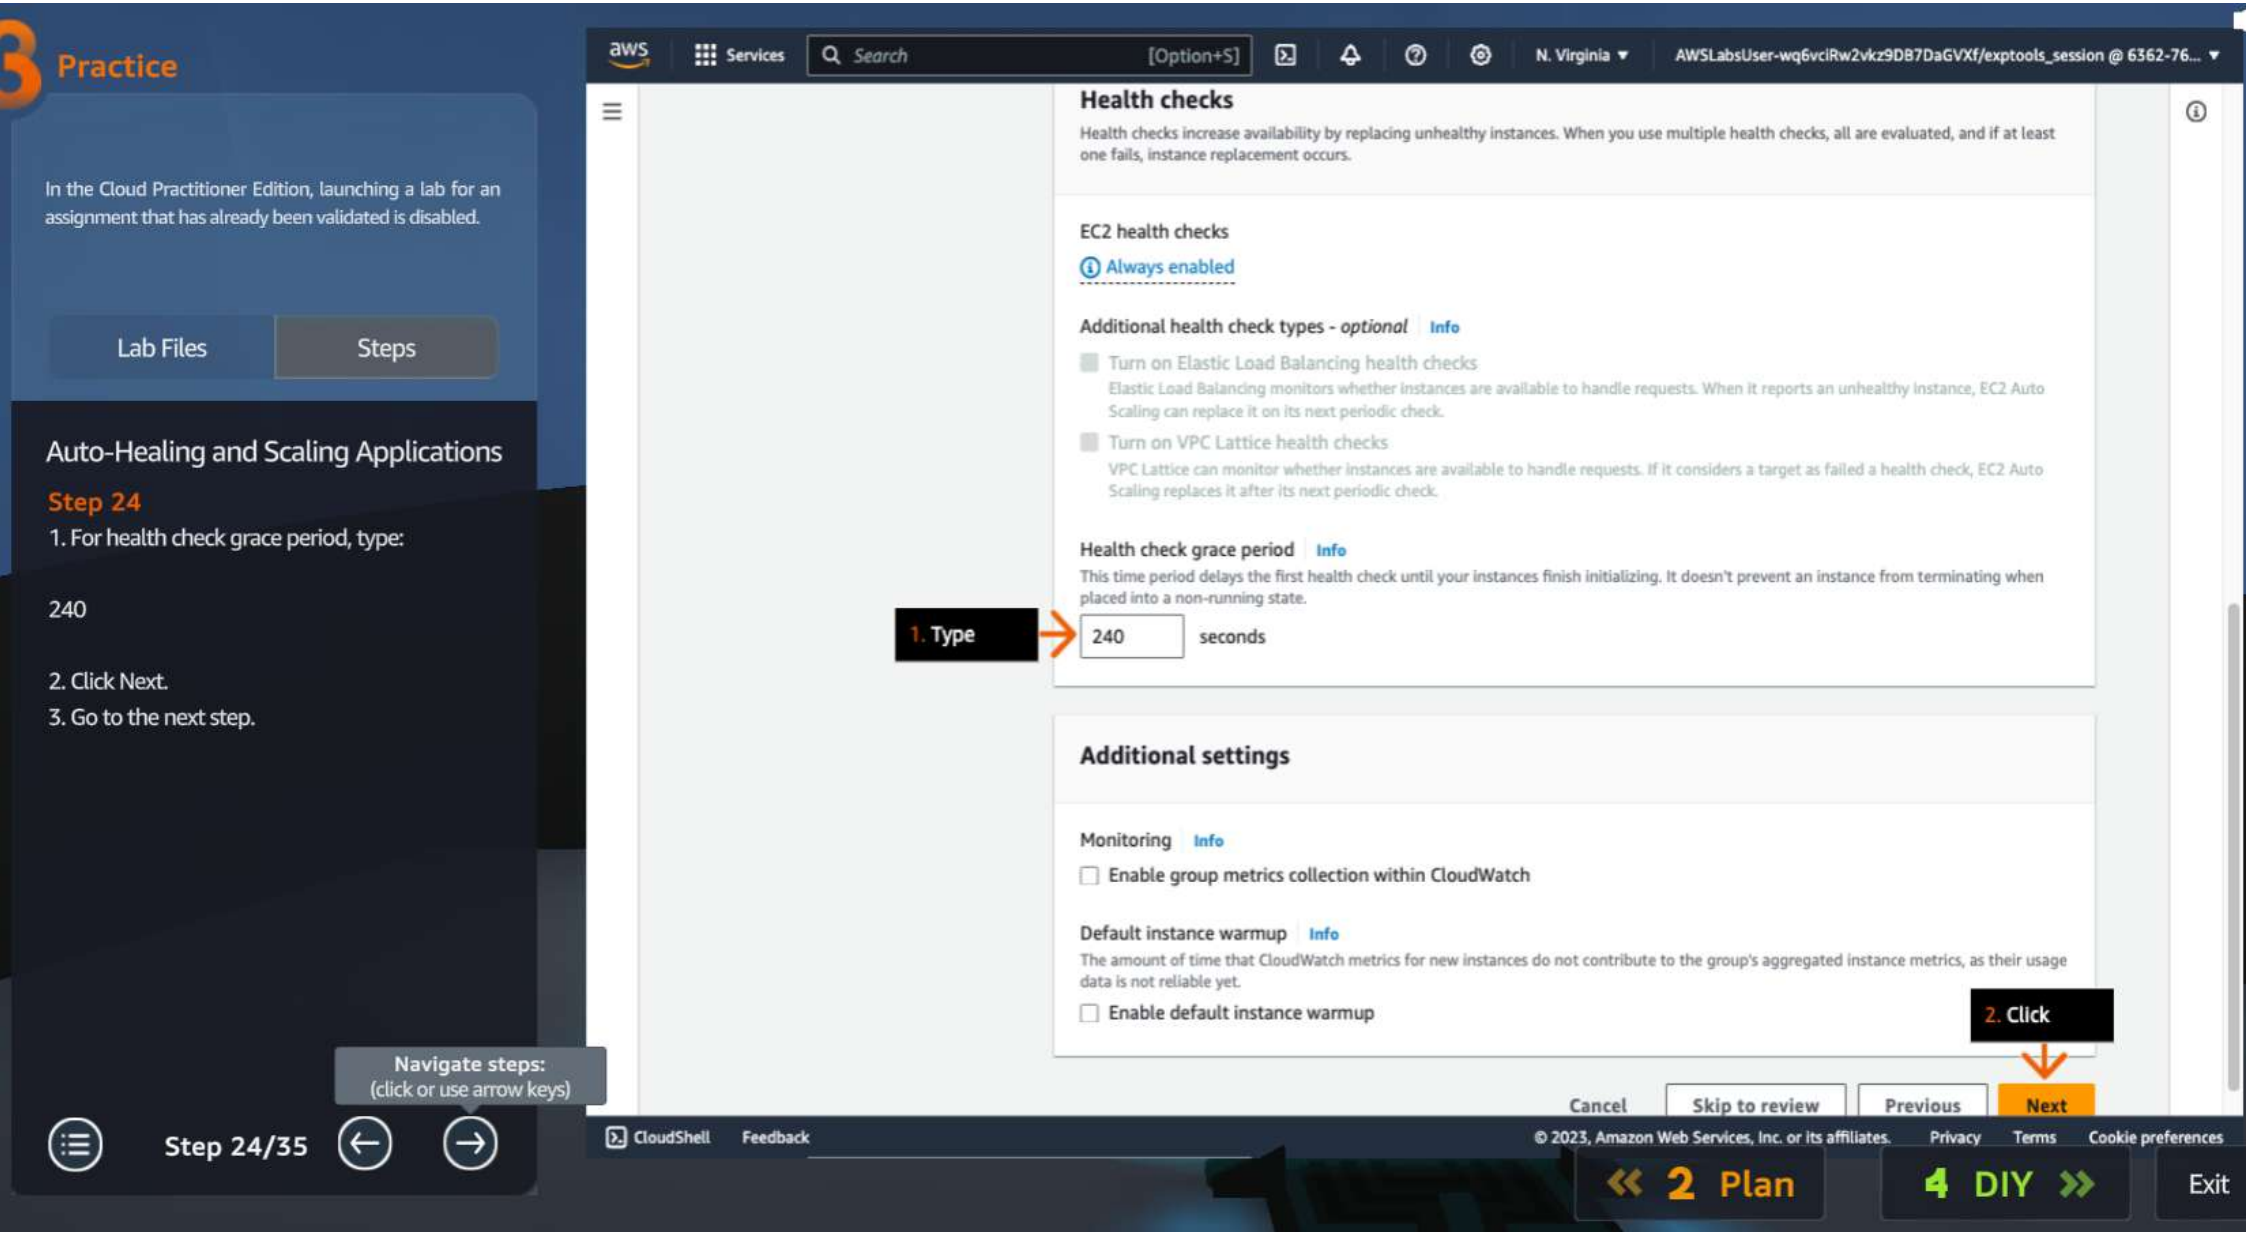
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 24
1. For health check grace period, type:
240
2. Click Next.
3. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 24/35
aws
Services
Search
[Option+S]
N. Virginia
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 6362-76...
Health checks
Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.
EC2 health checks
Always enabled
Additional health check types - optional
Info
Turn on Elastic Load Balancing health checks
Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.
Turn on VPC Lattice health checks
VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.
Health check grace period
Info
This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.
1. Type
240
seconds
Additional settings
Monitoring
Info
Enable group metrics collection within CloudWatch
Default instance warmup
Info
The amount of time that CloudWatch metrics for new instances do not contribute to the group's aggregated instance metrics, as their usage data is not reliable yet.
Enable default instance warmup
2. Click
Cancel
Skip to review
Previous
Next
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

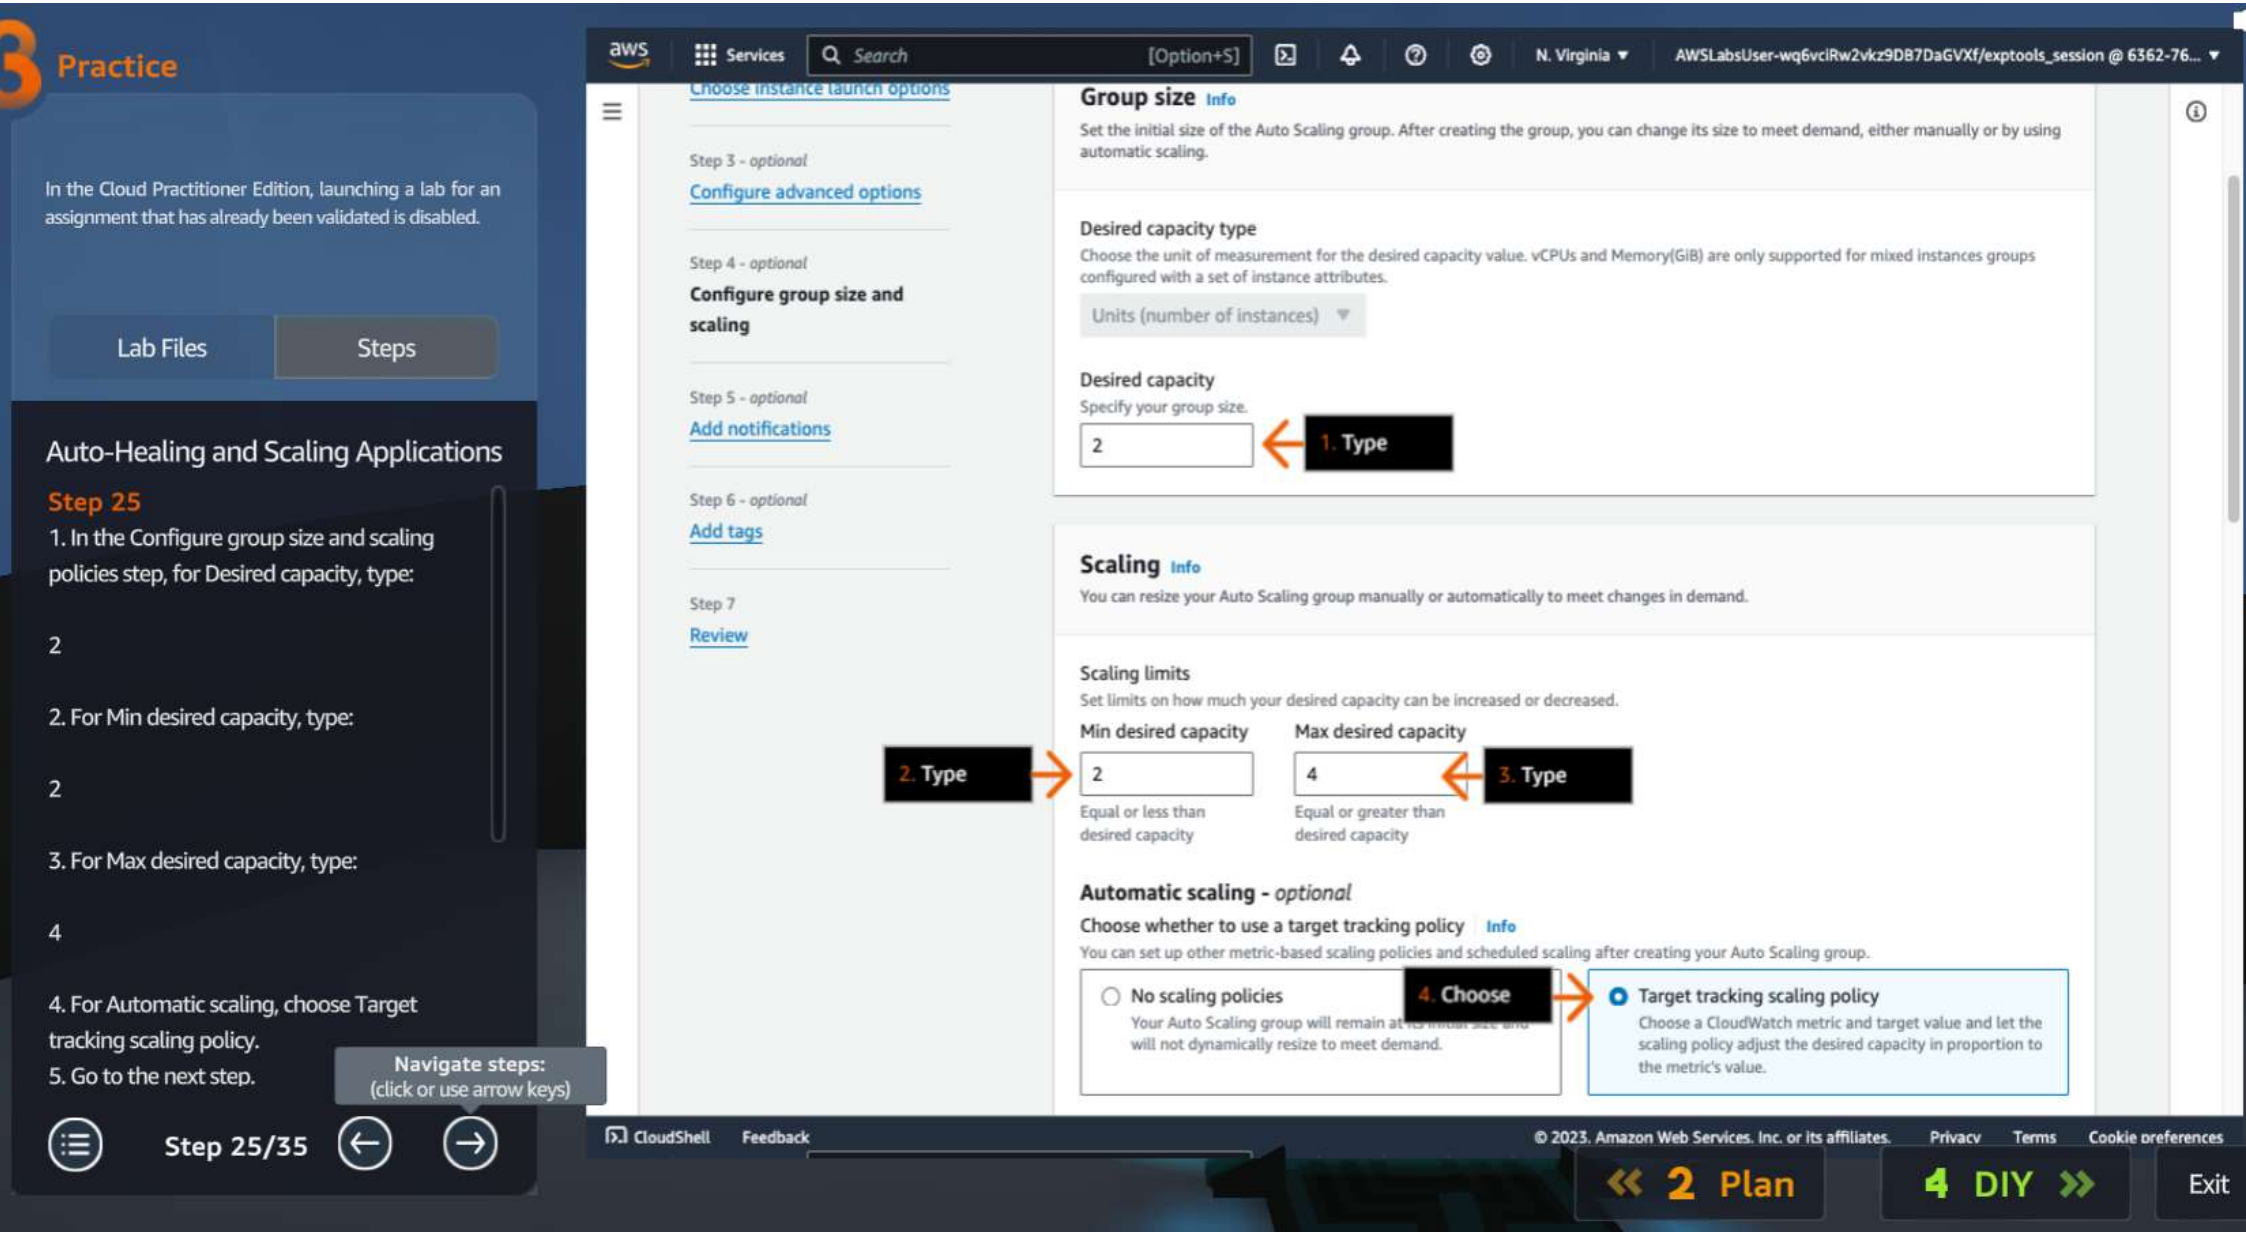
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 25
1. In the Configure group size and scaling policies step, for Desired capacity, type:
2
2. For Min desired capacity, type:
2
3. For Max desired capacity, type:
4
4. For Automatic scaling, choose Target tracking scaling policy.
5. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 25/35
aws
Services
Search
[Option+S]
N. Virginia
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 6362-76...
Choose instance launch options
Step 3 - optional
Configure advanced options
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Group size Info
Set the initial size of the Auto Scaling group. After creating the group, you can change its size to meet demand, either manually or by using automatic scaling.
Desired capacity type
Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed instances groups configured with a set of instance attributes.
Units (number of instances)
Desired capacity
Specify your group size.
2
1. Type
Scaling Info
You can resize your Auto Scaling group manually or automatically to meet changes in demand.
Scaling limits
Set limits on how much your desired capacity can be increased or decreased.
Min desired capacity
2. Type
2
Equal or less than desired capacity
Max desired capacity
3. Type
4
Equal or greater than desired capacity
Automatic scaling - optional
Choose whether to use a target tracking policy Info
You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.
No scaling policies
Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.
4. Choose
Target tracking scaling policy
Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

Practice

In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.

Lab Files | Steps

Auto-Healing and Scaling Applications

Step 26

1. Scroll down to Automatic scaling.
2. For Scaling policy name, type:

CPU Utilization

3. For Metric type, choose Average CPU utilization.
4. For Target value, type:

70

5. Scroll down to the bottom of the page, and then click Next (not shown).
6. Go to the next step.

Navigate steps: (click or use arrow keys)

Step 26/35

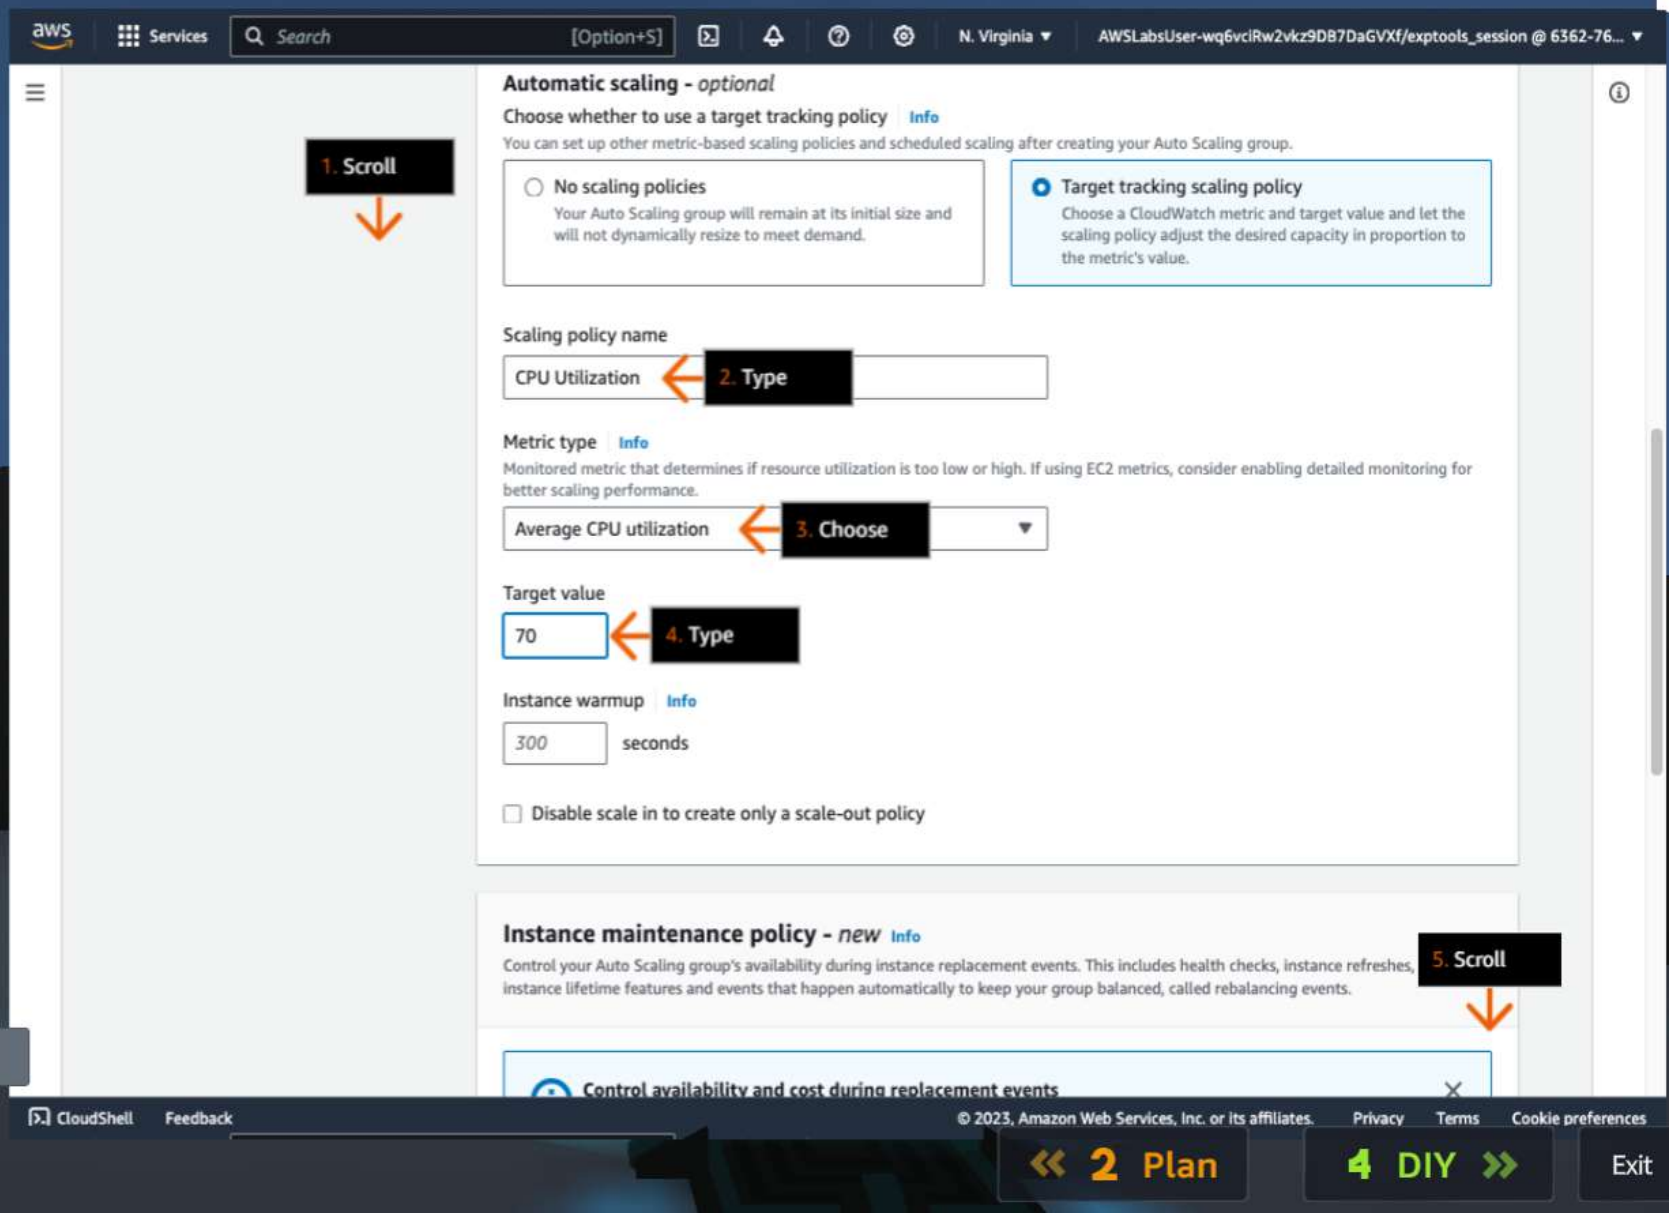

« 2 Plan | 4 DIY » | Exit

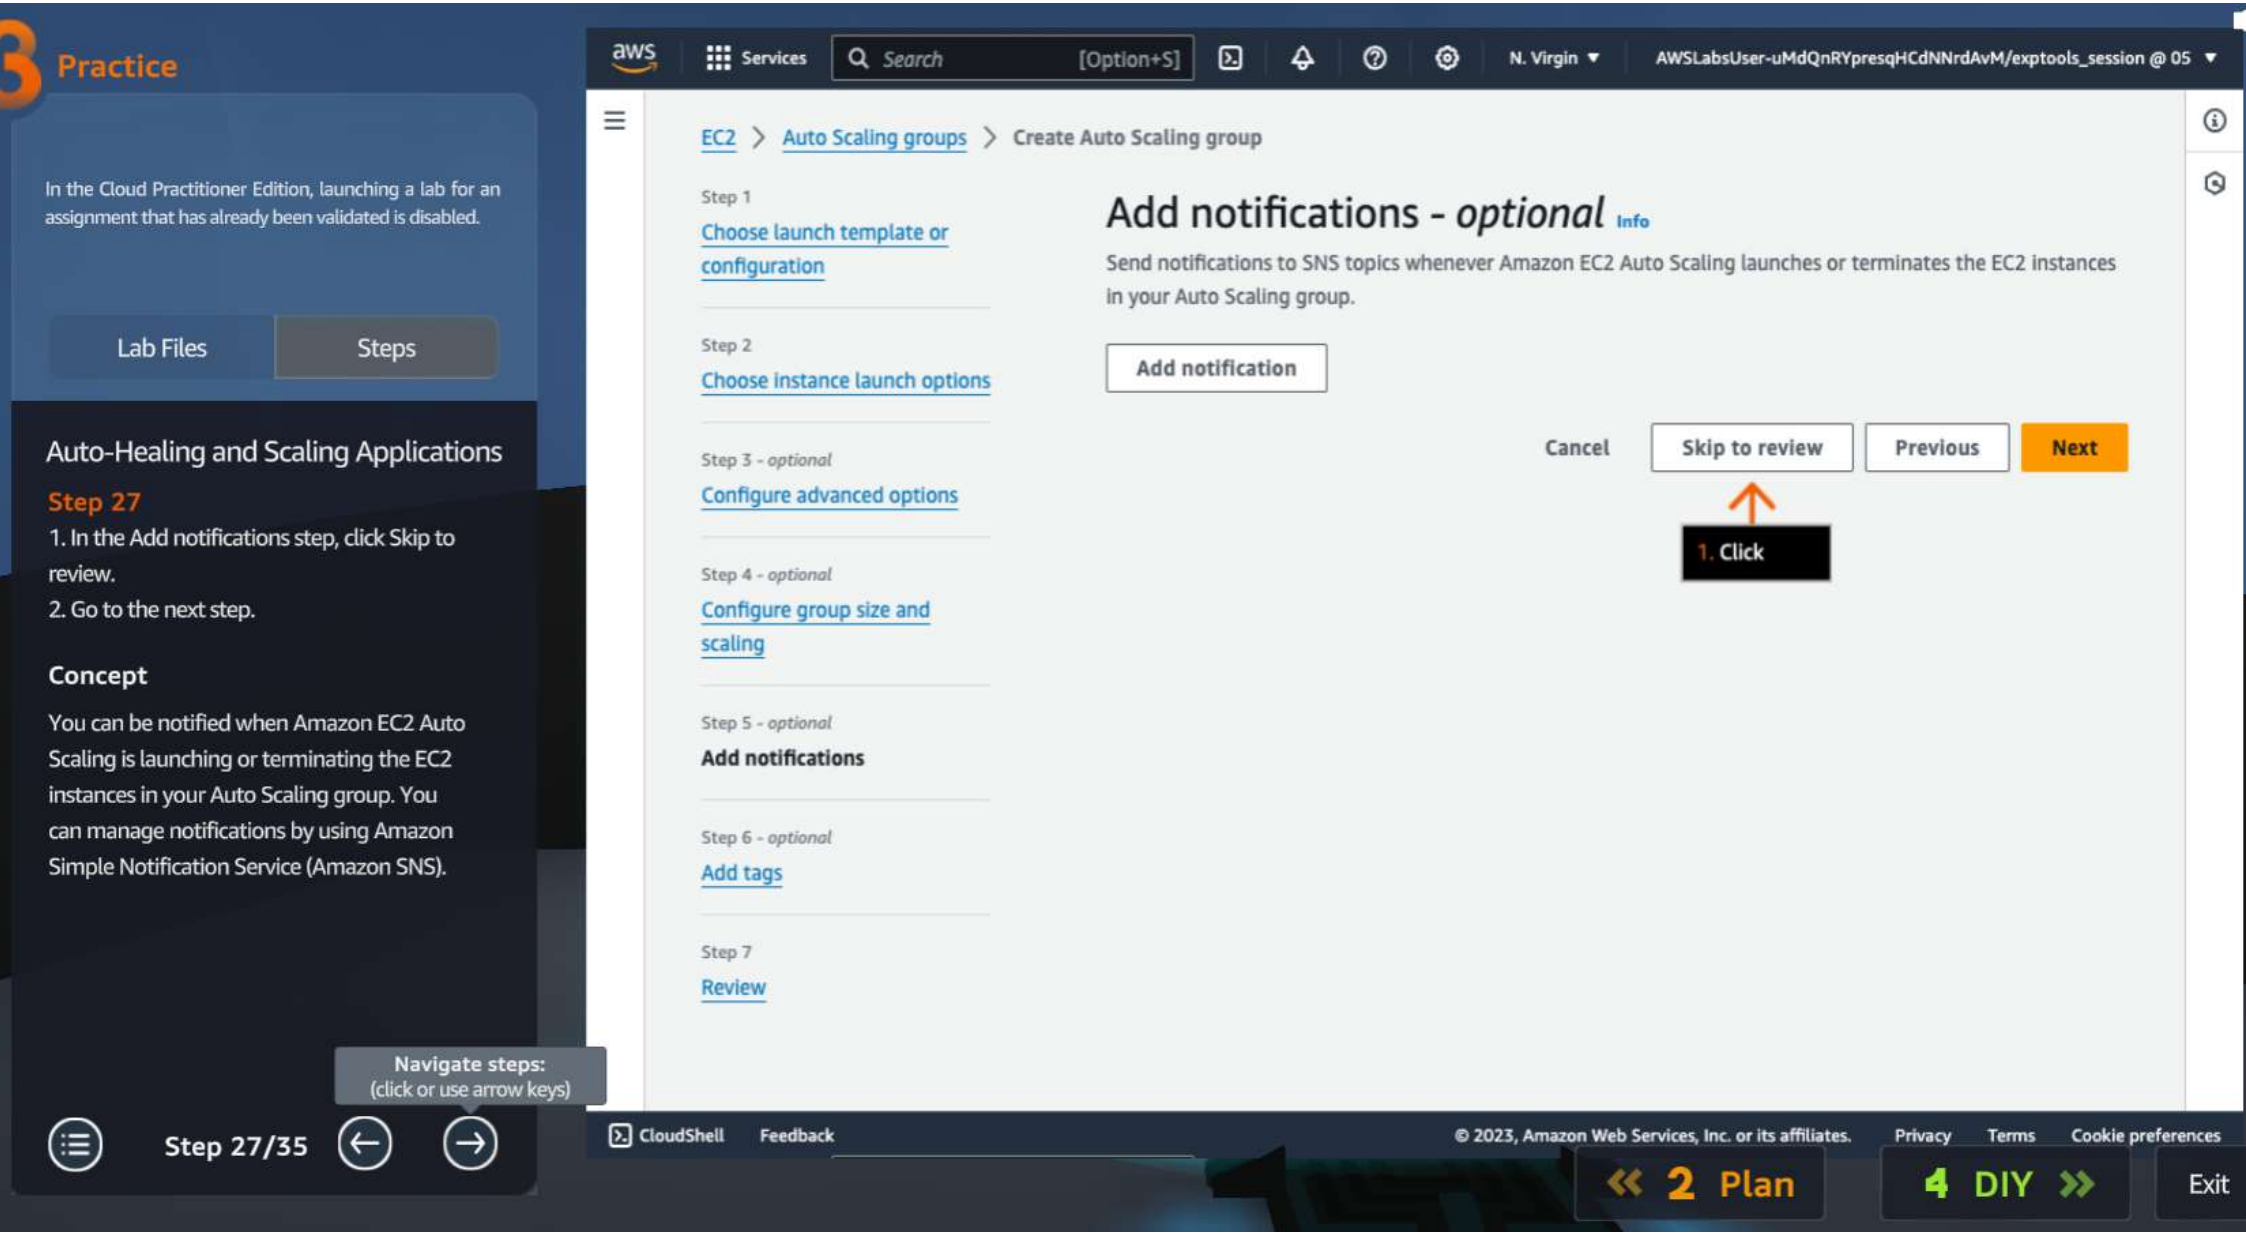
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 27
1. In the Add notifications step, click Skip to review.
2. Go to the next step.
Concept
You can be notified when Amazon EC2 Auto Scaling is launching or terminating the EC2 instances in your Auto Scaling group. You can manage notifications by using Amazon Simple Notification Service (Amazon SNS).
Navigate steps: (click or use arrow keys)
Step 27/35
aws
Services
Search
[Option+S]
N. Virgin
AWSLabsUser-uMdQnRYpresqHCdNNrdAvM/exptools_session @ 05
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template or configuration
Step 2
Choose instance launch options
Step 3 - optional
Configure advanced options
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Add notifications - optional Info
Send notifications to SNS topics whenever Amazon EC2 Auto Scaling launches or terminates the EC2 instances in your Auto Scaling group.
Add notification
Cancel
Skip to review
Previous
Next
1. Click
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

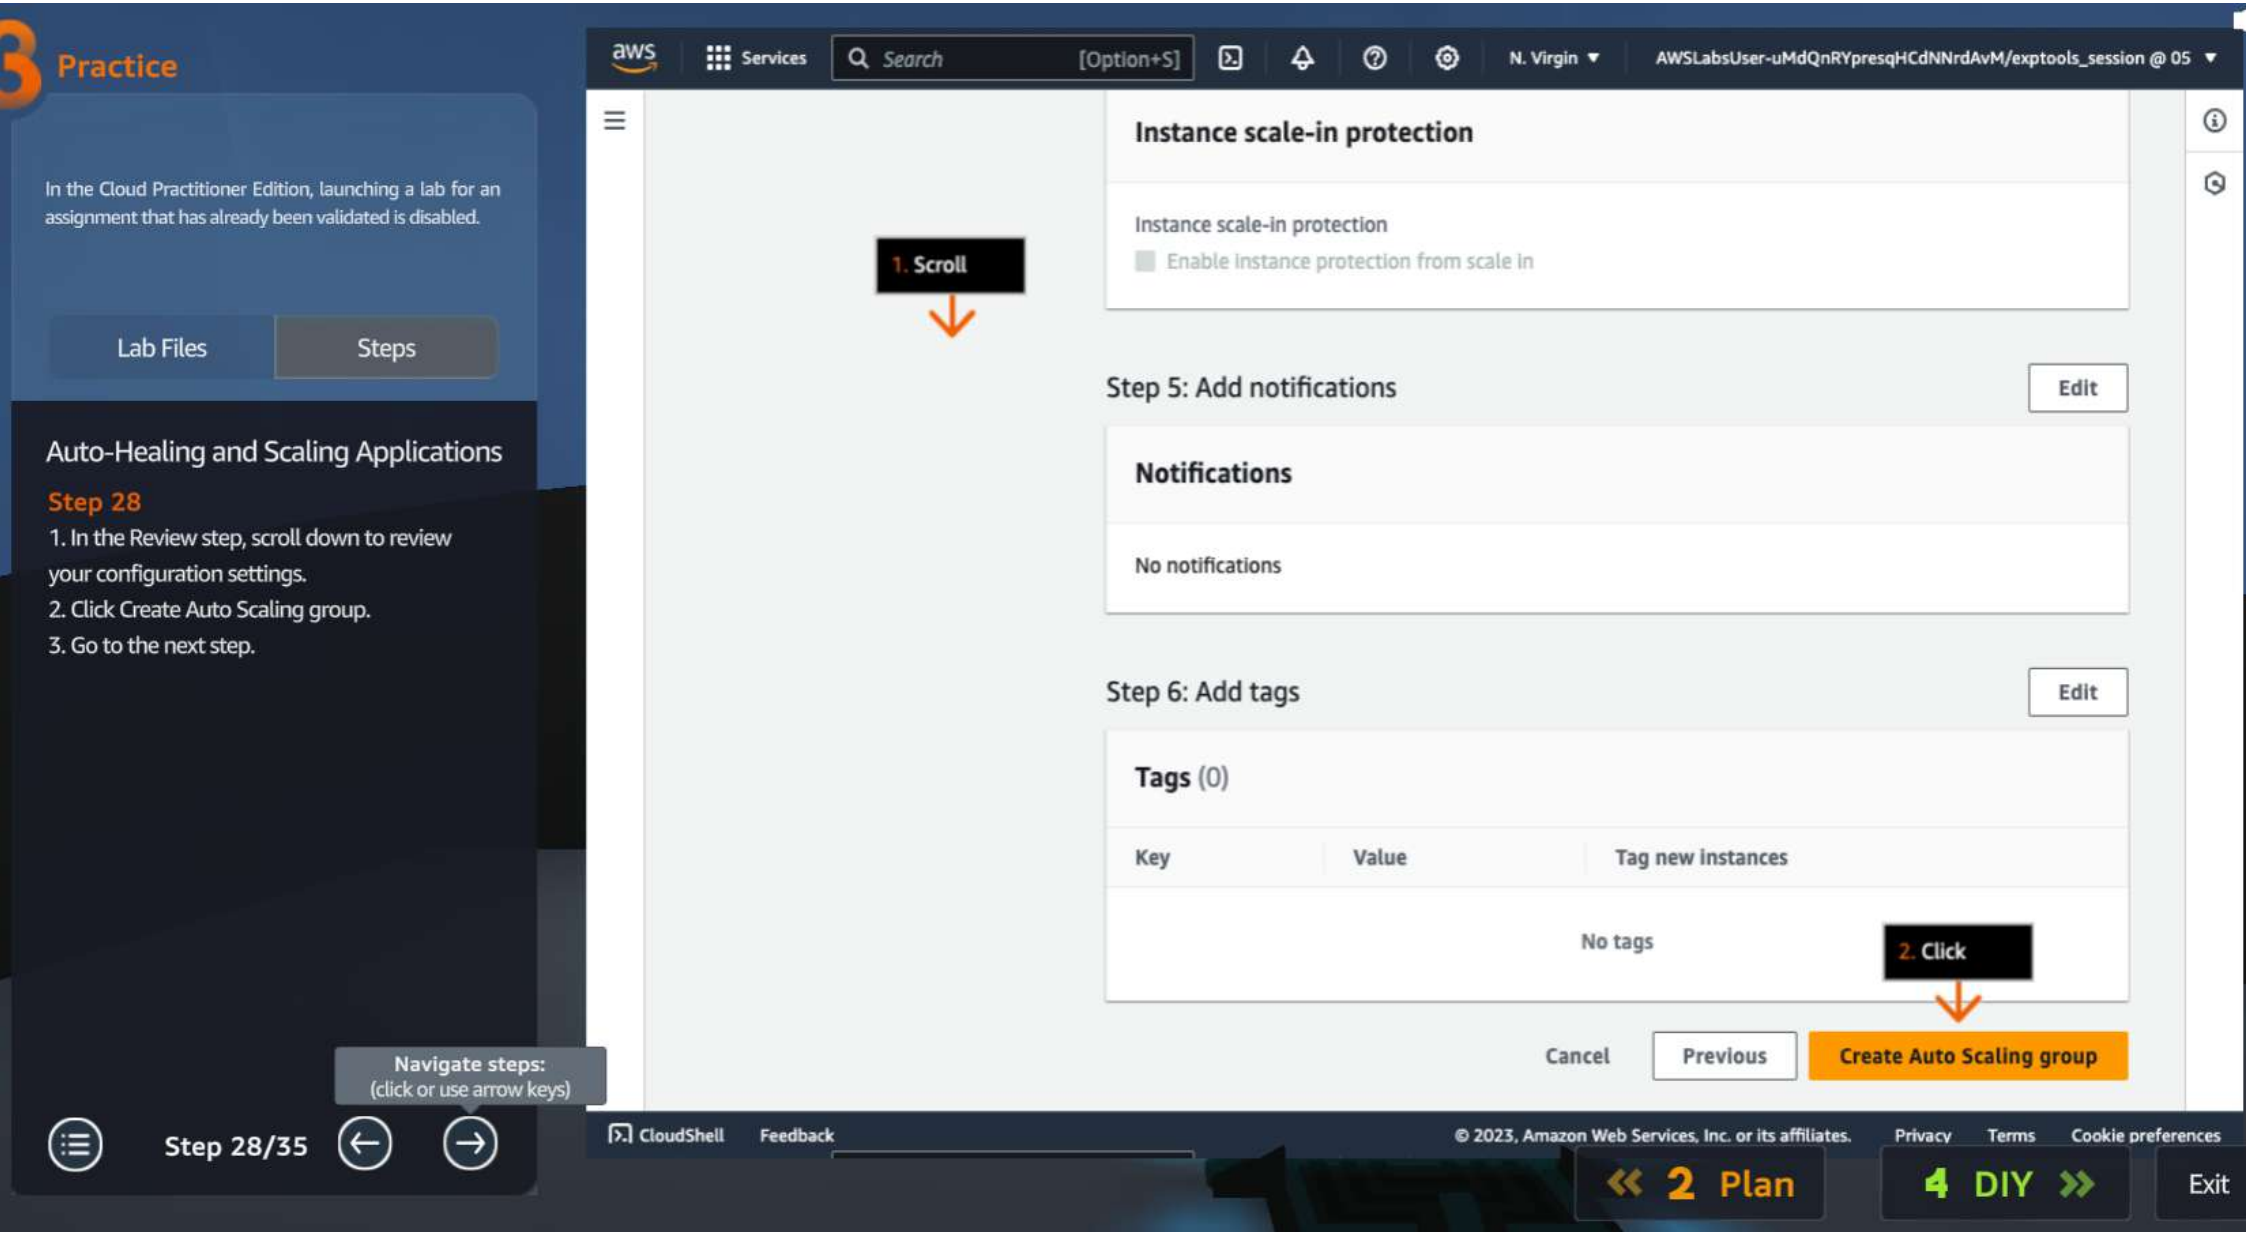
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 28
1. In the Review step, scroll down to review your configuration settings.
2. Click Create Auto Scaling group.
3. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 28/35
aws
Services
Search
[Option+S]
N. Virgin
AWSLabsUser-uMdQnRYpresqHCdNNrdAvM/exptools_session @ 05
1. Scroll
Instance scale-in protection
Instance scale-in protection
Enable instance protection from scale in
Step 5: Add notifications
Edit
Notifications
No notifications
Step 6: Add tags
Edit
Tags (0)
Key
Value
Tag new instances
No tags
2. Click
Cancel
Previous
Create Auto Scaling group
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

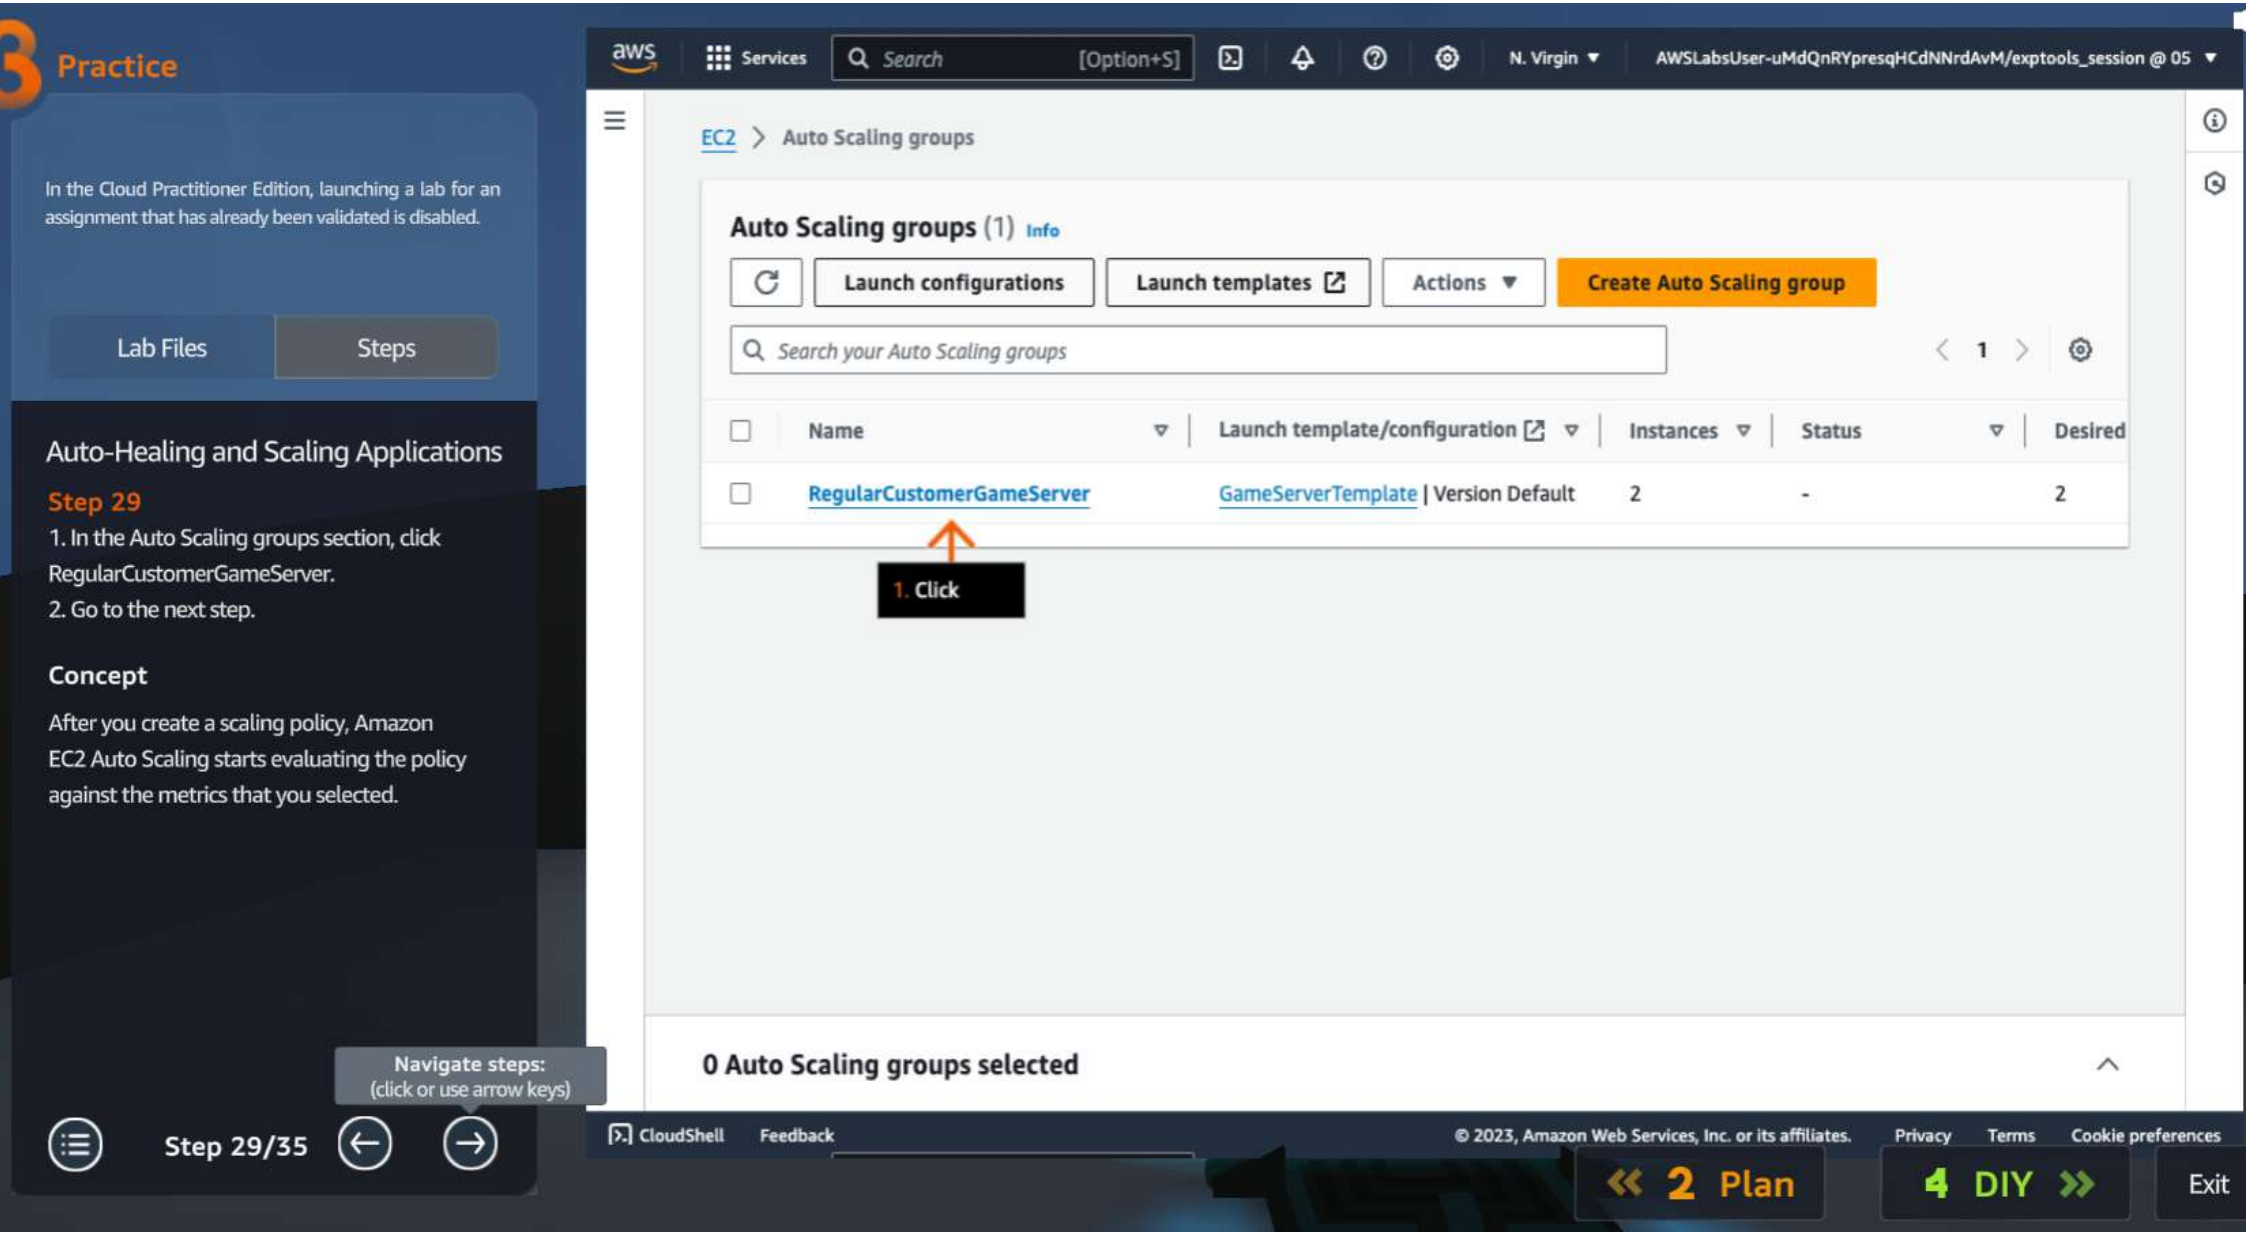
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 29
1. In the Auto Scaling groups section, click RegularCustomerGameServer.
2. Go to the next step.
Concept
After you create a scaling policy, Amazon EC2 Auto Scaling starts evaluating the policy against the metrics that you selected.
Navigate steps: (click or use arrow keys)
Step 29/35
aws
Services
Search
[Option+S]
N. Virgin
AWSLabsUser-uMdQnRYpresqHCdNNrdAvM/exptools_session @ 05
EC2
Auto Scaling groups
Auto Scaling groups (1) Info
Launch configurations
Launch templates
Actions
Create Auto Scaling group
Search your Auto Scaling groups
1
Name
Launch template/configuration
Instances
Status
Desired
RegularCustomerGameServer
GameServerTemplate | Version Default
2
-
2
1. Click
0 Auto Scaling groups selected
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

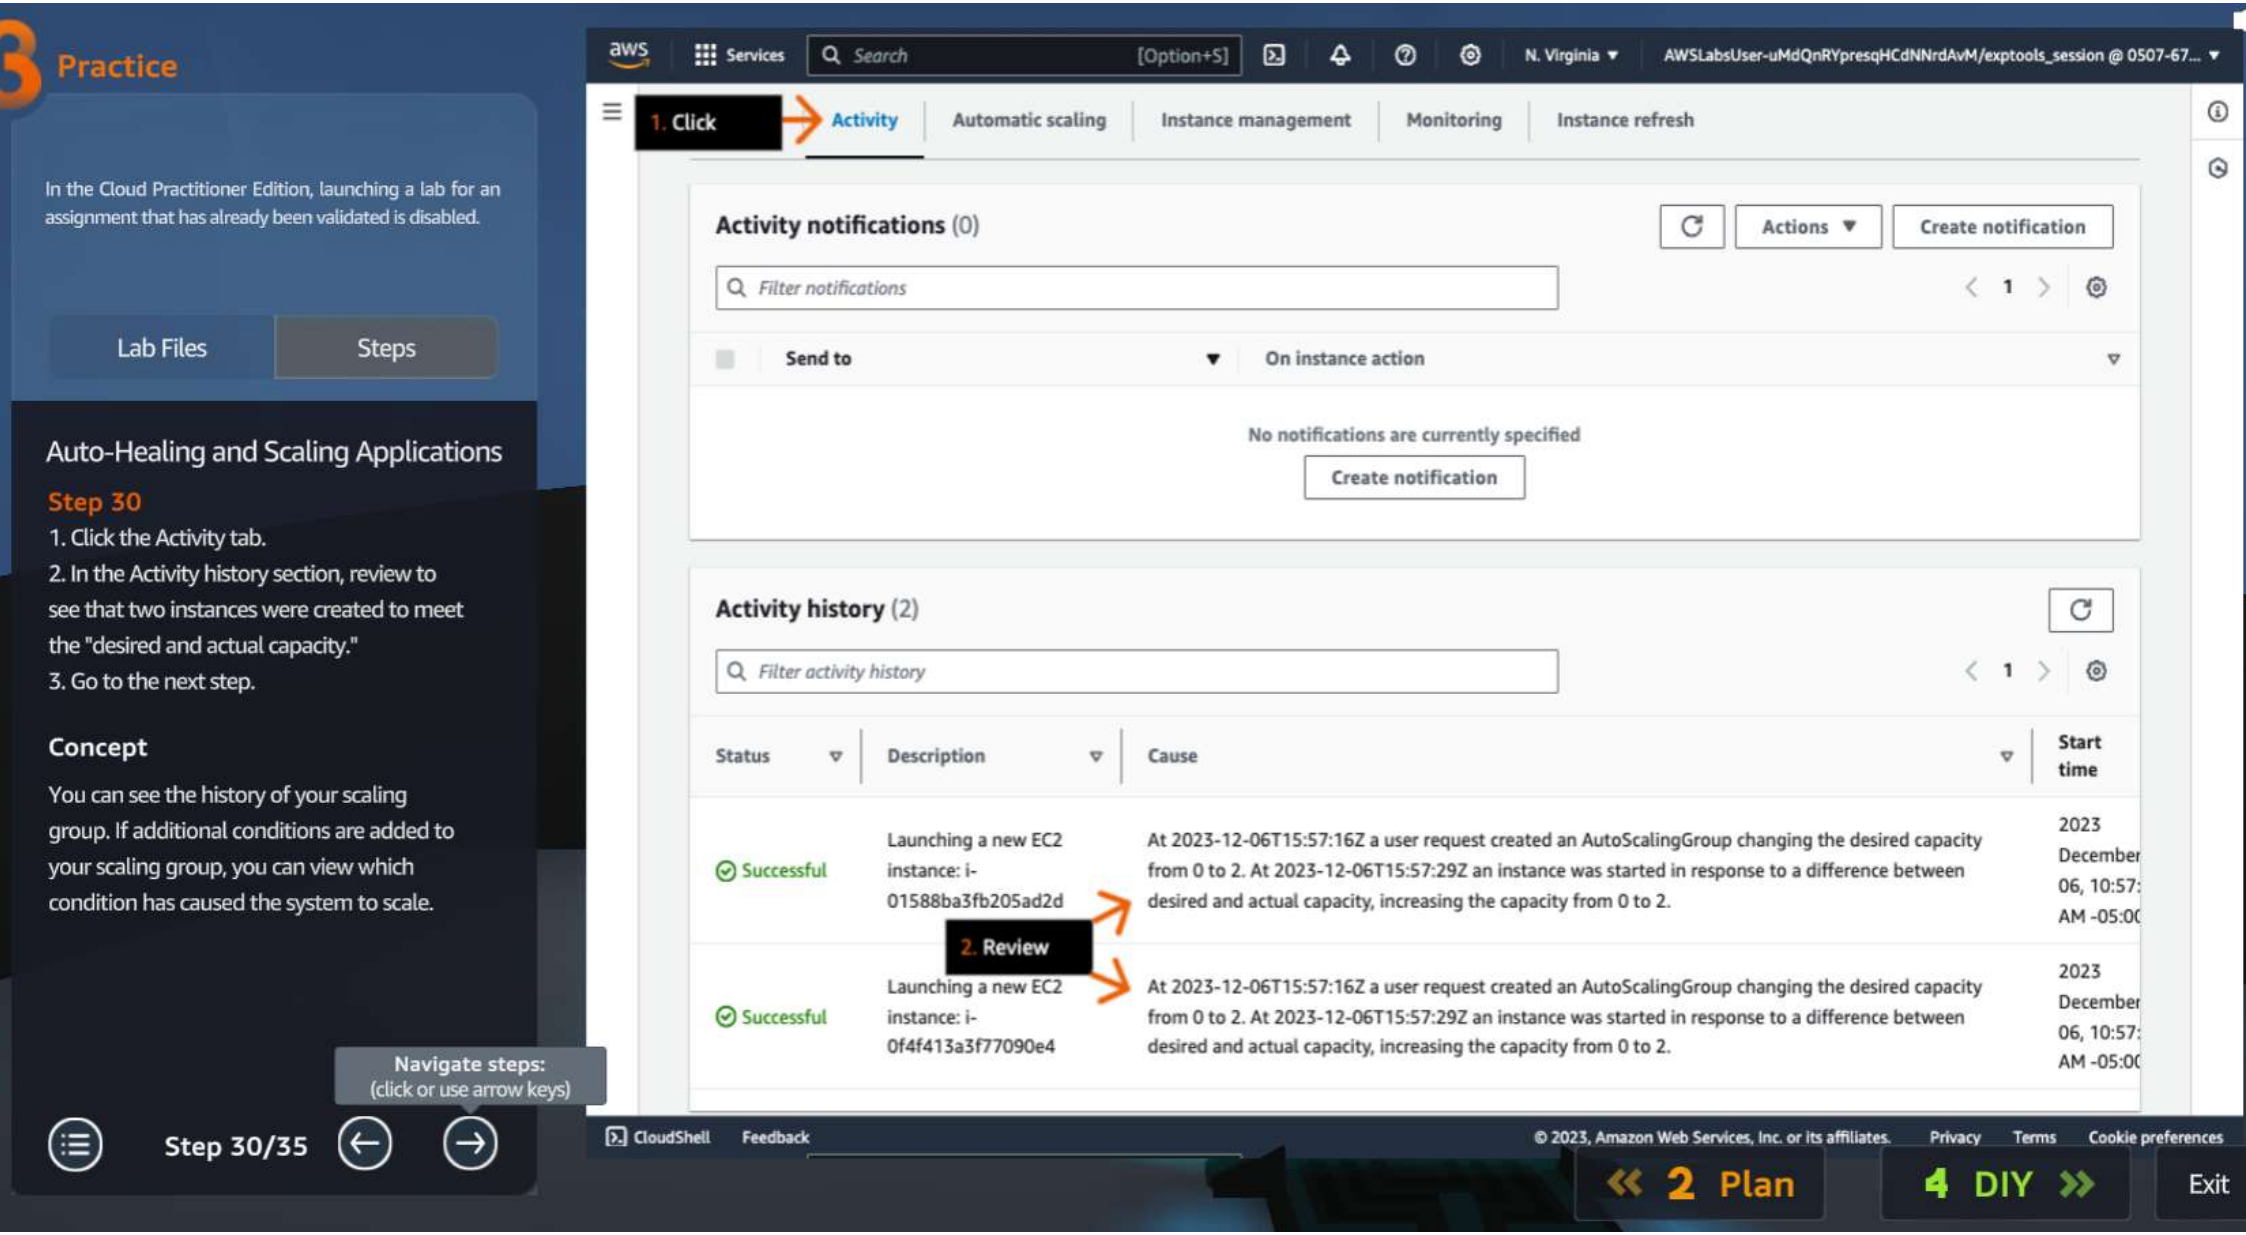
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 30
1. Click the Activity tab.
2. In the Activity history section, review to see that two instances were created to meet the "desired and actual capacity."
3. Go to the next step.
Concept
You can see the history of your scaling group. If additional conditions are added to your scaling group, you can view which condition has caused the system to scale.
Navigate steps: (click or use arrow keys)
Step 30/35
aws
Services
Search
[Option+S]
N. Virginia
AWSLabsUser-uMdQnRYpresqHCdNNrdAvM/exptools_session @ 0507-67...
1. Click
Activity
Automatic scaling
Instance management
Monitoring
Instance refresh
Activity notifications (0)
Actions
Create notification
Filter notifications
Send to
On instance action
No notifications are currently specified
Create notification
Activity history (2)
Filter activity history
Status
Description
Cause
Start time
Successful
Launching a new EC2 instance: i-01588ba3fb205ad2d
At 2023-12-06T15:57:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 2. At 2023-12-06T15:57:29Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 2.
2023 December 06, 10:57: AM -05:00
2. Review
Successful
Launching a new EC2 instance: i-0f4f413a3f77090e4
At 2023-12-06T15:57:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 2. At 2023-12-06T15:57:29Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 2.
2023 December 06, 10:57: AM -05:00
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

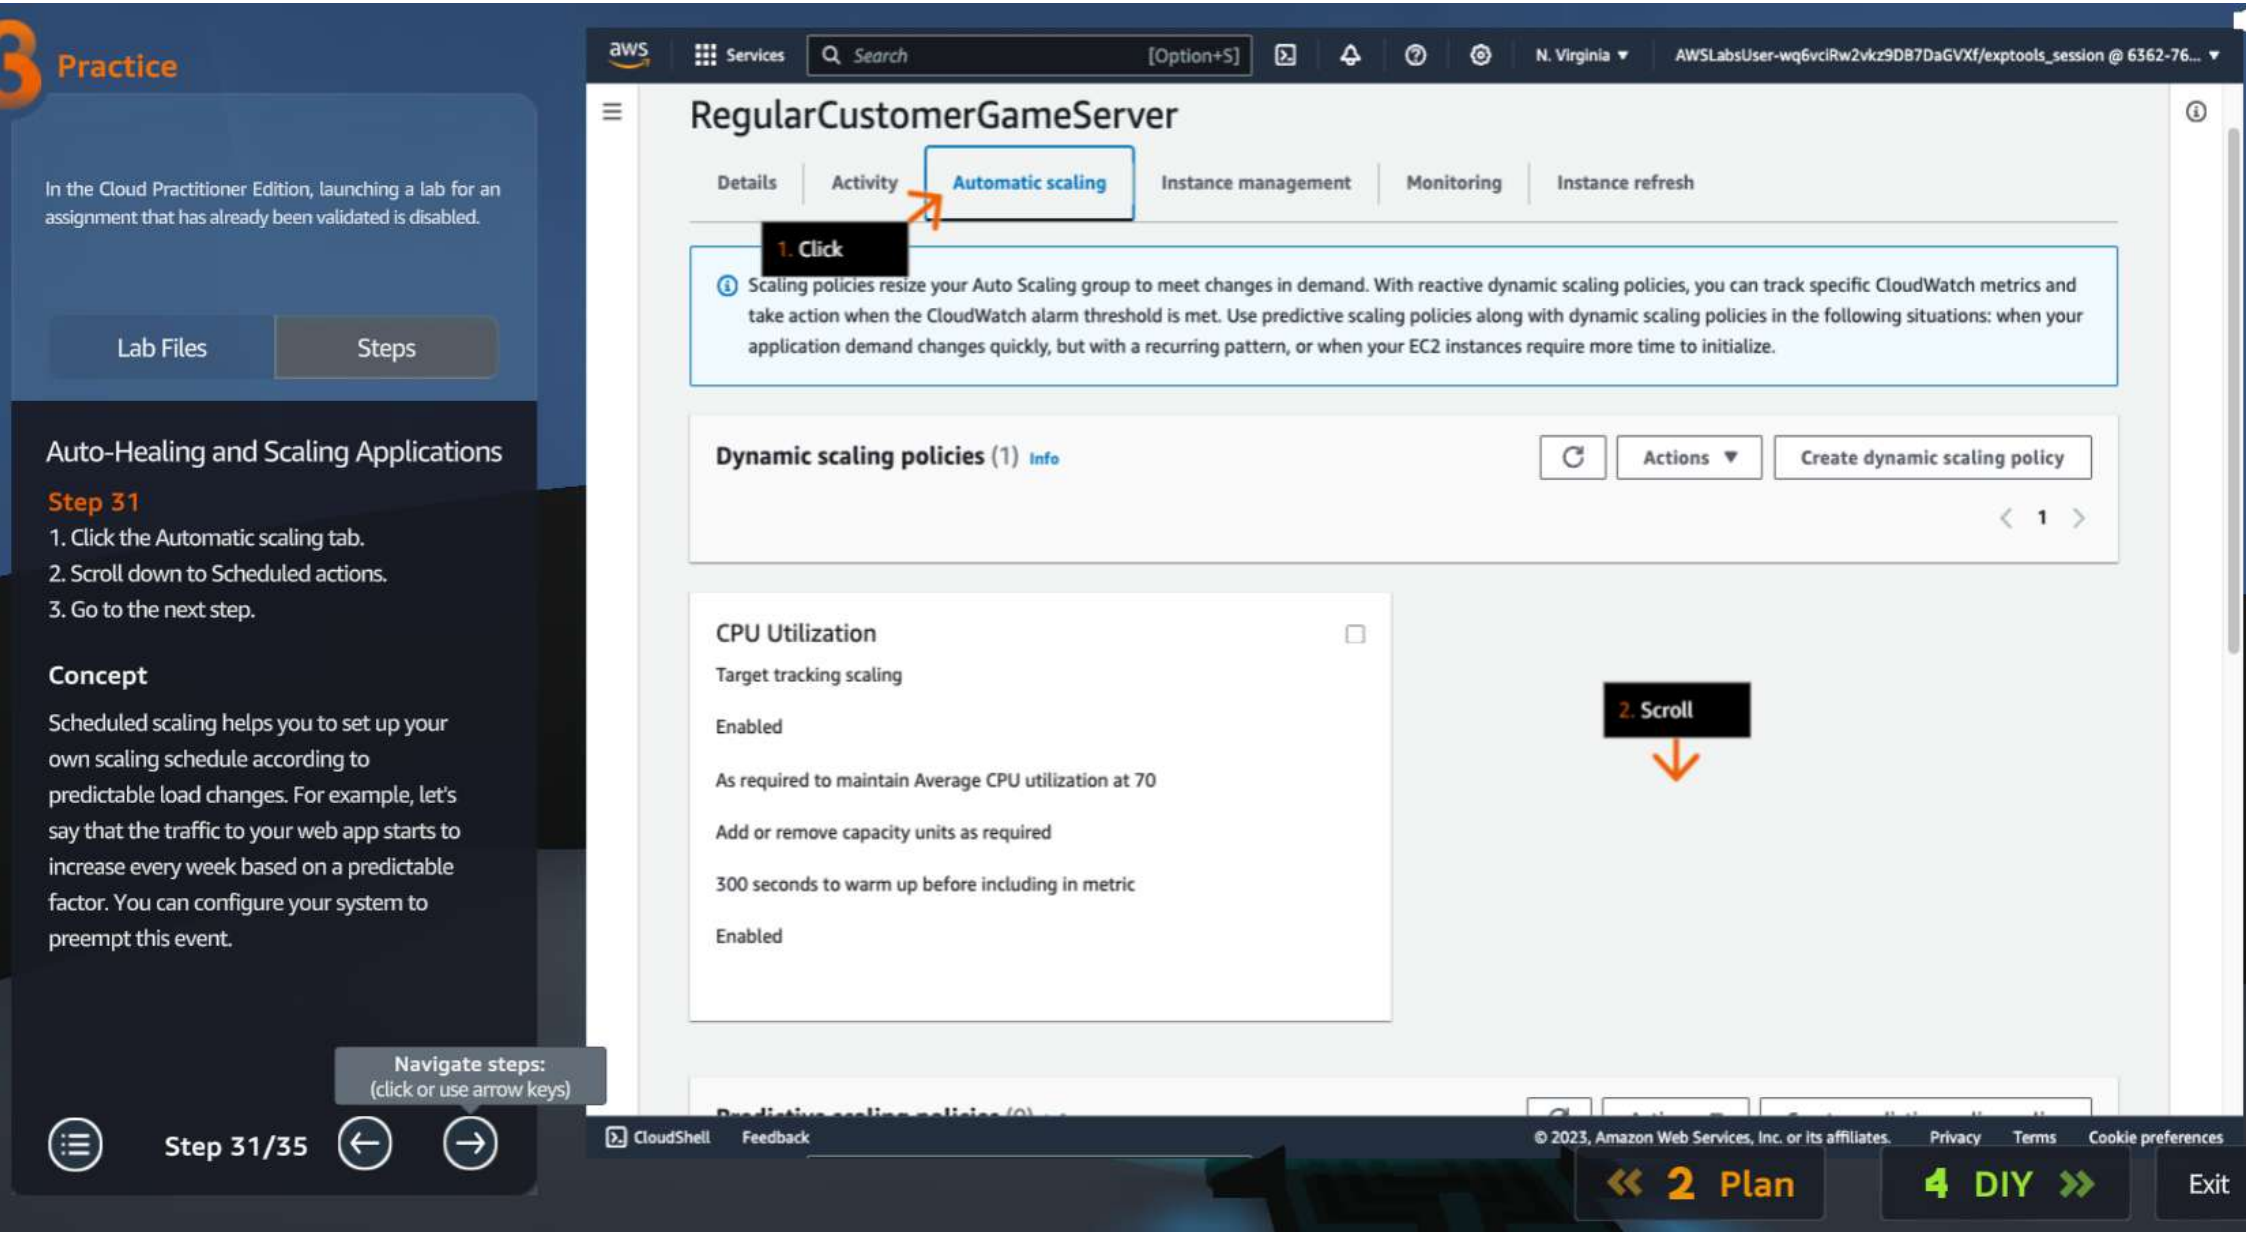
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 31
1. Click the Automatic scaling tab.
2. Scroll down to Scheduled actions.
3. Go to the next step.
Concept
Scheduled scaling helps you to set up your own scaling schedule according to predictable load changes. For example, let's say that the traffic to your web app starts to increase every week based on a predictable factor. You can configure your system to preempt this event.
Navigate steps: (click or use arrow keys)
Step 31/35
aws
Services
Search
[Option+S]
N. Virginia
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 6362-76...
RegularCustomerGameServer
Details
Activity
Automatic scaling
Instance management
Monitoring
Instance refresh
1. Click
Scaling policies resize your Auto Scaling group to meet changes in demand. With reactive dynamic scaling policies, you can track specific CloudWatch metrics and take action when the CloudWatch alarm threshold is met. Use predictive scaling policies along with dynamic scaling policies in the following situations: when your application demand changes quickly, but with a recurring pattern, or when your EC2 instances require more time to initialize.
Dynamic scaling policies (1) Info
Actions
Create dynamic scaling policy
1
CPU Utilization
Target tracking scaling
Enabled
As required to maintain Average CPU utilization at 70
Add or remove capacity units as required
300 seconds to warm up before including in metric
Enabled
2. Scroll
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

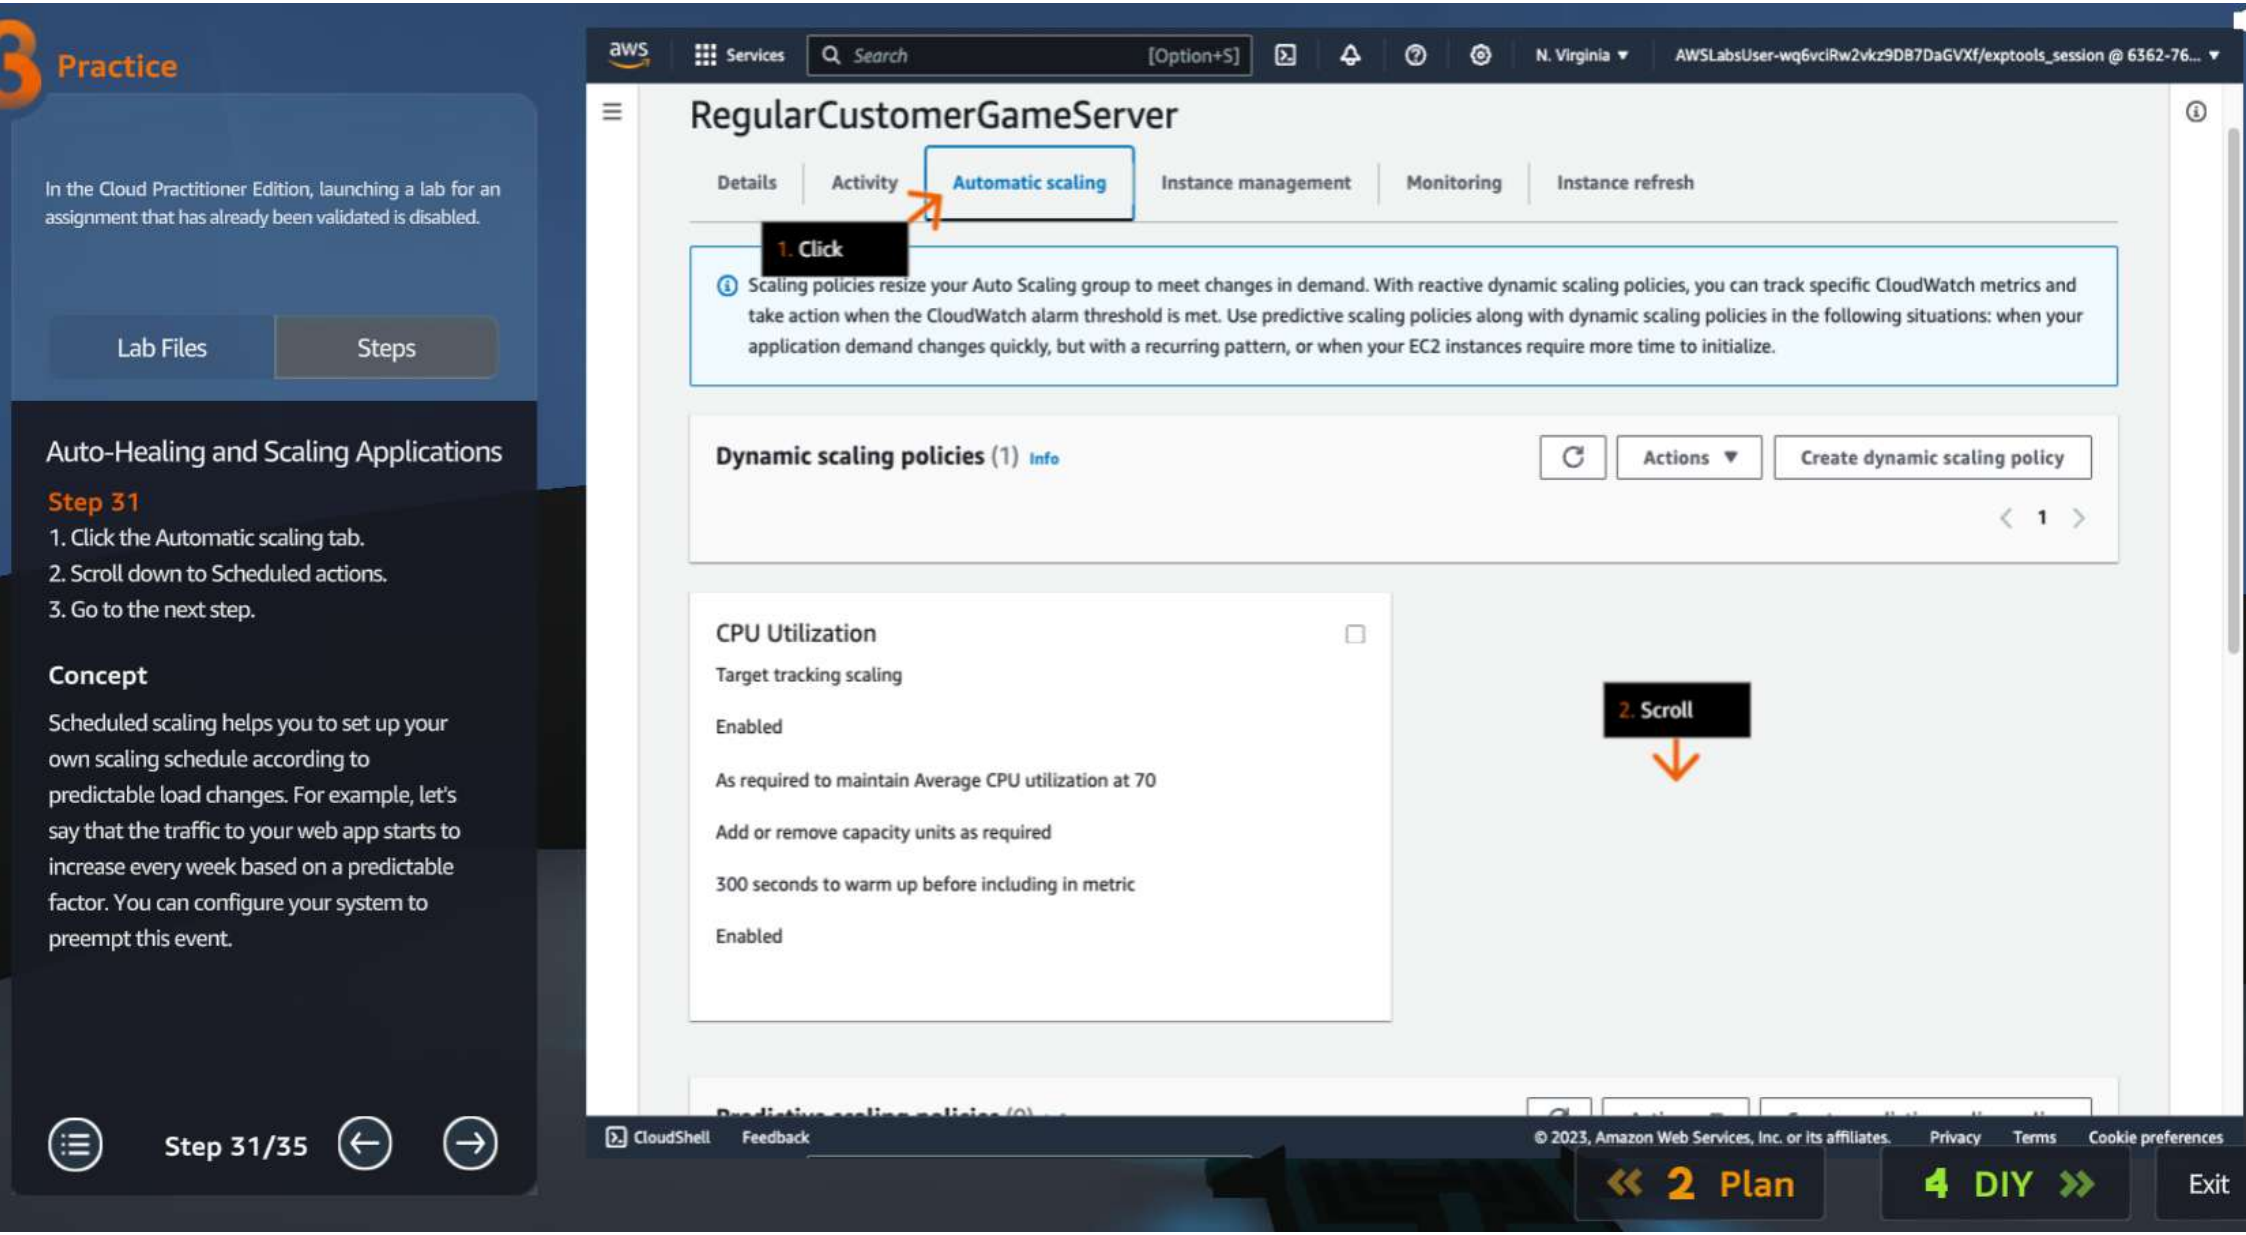
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 31
1. Click the Automatic scaling tab.
2. Scroll down to Scheduled actions.
3. Go to the next step.
Concept
Scheduled scaling helps you to set up your own scaling schedule according to predictable load changes. For example, let's say that the traffic to your web app starts to increase every week based on a predictable factor. You can configure your system to preempt this event.
Step 31/35
aws
Services
Search
[Option+S]
N. Virginia
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 6362-76...
RegularCustomerGameServer
Details
Activity
Automatic scaling
Instance management
Monitoring
Instance refresh
1. Click
Scaling policies resize your Auto Scaling group to meet changes in demand. With reactive dynamic scaling policies, you can track specific CloudWatch metrics and take action when the CloudWatch alarm threshold is met. Use predictive scaling policies along with dynamic scaling policies in the following situations: when your application demand changes quickly, but with a recurring pattern, or when your EC2 instances require more time to initialize.
Dynamic scaling policies (1) Info
Actions
Create dynamic scaling policy
1
CPU Utilization
Target tracking scaling
Enabled
As required to maintain Average CPU utilization at 70
Add or remove capacity units as required
300 seconds to warm up before including in metric
Enabled
2. Scroll
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

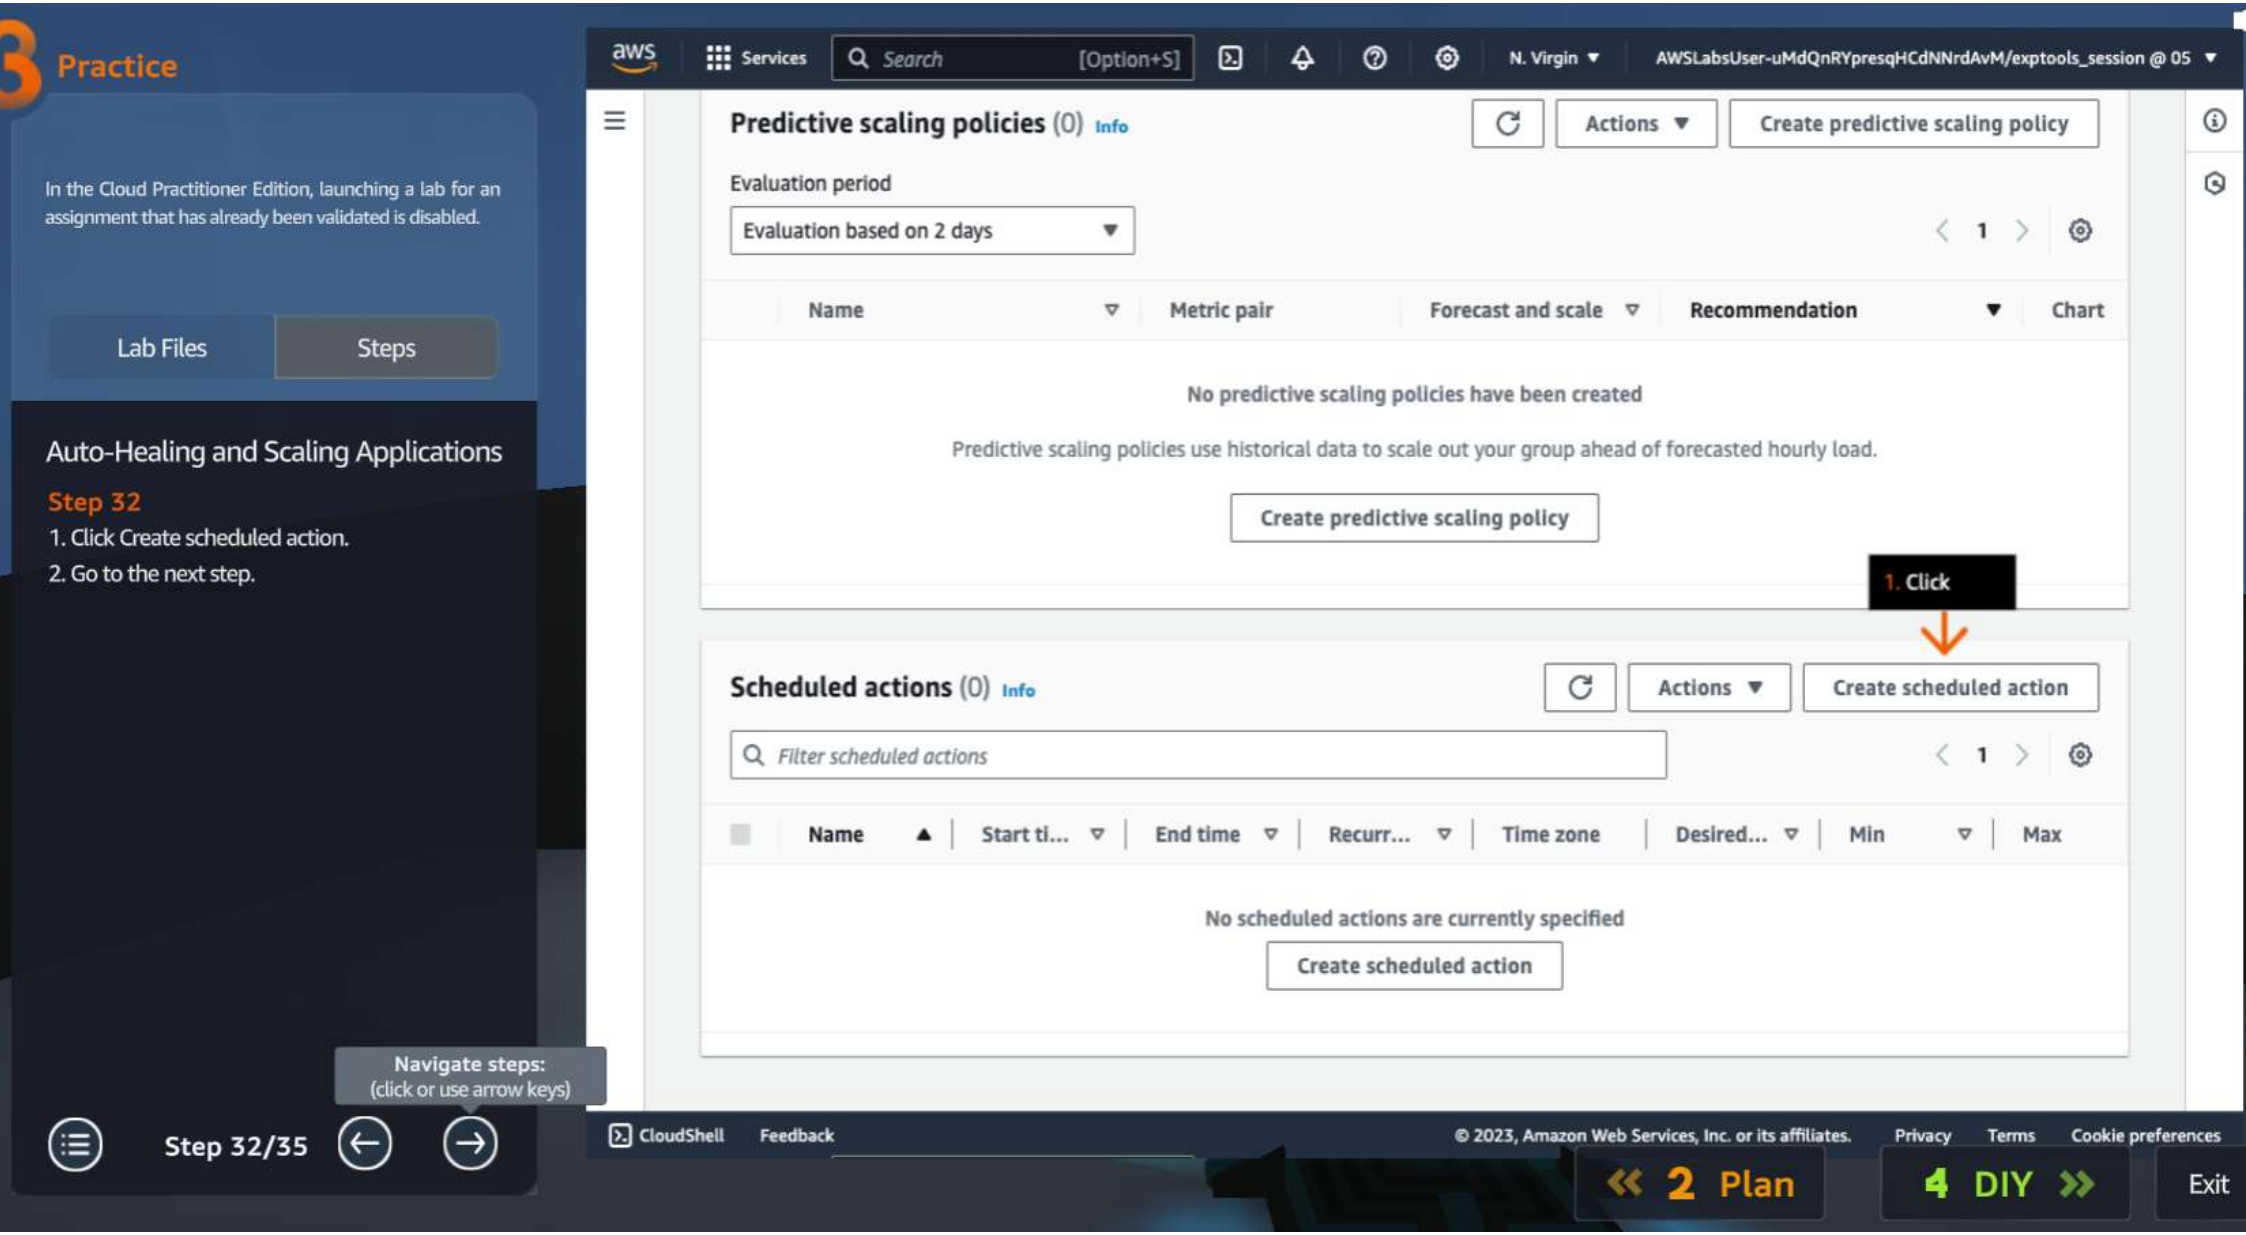
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 32
1. Click Create scheduled action.
2. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 32/35
aws
Services
Search
[Option+S]
N. Virgin
AWSLabsUser-uMdQnRYpresqHCdNNrdAvM/exptools_session @ 05
Predictive scaling policies (0) Info
Actions
Create predictive scaling policy
Evaluation period
Evaluation based on 2 days
Name
Metric pair
Forecast and scale
Recommendation
Chart
No predictive scaling policies have been created
Predictive scaling policies use historical data to scale out your group ahead of forecasted hourly load.
Create predictive scaling policy
1. Click
Scheduled actions (0) Info
Actions
Create scheduled action
Filter scheduled actions
Name
Start ti...
End time
Recurr...
Time zone
Desired...
Min
Max
No scheduled actions are currently specified
Create scheduled action
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

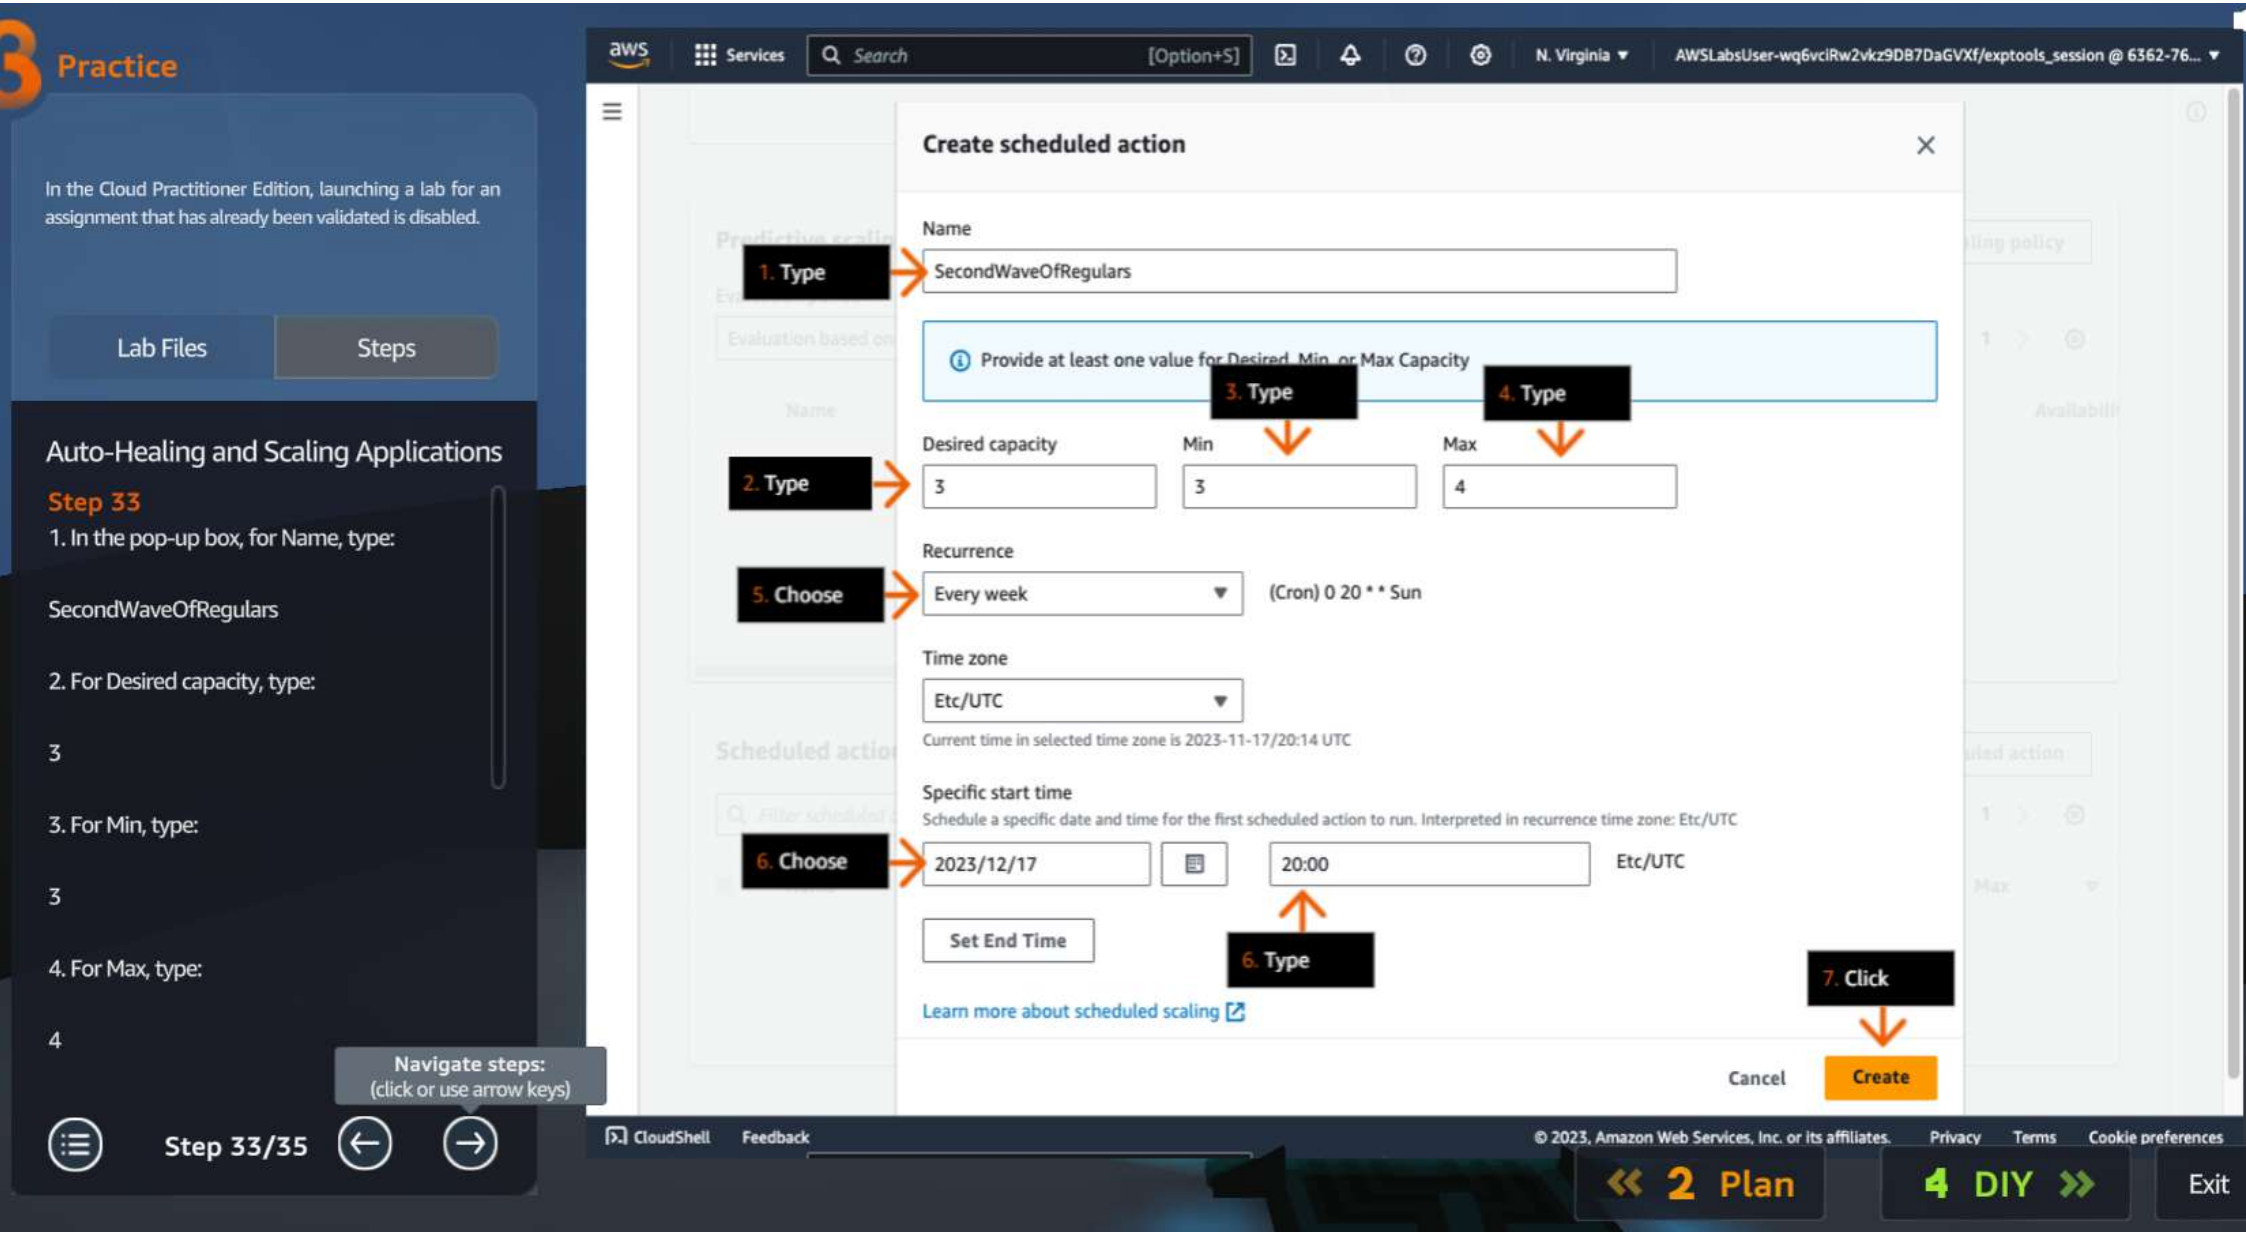
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 33
1. In the pop-up box, for Name, type:
SecondWaveOfRegulars
2. For Desired capacity, type:
3
3. For Min, type:
3
4. For Max, type:
4
Navigate steps: (click or use arrow keys)
Step 33/35
aws
Services
Search
[Option+S]
N. Virginia
AWSLabsUser-wq6vciRw2vkz9DB7DaGVXf/exptools_session @ 6362-76...
Create scheduled action
Name
1. Type
SecondWaveOfRegulars
Provide at least one value for Desired, Min, or Max Capacity
3. Type
4. Type
Desired capacity
Min
Max
2. Type
3
3
4
Recurrence
5. Choose
Every week
(Cron) 0 20 * * Sun
Time zone
Etc/UTC
Current time in selected time zone is 2023-11-17/20:14 UTC
Specific start time
Schedule a specific date and time for the first scheduled action to run. Interpreted in recurrence time zone: Etc/UTC
6. Choose
2023/12/17
20:00
Etc/UTC
Set End Time
6. Type
7. Click
Learn more about scheduled scaling
Cancel
Create
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit

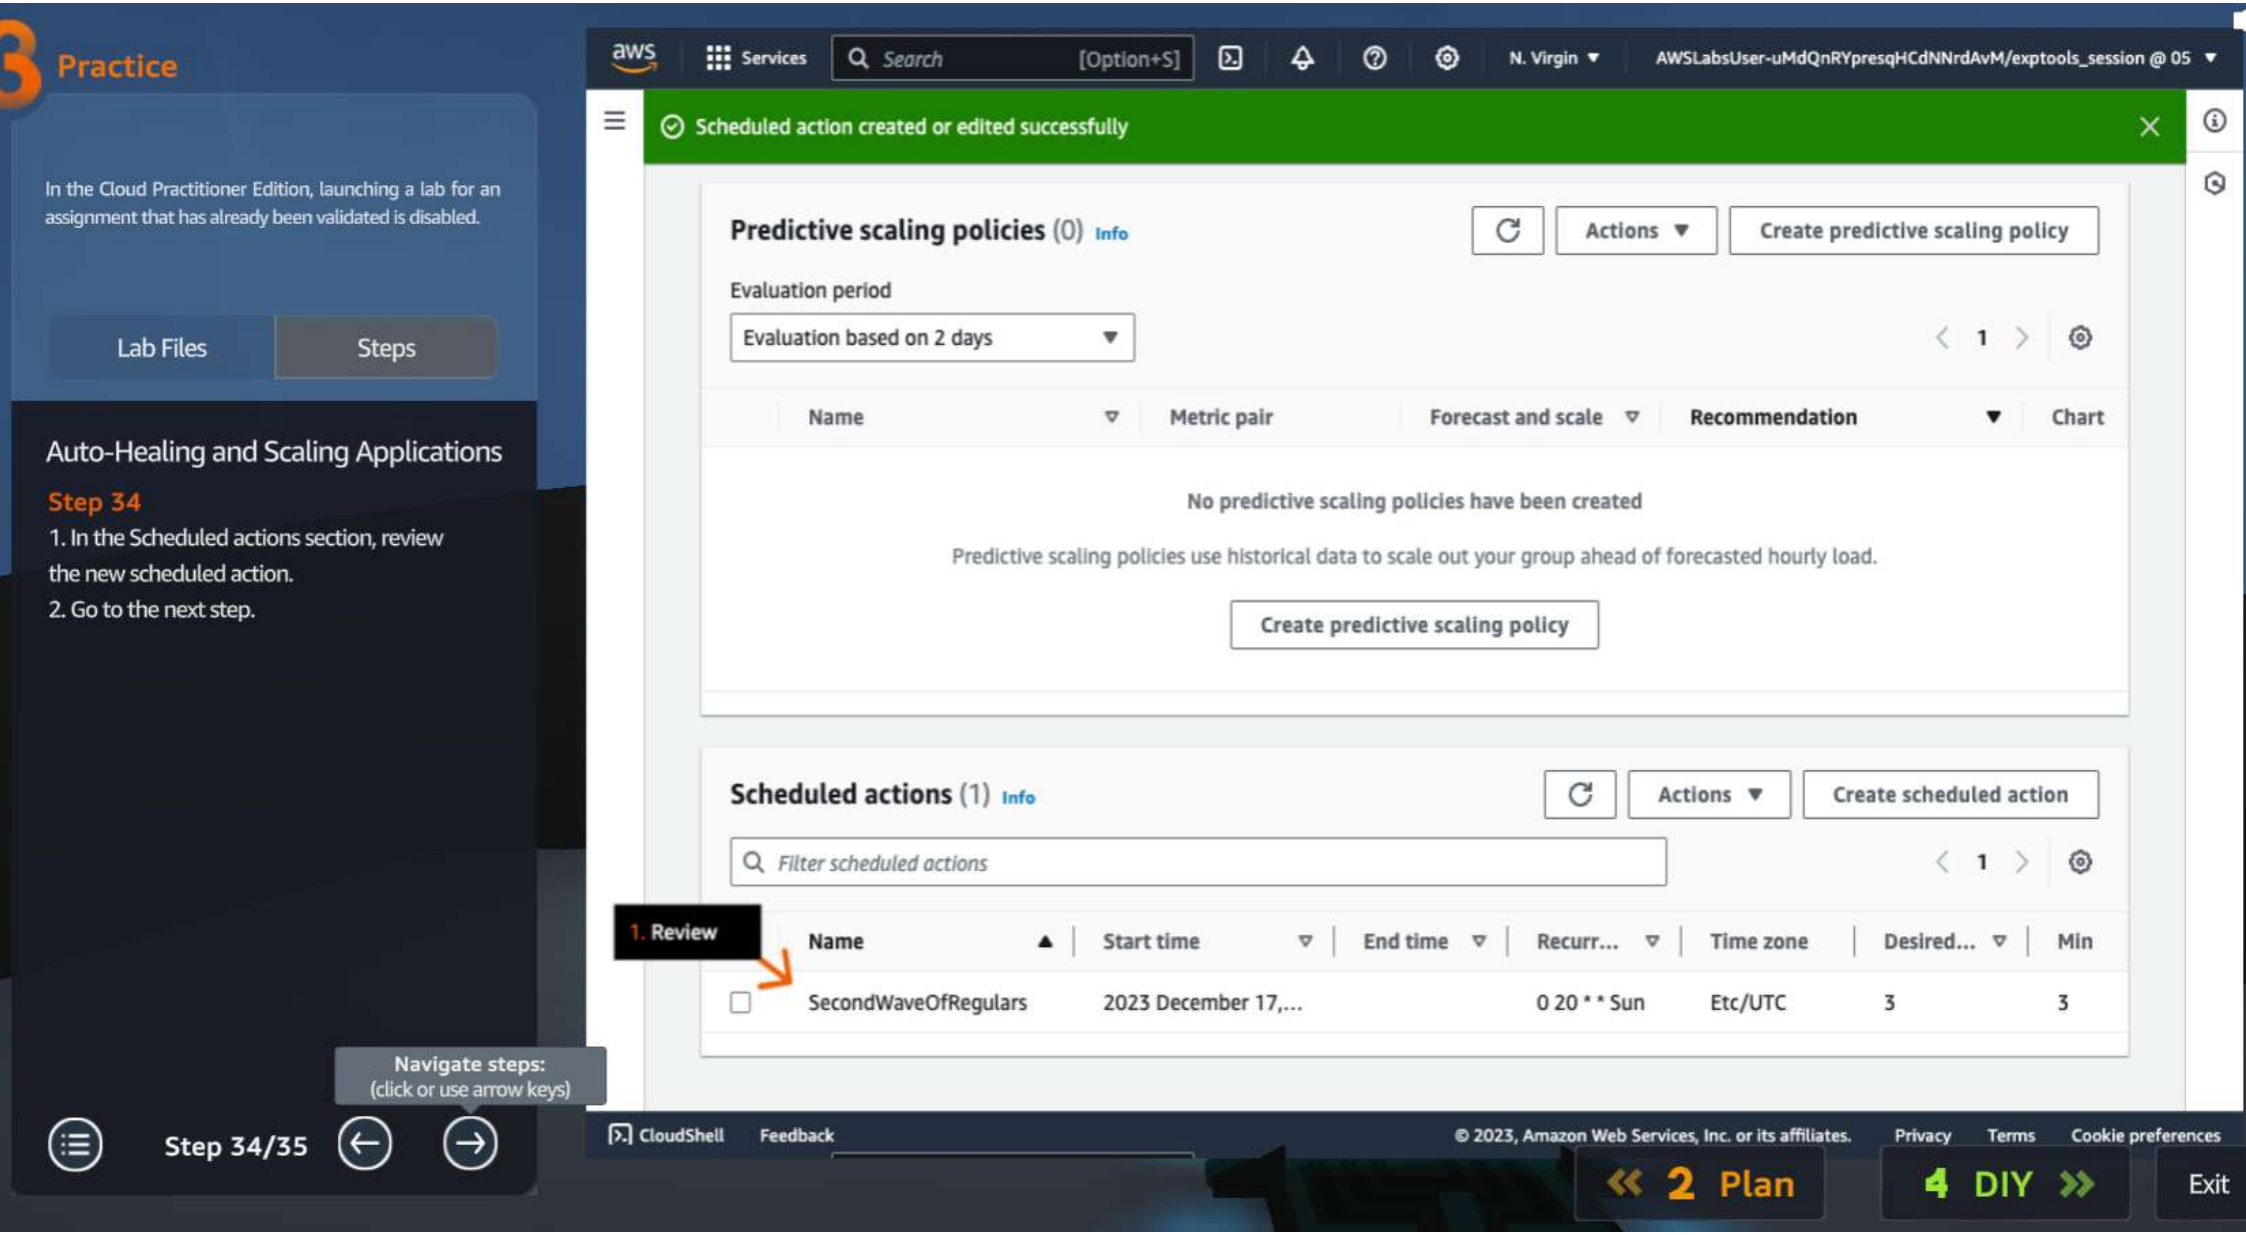
Practice
In the Cloud Practitioner Edition, launching a lab for an assignment that has already been validated is disabled.
Lab Files
Steps
Auto-Healing and Scaling Applications
Step 34
1. In the Scheduled actions section, review the new scheduled action.
2. Go to the next step.
Navigate steps: (click or use arrow keys)
Step 34/35
aws
Services
Search
[Option+S]
N. Virgin
AWSLabsUser-uMdQnRYpresqHCdNNrdAvM/exptools_session @ 05
Scheduled action created or edited successfully
Predictive scaling policies (0) Info
Actions
Create predictive scaling policy
Evaluation period
Evaluation based on 2 days
Name
Metric pair
Forecast and scale
Recommendation
Chart
No predictive scaling policies have been created
Predictive scaling policies use historical data to scale out your group ahead of forecasted hourly load.
Create predictive scaling policy
Scheduled actions (1) Info
Actions
Create scheduled action
Filter scheduled actions
1. Review
Name
Start time
End time
Recurr...
Time zone
Desired...
Min
SecondWaveOfRegulars
2023 December 17,...
0 20 * * Sun
Etc/UTC
3
3
CloudShell
Feedback
© 2023, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
2 Plan
4 DIY
Exit